

# Advanced SQL

Pieter Joubert

March 27, 2022

Advanced SQL



# Table of Contents

- 1 Associative Entities
- 2 Joins
- 3 Lecture summary

## NoSQL/GraphDB

Please note that the second session this week will focus on NoSQL/GraphDB, and will be given by another lecturer. Keep an eye on Canvas as I might upload some additional SQL videos there.



## Backup and Import

To be able to all start from the same page I've made a script available on Canvas that will import a clean and up-to-date version of the *todo\_list* database. The two commands below (run in the terminal or shell) export a database to a .sql file, and then import the same file into a database. Please note that the database must be created *first* when importing.

```
01 | pg_dump todolist > ~/joubertp/Documents/todo_list.sql
02 | psql todolist < ~/joubertp/Documents/todo_list.sql
```



# Section 1

## Associative Entities

# Auto Increment PRIMARY KEYS

- We have been manually inserting primary keys into our database.
- While this works, it is something that is prone to error, and something would be better to automate.
- The example below is use to make the *todo\_item\_id* into an auto-incremented primary key.

```
01 |  todo_item_id integer GENERATED ALWAYS AS IDENTITY  
    NOT NULL
```

- Note that this example would be best used when creating the table (and as such is included in the import script)



# Associative Entities

- So far we have only *one-to-many* relationships in our database.
- For example, one *owner* can own many *todo\_items*.
- As our database gets more complex we are also going to need to implement more complex relationships.
- One example of such a relationship could be a *many-to-many* relationship between *owners* and *projects*.
- We will now need to add a new table linking *owner* and *project*.



# Associative Entity SQL

- The code to create a new Associative Entity is shown below:

```
01 | CREATE TABLE owner_project (  
02 |     owner_id integer NOT NULL,  
03 |     project_id integer NOT NULL,  
04 |     PRIMARY KEY(owner_id, project_id)  
05 | );
```

- Note that we do not use auto-incremented values for the primary keys, and that we have a *compound* primary key made up of two fields.
- We can now insert some rows that link an owner to a project.





## Section 2

### Joins

- Now that we have all of these FK/PK links set up in our database we need to start using these relationships.
- What we essentially need to do is *join* together two or more tables, based on the links that we have created.
- The SQL here can get a little complex but this is where *set theory* actually becomes useful, as we essentially performing a UNION, INTERSECTION or COMPLEMENT on two or more sets (two or more tables).
- If you are unsure which kind of *join* to use, consider drawing out the relationships as a Venn diagram.



# INNER JOIN

- An INNER JOIN selects rows that match in both tables.
- For example we might join *owner* and *todo\_item* to get the owner name associated with a *todo\_item*, instead of the *owner\_id*:

```
01 | SELECT todo_item.description, owner.owner_name,  
    |        owner.owner_surname  
02 | FROM todo_item  
03 | INNER JOIN owner ON todo_item.owner_id=owner.  
    |        owner_id;
```



# Result

	description	owner_name	owner_surname
01			
02	-----+-----+-----		
03	Align telescope	Carl	Sagan
04	Align telescope	Neil	de Grasse Tyson
05	Write Report	Neil	de Grasse Tyson
06	Send Report	Neil	de Grasse Tyson
07	Send Another Report	Neil	de Grasse Tyson
08	Check Email	Neil	de Grasse Tyson
09	Check Email	Andrea	Ghaez
10	Align telescope	Carl	Sagan
11	Check telescope	Andrea	Ghaez



# INNER JOIN with all columns

- Expanding on the previous example to include all the Foreign Key Columns

```
01 | SELECT
02 |     priority.priority_name, todo_item.description
    |     , owner.owner_name, owner.owner_surname,
    |     project.project_name, status.status_name,
    |     context.context_name, due_date
03 | FROM todo_item
04 | INNER JOIN owner ON todo_item.owner_id =
    | owner.owner_id
05 | INNER JOIN project ON todo_item.project_id =
    | project.project_id
06 | INNER JOIN priority ON todo_item.priority_id
    | = priority.priority_id
07 | INNER JOIN context ON todo_item.context_id =
    | context.context_id
08 | INNER JOIN status ON todo_item.status_id =
    | status.status_id;
```



# LEFT/RIGHT JOIN

- In our example let's add one more owner to our owner table:

```
01 | INSERT INTO owner (owner_name, owner_surname)
    | VALUES ('Thomas', 'Kepler');
```

- Note that for now we don't link this new owner to a todo\_item.



# RIGHT JOIN

- The example below performs a *RIGHT* join between our `todo_item` and `owner` tables.

```
01 | SELECT todo_item.description, owner.owner_name,  
    |         owner.owner_surname  
02 | FROM todo_item  
03 | RIGHT JOIN owner ON todo_item.owner_id=owner.  
    |         owner_id;
```

- Note how we now have additional owner, that is not linked to a `todo_item`.



# FULL JOIN

- The example below performs a *FULL* join between our `todo_item` and `owner` tables.

```
01 | SELECT todo_item.description, owner.owner_name,
    |      owner.owner_surname
02 | FROM todo_item
03 | FULL JOIN owner ON todo_item.owner_id=owner.
    |      owner_id;
```

- In this example our results still looks the same. Remember that our `todo_item` has a foreign key relationship with the `owner` table, so a `todo_item` *must* have an owner and our *FULL* join will not have an additional `todo_item` entry.





- We've just looked at some basic examples of JOINS.
- The correct type of JOIN to use depends on the database structure you've design.
- Most often an *INNER JOIN* will solve your problem, but keep in mind that there are other options when performing a JOIN.
- And finally, we can combine any of the other functions and SQL code with a JOIN to create more complex and focused queries.



## Section 3

### Lecture summary

# Lecture summary

- Associative Entities
- JOINS

**Thank you! Questions?**