

Concurrency in Java

OOP week 10 lectures

Dr. **Madasar** Shah

November 2021

Course Overview

This Semester:

- Object Oriented Programming (Prof Reddy, Dr Chetty)
- **Programming Concurrency in Java**

This week:

- Motivating concurrency, Processes vs Threads, common pitfalls
- Thread programming, more pitfalls, concurrency patterns

Question: in all our programs so far statements have run sequentially. How do we run two or more statements in parallel to each other?

Motivation 0: a simple server in Java

- Example server without concurrency :
`socketserver.SimpleSocketServer`
- Usage: `SimpleSocketServer <port>`
- Set a port number in IntelliJ 'Edit configurations', near the ▶ button

Question: can we connect more than once to this server?


Question: can more than one person connect to the server simultaneously?

Processes vs Threads: Motivation I

Our programs **can access the entire memory space** allocated to the process.



Before: **un threaded, non-concurrent programs:** Follow the sequence of statements we write in-order, **one statement at a time, to modify memory.**




Processes vs Threads: Motivation II

Problem: in certain situations, we **want more than one statement to be active at any one time** within our process. e.g. GUI app or web server apps.



Solution: **multi threaded, concurrent programs:** follow a few specific design patterns to **allow concurrent statements to execute in parallel** in one program.



Caveat: concurrent programming **requires more design and development skill** than non-concurrent programming.

Concurrent Programming: Threads

Question: What are threads?

Concurrent Programming: Threads

What are threads?

- One means to **achieve concurrent processing**
- A standard **library** in Java (`java.lang.Thread`)
- An OO **design pattern**: *Active Object*
- A **conventional API** in the operating-system kernel, also: **architectural pattern**
- A **feature** supported by many CPU designs
- Take **extra effort to develop**, test and debug (non-deterministic)
- **Difficult to reason** about mechanically
- In order to write threaded programs we must **discuss the compilation of our non-threaded code**

Concurrent Programming: Key Challenge

Always consistent modification of shared memory using **synchronised access** to variables and objects.

Shared data and concurrent processing

- **Threads** used to achieve parallel execution
- Threads can be **interrupted**
- **Non-atomic operations on shared data become interleaved when using multiple threads**
- Result: data can become corrupted

Non-atomic operations

`balance++` is a **non-atomic**, **stack-based** operation.

- `balance++` compiled as:
 - **load** `balance` into cpu register `r`
 - **increment** register `r`
 - **store** `r` into `balance`

Inherent problem with concurrent processing that cannot be solved in hardware.
Software-based **synchronisation** of resource updates needed.

Example: `iplusplus`

Case Study: Concurrent Cashier System

Let's start with an example of buggy, concurrent code:

- Account.java (one object, shared by cashiers)
- Cashier.java (zero or more *active* objects)
- Manager.java (one object)

Shared Data Structure: Account.java

```
1 class Account {  
2  
3     private int balance = 0;  
4  
5     public void increment() {  
6         balance++;  
7     }  
8     ...  
9 }
```

try it on codesnip!

Non-solution: `volatile` keyword

- `volatile` keyword: variable reads and variable writes become atomic
- declare `int` `balance` as `volatile`

```
1 class Account {  
2  
3     private volatile int balance = 0;  
4     ...  
5 }
```

try it on codesnip!

Does not fix the issue in this code- why?

Possible solution 2a: `synchronized` keyword

- `synchronized` keyword: marks whole method as a *critical section*
- declare `increment()` and `decrement()` as `synchronized`

```
1 class Account {  
2     ...  
3     public synchronized void increment()  
4     ...  
5     public synchronized void decrement()  
6     ...  
7 }
```

try it on codesnip!

Fixed!

Problem: overkill for all but the most simple methods.

Possible solution 2b: `synchronized` block

- `synchronized` block: marks selected statements as a *critical section*
- mark `balance++` and `balance--` as synchronized

```
1 class Account {  
2     ...  
3     public void increment(){  
4         synchronized(this){  
5             balance++;  
6         }  
7     }  
8     ...  
9 }
```

try it on codesnip!

Question: what is `this` and how is it an implicit *lock*?

Problem: `synchronized` not suitable with multiple exclusive variables.

Possible solution 3: Lock object

- Create an **explicit** Lock object for each shared resource
- lock() and unlock() around balance++ and balance--

```
1  class Account {  
2  
3      private Lock balanceLock = new ...  
4  
5      public void increment(){  
6          balanceLock.lock();  
7          balance++;  
8          balanceLock.unlock();  
9      }  
10     ...  
11 }
```

try it on codesnip!

result: fine-grained access control, standard Java syntax, lock per exclusive variables

Next time: more on concurrency patterns in Java

- Writing threaded programs in Java
- Anti-patterns: spinlock, deadlock, starvation, race conditions, live lock
- Patterns: Active Object, Producer-Consumer, ThreadPool patterns and libraries

End of Section

- Questions?

With thanks to Prof. Martín Escardó.