

# Week 1 Note

---

## Introduction

---

### Class Plan

- **Introduction/definition of learning**
- **Learning approaches for classification**
  - Logistic Regression
  - Nonlinear Transformations
  - Support Vector Machines
- **Learning approaches for regression:**
  - Linear regression and its closed form solution
  - Support Vector Regression
- **Kernels as similarity functions**
  - Other kernel-based methods
- **Optimisation algorithms:**
  - Gradient Descent and its weaknesses
  - Newton Raphson and Iterative Reweighted Least Squares
  - Sequential Minimal Optimisation
- **Fundamentals:**
  - Is learning feasible? Learning theory
  - Bias and variance
  - Generalisation, overfitting and regularisation
  - Validation and model selection

### Assessment

- Graded quizzes(**20%**)
  - Two Summative canvas quizzes worth 10% each
  - Timed for 1h max
  - Can start at any time between the release and due times
  - Your assignment will be automatically submitted at the due time, even if you started less than 1h before
  - deadline is **strict**
  - Marks and feedback released in the week after the quiz
- Exam(**80%**)
  - Main summer assessment period(May maybe)

## Office Hours

- Room UG39
- Friday 14:00
- Friday 15:00

## Supervised Learning

---

- Learns a mapping from inputs  $\vec{x} = (x_1, \dots, x_d)^T \in X$  to outputs  $y \in Y$ , given a training set of input-output pairs  $J = \{(\vec{x}^1, y^1), (\vec{x}^2, y^2), \dots, (\vec{x}^n, y^n)\}$

$$J = \{(\vec{x}^i, y^i)\}_{i=1}^N$$

- The Output Space  $y$ 
  - **Regression:**  $y = R$
  - **Classification:**  $Y$  is a set of categories
    - 2 categories: binary classification
    - 2 categories: multi-class classification
- The Input Space  $X$ 
  - $d$ -dimensional space, where each dimension can be:
    - Numeric
    - Ordinal
    - Categorical

Use Training Examples  $J = \{(\vec{x}^i, y^i)\}_{i=1}^N$  and Hypothesis Set  $h(\vec{x}) = \vec{a}^T \vec{x} + b, \forall \vec{a} \in \mathbb{R}^d, b \in \mathbb{R}$  to Learn Algorithm and get Final Hypothesis  $g \approx f$

- Problem Setting
  - Given a set of training examples  $J = \{(\vec{x}^1, y^1), (\vec{x}^2, y^2), \dots, (\vec{x}^n, y^n)\}$  where  $(\vec{x}^i, y^i) \in X \times Y$  are drawn from a fixed albeit unknown joint probability distribution  $p(\vec{x}, y) = p(y|\vec{x})p(\vec{x})$
  - Goal: to learn a function  $g : X \rightarrow Y$  able to generalise to unseen (test) examples of the same probability distribution  $p(\vec{x}, y)$ 
    - $g : X \rightarrow Y$ , mapping input space to output space
    - $g$  as a probability distribution approximating  $P(y|\vec{x})$
  - Generalisation: minimise  $E(g) = P_{(\vec{x}, y) \sim p[g(\vec{x}) \neq y]}$  or  $E(g) = E[(g(\vec{x}) - y)^2]$

## Logistic Regression

- In Logistic Regression, we will model the probability (actually the odds) of an instance to belong to a given class as a linear combination of the inputs.
- Odds and Logit

- **Odds**: ratio of probabilities of two possible outcomes:

$$o_1 = \frac{p_1}{p_0} = \frac{p_1}{1 - p_1}$$

If  $o_1 \geq 1$ , predict class 1  
 If  $o_1 < 1$ , predict class 0

- **Logit(aka. log-odds)**: the logarithm of the odds:

$$\text{Logit}(p_1) = \vec{w}^T \vec{x}$$

$$\text{where } \text{logit}(p_1) = \ln\left(\frac{p_1}{1-p_1}\right)$$

Logit enables us to map from **[0,1]** to  $[-\infty, \infty]$   
 If  $\text{logit}(p_1) \geq 0$ , predict class 1  
 If  $\text{logit}(p_1) < 0$ , predict class 0

- In the case above, we know:

$$p_1 = \frac{e^{\vec{w}^T \vec{x}}}{1 + e^{\vec{w}^T \vec{x}}}$$

$$p_0 = 1 - p_1 = \frac{1}{1 + e^{\vec{w}^T \vec{x}}}$$

$$h(\vec{x}) = p_1 = p(1|\vec{x}, \vec{w}), \forall \vec{w} \in \mathbb{R}^{d+1}$$

## Likelihood

- **Likelihood function**

$$\prod_{i=1}^N P_{y^i} = \prod_{i=1}^N p(y^i | x^i; \vec{w}) = p(\vec{y} | \vec{X}, \vec{w}) = L(\vec{w}) \quad (1)$$

$$= \prod_{i=1}^N p(1|\vec{x}^i, \vec{w})^{(y_i)} (1 - p(1|\vec{x}^i, \vec{w}))^{(1-y_i)} \quad (2)$$

这一段使用了Bernoulli distribution进行转换

把所有例子  $\vec{X}$  与对应的结果  $\vec{y}$ , 根据代求的参数  $\vec{w}$  全部汇总为一个式子  $L(\vec{w})$   
 $\vec{x}$  代表一个例子所需要的输入,  $\vec{X}$  代表所有例子的输入(汇总)

$$\vec{X} = \begin{pmatrix} x_1^1 & x_2^1 & \dots & x_d^1 \\ x_1^2 & x_2^2 & \dots & x_d^2 \\ \dots & \dots & \dots & \dots \\ x_1^N & x_2^N & \dots & x_d^N \end{pmatrix}$$

$$\vec{y} = \begin{pmatrix} y^1 \\ y^2 \\ \dots \\ y^N \end{pmatrix}$$

- **Log-Likelihood**

$$\ln(L(\vec{w})) = \ln \prod_{i=1}^N P_{y^i} = \sum_{i=1}^N \ln P_{y^i}$$

把求  $\arg \max_w L(\vec{w})$  转换为  $\arg \max_w \ln L(\vec{w})$

作用:  $\ln$  函数单调递增, 使数值更为稳定(stable)

- **Loss Function**

$$E(\vec{w}) = -\ln(L(\vec{w})) = -\sum_{i=1}^N \ln P_{y^i} \quad (3)$$

$$= -\sum_{i=1}^N y^i \ln p(1|\vec{x}^i, \vec{w}) + (1 - y^i) \ln(1 - p(1|\vec{x}^i, \vec{w})) \quad (4)$$

把求  $\arg \max_w \ln L(\vec{w})$  转换为  $\arg \min_w E(\vec{w})$

作用: 把问题转换为求最小值问题(而非最大值), 符合尝试

## Gradient Descent

- Gradient descent adjusts  $\vec{w}$  iteratively in the direction that leads to the biggest decrease (steepest descent) in  $E(\vec{w})$ .

$$x := x - \eta \frac{df}{dx}$$

即

$$\vec{w} = \vec{w} - \eta \nabla E(\vec{w})$$

where  $\eta > 0$  and  $\nabla E(\vec{w}) = \sum_{i=1}^N (p(1|\vec{x}^i, \vec{w}) - y^i) \vec{x}^i$

$\eta$  是学习率(Learning rate), 是超参数(Hyperparameter), 由使用者决定  
过高的  $\eta$  可能导致无法找到最低点 过低的  $\eta$  可能导致找的效率变低

- Gradient descent is a general purpose optimisation algorithm. But is likely to get stuck in local minima

For **logistic regression** using **cross-entropy loss**, this is **not a problem**, as its  $E(\vec{w})$  is strictly convex with respect to  $\vec{w}$ , having a single unique minimum.

## Addition

### Equivalent Terms

- $\vec{x}$ : input attribute, input feature, independent variable, input variable.
- $y$ : output attribute, output variable, dependent variable, label (for classification).
- mapping: learned function, predictive model, classifier (for classification).
- Learning a model, training a model, building a model.
- $J$ : set of training examples, training data.
- $(\vec{x}, y)$ : example, observation, data point, instance (more frequently used for examples with unknown outputs).
- Different people and books will use different notations!

## Notation

- Scalar: lower case, e.g.,  $b$
- Column Vector: lower case, bold, e.g.,  $\vec{x}$
- Vector element: lower case with subscript, e.g.,  $x_i$
- Matrix: upper case, bold, e.g.,  $X$
- Matrix element: upper case with subscripts, e.g.,  $X_{i,j}$
- If enumerating these (e.g., having multiple vectors), superscript will be used to differentiate this from indices, e.g.,  $\vec{x}^i$