


Week 3 Note

Support Vector Machines(SVMs)



- The perpendicular distance γ between the decision boundary and the closest training example is called the margin
- The decision boundary can be chosen so as to maximise the margin
- Training examples that are exactly on the margin are called support vectors
- Given a set of training examples

$$J = \{(\vec{x}^{(1)}, y^{(1)}), (\vec{x}^{(2)}, y^{(2)}), \dots, (\vec{x}^{(N)}, y^{(N)})\}$$

where $(\vec{x}^{(i)}, y^{(i)}) \in X \times Y$ are drawn from a fixed albeit unknown joint probability distribution $P(\vec{x}, y) = P(y|\vec{x})P(\vec{x})$

- Goal: to learn a function g able to generalise to unseen(test) examples of the same probability distribution $P(\vec{x}, y)$
 - $g: X \rightarrow Y$, mapping input space to output space
 - g as a probability distribution approximating $P(y|\vec{x})$

Hypothesis Set

$$h(\vec{x}) = \begin{cases} +1 & \text{if } \vec{w}^T \vec{x} + b > 0 \\ -1 & \text{if } \vec{w}^T \vec{x} + b < 0 \end{cases}, \forall \vec{w} \in \mathbb{R}^d, \forall b \in \mathbb{R}$$

- Perpendicular Distance From a Point $\vec{x}^{(n)}$ to a Hyperplane $h(\vec{x}) = 0$

$$\text{dist}(h, \vec{x}^{(n)}) = \frac{|h(\vec{x}^{(n)})|}{\|\vec{w}\|} = \frac{y^{(n)} h(\vec{x}^{(n)})}{\|\vec{w}\|}$$

where $\|\vec{w}\| = \sqrt{\vec{w}^T \vec{w}}$ is the Euclidean norm (the length of the vector \vec{w})

- Find \vec{w} and b that maximise the margin
- Constraint: all training examples must be correctly classified

$$\begin{aligned} & \min_n \text{dist}(h, \vec{x}^{(n)}) \\ & \downarrow \\ & \arg \max_{\vec{w}, b} \{ \min_n \text{dist}(h, \vec{x}^{(n)}) \} \end{aligned}$$

Constraint:



$$\text{Subject to } y^{(n)} h(\vec{x}^{(n)}) > 0, \forall (\vec{x}^{(n)}, y^{(n)}) \in J$$

$$\begin{aligned} & \arg \max_{\vec{w}, b} \{ \min_n \left(\frac{y^{(n)} h(\vec{x}^{(n)})}{\|\vec{w}\|} \right) \} \\ & \downarrow \\ & \arg \max_{\vec{w}, b} \left\{ \frac{1}{\|\vec{w}\|} \min_n (y^{(n)} h(\vec{x}^{(n)})) \right\} \end{aligned}$$

Constraint:

$$\begin{aligned} & \text{Subject to } y^{(n)} h(\vec{x}^{(n)}) > 0, \forall (\vec{x}^{(n)}, y^{(n)}) \in J \\ & \text{Subject to } \min_n y^{(n)} h(\vec{x}^{(n)}) = 1, \forall (\vec{x}^{(n)}, y^{(n)}) \in J \end{aligned}$$

- Why are these two constraints equivalent?
 - Rescaling \vec{w} and b does not change the position of the hyperplane, nor the distances of the training examples to it



- If there is a hyperplane that can separate the training examples, its \vec{w} and b can be divided by $\min_n y^{(n)} h(\vec{x}^{(n)})$ so that $y^{(n)} h(\vec{x}^{(n)}) = 1$ for the closet example

$$\arg \max_{\vec{w}, b} \left\{ \frac{1}{\|\vec{w}\|} \right\}$$

$$\downarrow$$

$$\arg \min_{\vec{w}, b} \{\|\vec{w}\|\}$$

Constraint:

Subject to $\min_n y^{(n)} h(\vec{x}^{(n)}) = 1, \forall (\vec{x}^{(n)}, y^{(n)}) \in J$ stricter

Subject to $y^{(n)} h(\vec{x}^{(n)}) \geq 1, \forall (\vec{x}^{(n)}, y^{(n)}) \in J$ looser

The optimal solution will satisfy the equality in $y^{(n)} h(\vec{x}^{(n)}) \geq 1$ for at least one training example

$$\arg \min_{\vec{w}, b} \left\{ \frac{1}{2} \|\vec{w}\|^2 \right\}$$

Constraint:

Subject to $y^{(n)} (\vec{w}^T \phi(\vec{x}^{(n)}) + b) \geq 1, \forall (\vec{x}^{(n)}, y^{(n)}) \in J$

Convexity

- Convex Sets

- A set C is convex if the line segment between any two points in C lies in C
- For any two points $\vec{x}^{(1)}, \vec{x}^{(2)} \in C$ and any $\lambda \in (0, 1)$, we have:

$$\lambda \vec{x}^{(1)} + (1 - \lambda) \vec{x}^{(2)} \in C$$





- Convex Functions

- A convex function $f(\vec{x})$ is a function with a convex domain C that satisfies the following condition for any $\vec{x}^{(1)}, \vec{x}^{(2)} \in C$ and $\lambda \in (0, 1)$

$$f\left(\lambda\vec{x}^{(1)} + (1 - \lambda)\vec{x}^{(2)}\right) \leq \lambda f\left(\vec{x}^{(1)}\right) + (1 - \lambda)f\left(\vec{x}^{(2)}\right)$$

- Strictly convex: satisfies the condition with $<$



- Importance of Convexity in Machine Learning/Optimisation
 - Any minimum in a convex function is a global minimum
 - A strictly convex function has at most one stationary (critical) point. If such a point exists, it is a global minimum


- Concave: A function $f(\vec{x})$ is concave if $-f(\vec{x})$ is convex

- First-Derivative Characterisation of Convexity

- A differentiable function $f(\vec{x})$ is convex iff its domain C is convex and it satisfies the following condition for any pair $\vec{x}^{(0)}, \vec{x} \in C$

$$f(\vec{x}) \geq \underbrace{f(\vec{x}^{(0)}) + \nabla f(\vec{x}^{(0)}) \cdot (\vec{x} - \vec{x}^{(0)})}_{\text{Equation of the tangent line}}$$

- Strictly convex: satisfies the condition with $>$ for any $\vec{x}^{(1)} \neq \vec{x}^{(2)}$



- Second-Derivative Characterisation of Convexity

- A twice differentiable function $f(\vec{x})$ is convex iff:
 - Its domain C is a convex set and
 - Its Hessian $H_f(\vec{x})$ is positive semidefinite for all $\vec{x} \in C$
- If a twice differentiable function $f(\vec{x})$
 - has a convex set C as its domain and
 - its Hessian $H_f(\vec{x})$ is positive definite for all $\vec{x} \in C$
- It is a strictly convex function.(sufficient but not necessary condition)

- First-Order(Partial) Derivatives

- (First-order) derivatives tell us the rate of change of $f(x)$ as we increase x

$$\frac{d}{dx} f(x) = \frac{df}{dx} = f'(x) = f^{(1)}(x)$$

- (First-order) partial derivatives tell us the rate of change of $f(\vec{x})$ as we increase a specific variable x_i

$$\frac{\partial f}{\partial x_i}$$

- (Partial) derivatives tell us whether $f(\vec{x})$ is increasing /decreasing (along a specific axis) and how rapidly

- Second-Order(Partial) Derivatives

- Second-order derivative: This is the derivative of the derivative of $f(x)$, denoted as $\frac{d^2 f(x)}{dx^2}$. In simpler terms, it gives the rate of change of the slope $f'(x)$.

$$\frac{d^2}{dx^2} f(x) = \frac{d}{dx} \left(\frac{df}{dx} \right) = \frac{d^2 f}{dx^2} = f''(x) = f^{(2)}$$

- `Second-order partial derivative`: This is the partial derivative of the partial derivative of $f(x)$. It shows the rate of change of the slope along a specific axis, relative to that same axis or another one.

$$\begin{aligned} \frac{\partial^2 f}{\partial x_i^2} &= \frac{\partial}{\partial x_i} \left(\frac{\partial f}{\partial x_i} \right) \\ \frac{\partial^2 f}{\partial x_i \partial x_j} &= \frac{\partial}{\partial x_i} \left(\frac{\partial f}{\partial x_j} \right) \end{aligned}$$

- Hessian - Matrix of Second-Order Partial Derivatives

- Consider $f(\vec{x})$, where $\vec{x} = (x_0, x_1, \dots, x_d)^T$

$$H(f(\vec{x})) = H_f(\vec{x}) = \begin{bmatrix} \frac{\partial^2 f}{\partial x_0^2} & \frac{\partial^2 f}{\partial x_0 \partial x_1} & \cdots & \frac{\partial^2 f}{\partial x_0 \partial x_n} \\ \frac{\partial^2 f}{\partial x_1 \partial x_0} & \frac{\partial^2 f}{\partial x_1^2} & \cdots & \frac{\partial^2 f}{\partial x_1 \partial x_n} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial^2 f}{\partial x_d \partial x_1} & \frac{\partial^2 f}{\partial x_d \partial x_2} & \cdots & \frac{\partial^2 f}{\partial x_d^2} \end{bmatrix}$$

- Univariate Case

- The function is convex iff $f''(x) \geq 0$ for all x
- If $f''(0) > 0$ for all x , a function is strictly convex (sufficient but not necessary condition)
- If $f'(x) = 0$ and $f''(x) > 0$, then x is a (local) minimum (sufficient but not necessary condition)
- The function is concave iff $f''(x) \leq 0$ for all x
- If $f''(0) < 0$ for all x , a function is strictly concave (sufficient but not necessary condition)
- If $f'(x) = 0$ and $f''(x) < 0$, then x is a (local) maximum (sufficient but not necessary condition)

- Multivariate Case

- The function is convex iff $H_f(\vec{x}) \geq 0$ (positive semidefinite) for all \vec{x}
- If $H_f(\vec{x}) > 0$ (positive definite) for all \vec{x} , a function is strictly convex (sufficient but not necessary condition)
- if $\nabla f(\vec{x}) = 0$ and $H_f(\vec{x}) > 0$, then x is a (local) minimum (sufficient but not necessary condition)

- Positive Semidefinite Matrix

- A $d \times d$ symmetric matrix A is positive semidefinite iff for any non-zero vector $\vec{z} \in \mathbb{R}^d$, the following is true:

$$\vec{z}^T A \vec{z} \geq 0$$

e.g.

$$\vec{z}^T A \vec{z} = \begin{pmatrix} z_1 & z_2 \end{pmatrix} \begin{pmatrix} 1 & 0 \\ 0 & 1 \end{pmatrix} \begin{pmatrix} z_1 \\ z_2 \end{pmatrix} = \begin{pmatrix} z_1 & z_2 \end{pmatrix} \begin{pmatrix} z_1 \\ z_2 \end{pmatrix} = z_1^2 + z_2^2$$

Satisfying the above with $>$ defines a "positive definite" matrix

- Second-Derivative Characterisation of Convexity

- A twice differentiable function $f(\vec{x})$ is convex iff:
 - its domain C is a convex set and
 - its Hessian $H_f(\vec{x})$ is positive semidefinite for all $\vec{x} \in C$
- For any \vec{z}, \vec{x} , we have $\vec{z}^T H_f(\vec{x}) \vec{z} \geq 0$

- Eigenvalues and Eigenvectors

- The eigenvalues of H capture the direction of the principal curvatures of the function $f(\vec{x})$, where the curvature is most pronounced
- The eigenvalues of H capture the curvature itself
- If all eigenvalues are ≥ 0 , the curvature is always positive, "upwards"
- The eigenvalues are ≥ 0 iff $H_f(\vec{x}) \geq 0$

The Dual Representation for SVM

Dual representation of SVM

- Primal Representation

$$\arg \min_{\vec{w}, b} \left\{ \frac{1}{2} \|\vec{w}\|^2 \right\}$$

Subject to: $y^{(n)}(\vec{w}^T \phi(\vec{x}^{(n)}) + b) \geq 1 \quad \forall (\vec{x}^{(n)}, y^{(n)} \in J)$

- Dual Representation

$$\arg \max_a \tilde{L}(\vec{a}) = \sum_{n=1}^N a^{(n)} - \frac{1}{2} \sum_{n=1}^N \sum_{m=1}^N a^{(n)} a^{(m)} y^{(n)} y^{(m)} k(\vec{x}^{(n)}, \vec{x}^{(m)})$$

where: $k(\vec{x}^{(n)}, \vec{x}^{(m)}) = \phi(\vec{x}^{(n)})^T \phi(\vec{x}^{(m)})$

Subject to: $a^{(n)} \leq 0, \forall n \in \{1, \dots, N\} \quad \sum_{n=1}^N a^{(n)} y^{(n)} = 0$

Kernel trick

- There is a way to compute $k(\vec{x}^{(n)}, \vec{x}^{(m)}) = \phi(\vec{x}^{(n)})^T \phi(\vec{x}^{(m)})$ without having to ever compute $\phi(x)$. This is called the Kernel Trick
- This calculation can be generalised to basis expansions composed of all terms of order up to p

$$k(\vec{x}, \vec{z}) = \phi(\vec{x})^T \phi(\vec{z}) = (1 + \vec{x}^T \vec{z})^p$$

- Mercer's Condition
 - Consider any finite set of points $\vec{x}^{(1)}, \dots, \vec{x}^{(M)}$ (not necessarily the training set)
 - Gram matrix: An $M \times M$ similarity matrix K , whose elements are given by $K_{i,j} = k(\vec{x}^{(i)}, \vec{x}^{(j)})$
 - Mercer's condition states that K must be symmetric and positive semidefinite.
 - Symmetric: $k(\vec{x}^{(i)}, \vec{x}^{(j)}) = k(\vec{x}^{(j)}, \vec{x}^{(i)})$
 - Positive semidefinite: $\vec{z}^T K \vec{z} \geq 0 \quad \forall \vec{z} \in \mathbb{R}^M$

If these conditions are satisfied, the inner product defined by the kernel in the feature space respects the properties of inner products.

- Given valid kernels $k_1(x, z)$ and $k_2(x, z)$, the following will also be valid kernels:
 - $k(x, z) = ck_1(x, z)$
 - where $c \geq 0$ is a constant.
 - $k(x, z) = f(x)k_1(x, z)f(z)$
 - where $f(\cdot)$ is any function.
 - $k(x, z) = q(k_1(x, z))$
 - where $q(\cdot)$ is a polynomial with non-negative coefficients.
 - $k(x, z) = e^{k_1(x, z)}$
 - $k(x, z) = k_1(x, z) + k_2(x, z)$
 - $k(x, z) = k_1(x, z)k_2(x, z)$

- Gaussian kernel, a.k.a. Radial Basis Function (RBF) kernel

$$k(\vec{x}, \vec{x}^{(n)}) = e^{-\frac{\|\vec{x} - \vec{x}^{(n)}\|^2}{2\sigma^2}}$$

The embedding ϕ is infinite dimensional