

NLP Tasks and Applications Using Word Embeddings End-to-end problem solving

Venelin Kovatchev

Lecturer in Computer Science

v.o.kovatchev@bham.ac.uk

Outline

- NLP Applications
- Using embeddings. Compositionality.
- End to end neural models

NLP Applications

Classical NLP Pipeline and Features

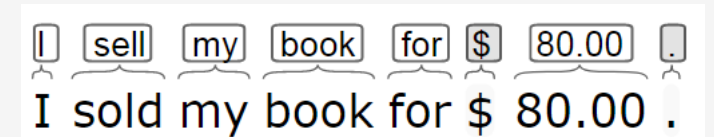
What we have learned so far

- Tokenization
- Pos tagging
- Chunking / Syntactic parsing / Dependency parsing
- Word co-occurrence and distributional semantic models
- Word embeddings

“Linguistic” tasks: Text tokenization and POS tagging

- Tokenization: output is a text segmented in tokens

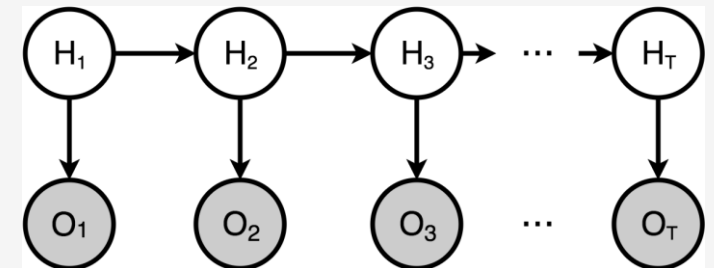
- Regular expressions, BPE



- POS tagging: output is a sequence of hidden states:

- noun, verb, adjective

- Hidden Markov Models (HMM)



“Linguistic” tasks: Chunking and constituent analysis

- Grouping words together based on their shared “function” in the text
- Find all groups that “function together” in the sentence
- I went to the movies with a friend who I know from high school.
 - [I] went to the movies with a friend who I know from high school.
 - I [went] to the movies with a friend who I know from high school.
 - I went [to the movies] with a friend who I know from high school.
 - I went to the movies [with a friend who I know from high school] .

“Linguistic” tasks: The problem of syntax

“What combinations can we get with the constituents “dog”, “human”, and “bites”

- “Dog bites human” – (statistically) most common
- “Human bites dog” – meaningful, possible, but unlikely
- “Bites dog human”, “Human dog bites”, etc. – ungrammatical
- Same constituents, different rules -> different (im-)possible complex expression



"Linguistic" tasks: Full syntactic parsing

- "Colorless green ideas sleep furiously in love"

S -> NP VP

NP -> A N

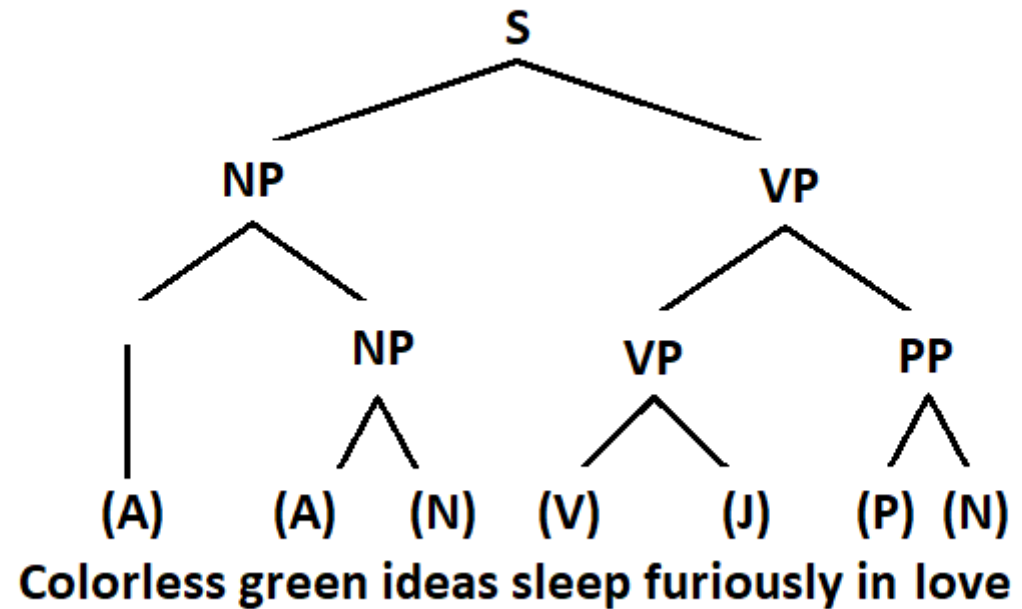
NP -> A NP

VP -> V J

VP -> VP PP

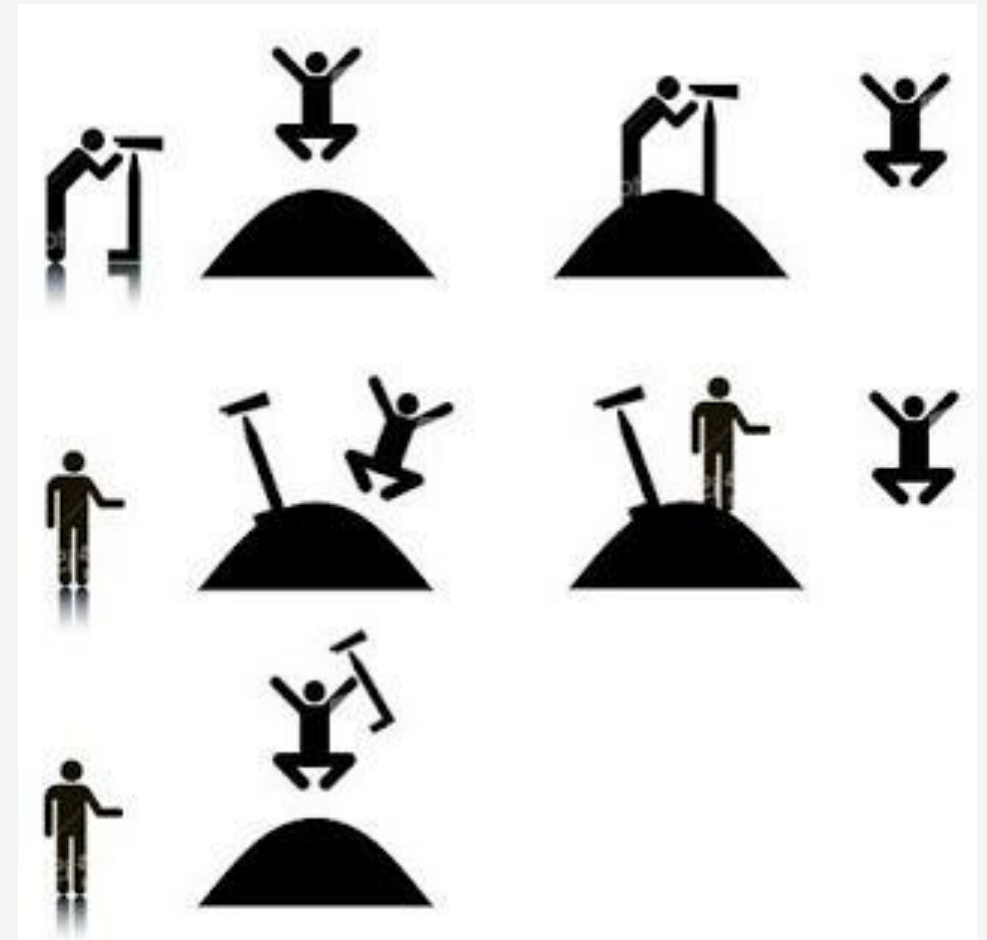
PP -> P N

(NP -> N)



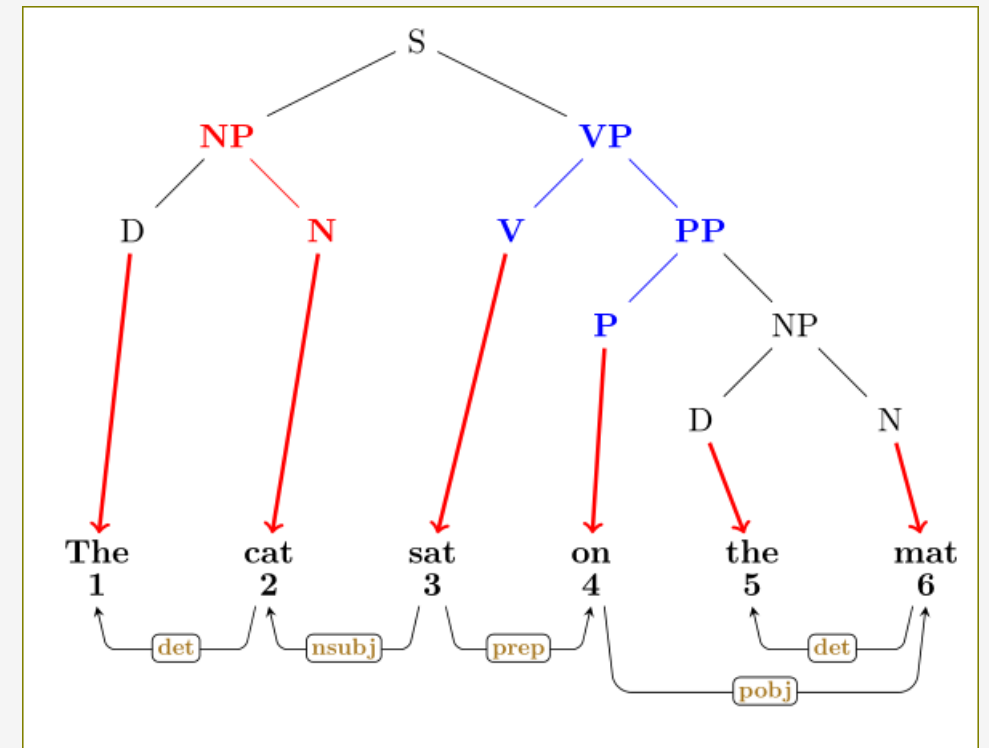
Syntactic ambiguity

- Sentences are often ambiguous
- How many interpretations for the following sentence:
 - "I saw a man in the park with the telescope"



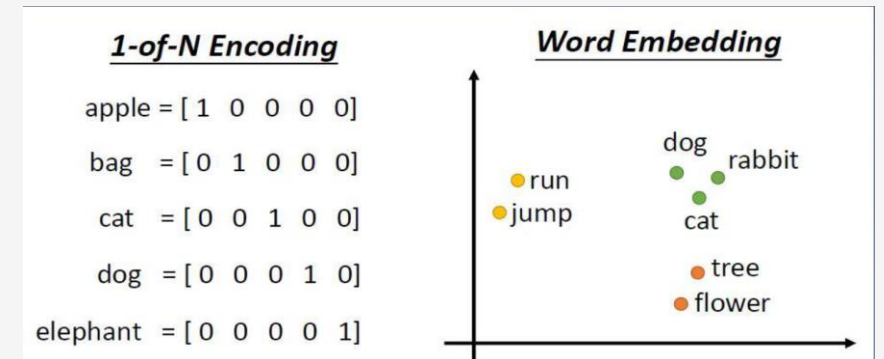
“Linguistic” tasks: Dependency parsing

- The – (det) -> cat
- cat – (subj) -> sat
- on – (prep) -> sat
- mat – (pobj) -> on
- the – (det) mat



"Linguistic" tasks: word embeddings

- What do the words "mean"?
- How can we "measure" the meaning?
- Distributional semantics:
 - the meaning is a function of the context
- Word vectors: one-hot, count based, embeddings



“Linguistic” NLP tasks

- Why do we need linguistic tasks?
 - To help computers make sense of language
 - Pre-processing and feature extraction
 - To help humans make sense of language
 - Linguistic and cognitive science experiments
 - For some problems “linguistic” NLP tasks are end goals

Linguistic tasks and machine learning

- Linguistic tasks are a goal on their own
- Linguistic tasks are tools in the (classical) NLP toolbox
- The bi-direction interaction between ML and linguistic tasks and data
 - Many linguistic tasks require ML
 - The output of linguistic analysis is used in practical applications
 - Word embeddings and transfer learning

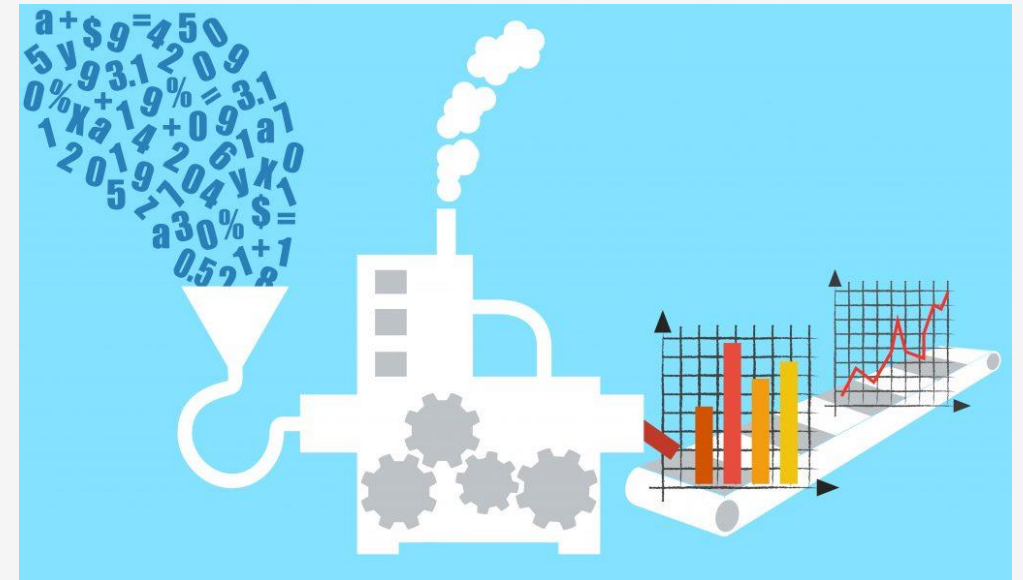
Everyday tasks that use language / NLP

- Marketing
 - Sentiment analysis
 - Recommender systems
 - Content creation



Everyday tasks that use language / NLP (2)

- News articles, social media, and search
- Information extraction
- Question answering
- Inference and Fact checking
- Content moderation



Everyday tasks that use language / NLP

- User experience and assistance
 - Conversational Agents / Chat bots
 - Machine translation
 - Personal assistants
 - Autocomplete, copilot



Extra-linguistic tasks

- Extra-linguistic tasks are not “internal” to language
 - Part-of-speech tagging vs book recommendation
- Language can be used to solve extra-linguistic tasks (in part)
- Rest of the module: taking a more practical direction
 - How do we solve problems using language
 - How do we define (and evaluate) problems and solutions
 - Feature-based solutions vs end-to-end solutions

Text classification

Feature Engineering

Machine learning and NLP – supervised and unsupervised NLP

- Types of machine learning problems: supervised, unsupervised, reinforcement
- Pop quiz: can you name ML problems of each of the three types?
- NLP problems are predominantly (represented as) supervised
 - Text classification: Sentiment analysis, Textual Inference, Fact checking, Toxic language detection
 - Text generation: Question answering, Chatbots, Machine translation

Text classification

- Observations are independent from each other (e.g. single tweet)
- Observations are preprocessed and fed into (a trained) classifier
- The classifier assigns the correct class-label to the observation
- Classes are discrete and often disjointed:
 - an email is either spam or ham, never both
- Different types of classification: binary, multi-class, multi-label

Extra-linguistic problems and feature engineering

- Historically, we approached extra-linguistic problems via feature engineering
- Feature engineering
 - Converting language data into relevant data that is easy to process for machines
 - A “faulty” translation from “human” to “computer”
 - Loses some (often a lot of) information
 - Requires a lot of human intervention

Feature engineering

- Analyze the problem, the input, and the desired outcome
- Explore existing resources and processing techniques
- Select the most relevant features and feature-extraction methods
- Empirically test what works best

Feature engineering

- Various features can be extracted from texts
- Bag-of-words (+ tf-idf)
- N-grams
- Part-of-speech tags
- Named entities
- Sentiment words
- Stop words
- Length

Feature engineering (example)

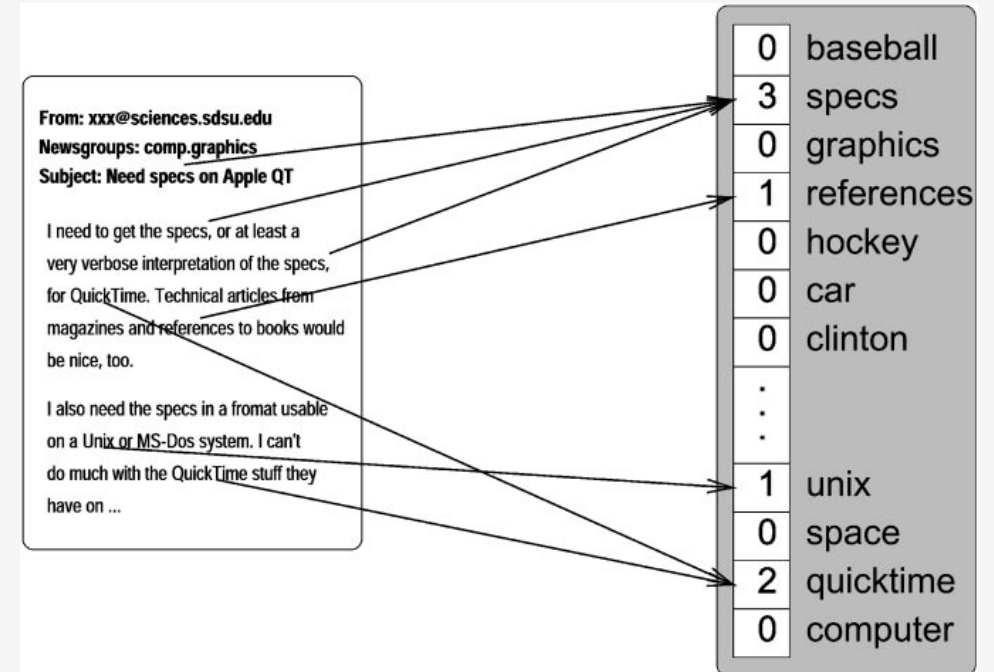
- Consider the task of sentiment analysis
- What features can you extract from the following examples?
- Are these tweets positive or negative?

1." 🌟 Absolutely in love with my new headphones from SoundWave! The sound quality is top-notch, crystal clear, and the noise cancellation is a game-changer! 🎧 ✨ Highly recommend to all music lovers out there! #SoundWave #MusicLife 👍 😊 "

2."Really disappointed with my purchase from QuickTech. The laptop crashes constantly and the battery life is a joke. 😡 💻 Worst customer service ever - they just don't care. Totally regret this buy. #QuickTechFail #Frustrated 😞 👎 "

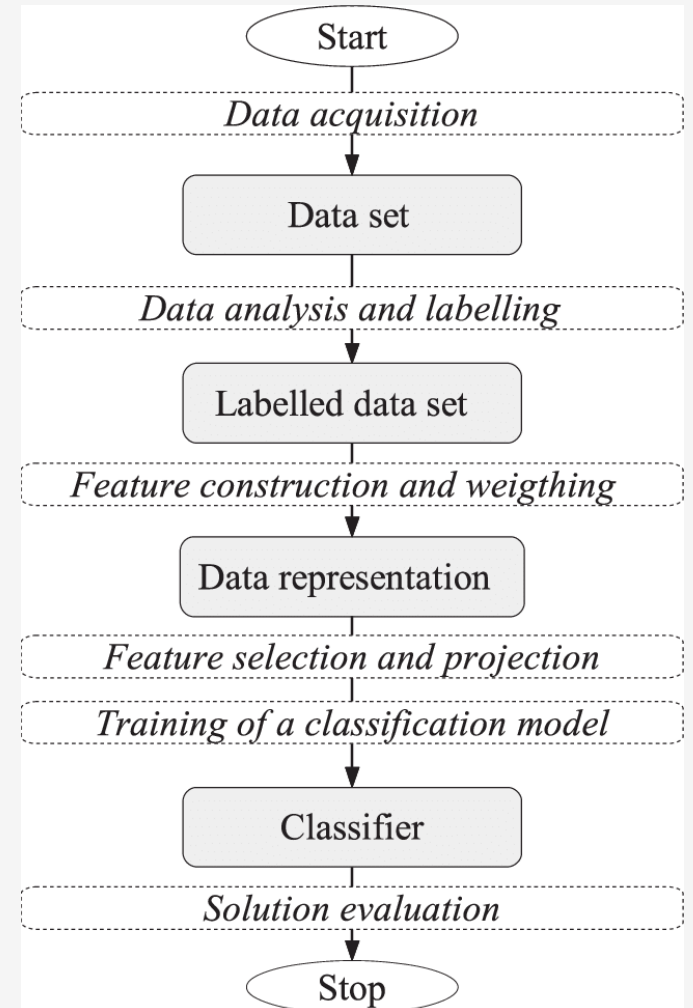
Feature extraction

- Define the set of relevant features
- Train (or program) algorithms to process the text
- Extract features
- Represent the text via the features



Text classification using features

- Step by step process
- Involves active human engagement
- Feature selection and extraction
- Data is fed into a classifier (Logistic, NB, SVM)
- Iteratively improve feature selection and model (hyper) parameters



The shift towards data-driven approaches

- DSM and embeddings extract (relevant) features from the data
 - Minimal supervision
 - Easier maintenance
- Allow computers to “read” and use language in a more direct way
- The cost is loss of control and transparency

New NLP paradigm – the rise of end-to-end neural models

- Word embeddings mark a major shift in NLP
- What do you think is an “end-to-end” neural model?
- End-to-end neural model represent the complete target system
 - No external preprocessing
 - No explicit pipelines
 - Input-output mapping
 - Training for a specific task (with some transfer learning)
- What would be some limitations of end-to-end models?



Using embeddings Compositionality

Use of word embeddings

- Embeddings are high dimensional vector representations of words
- All words in the vocabulary can be represented as a high dimensional vector
- What can we use embeddings for?
 - Write five different applications

Querying embeddings for lexical information

- Embeddings are originally a lexical resource
- Performing operations at word level
 - Find words that are semantically similar to a given word
 - Expand existing dictionaries by automatically searching for similar words
 - Explore relational similarity (e.g., UK – London: France – Paris)
 - Compare the meaning of a word to its context
 - Automatic error correction
-

Querying embeddings for lexical information

- Learning from embeddings
 - Learn specific semantic relations (e.g. hypernymy) (Shwartz, et al. 2016)
 - Learn compositionality rules (Baroni and Zampareli, 2010, Socher et al. 2013)
 - Learn representations for phrases and compare with words
- Clustering (Kovatchev et al. 2016)
 - Topics
 - Part of speech

Evaluating word embeddings

- Embeddings are trained on word prediction (or noise detection)
- Embedding algorithms converge to the best (mathematical) solution
- We throw away the original task!
- How do we evaluate embeddings?
 - Word similarity and word analogy
 - Downstream tasks

From words to text

- Embeddings represent words
- NLP is about processing text
- "The cat sat on a mat" vs "the mat sat on a cat"



Compositionality of meaning

- “The meaning of a complex expression is determined by the meanings of its constituent expressions and the rules used to combine them”
- Two key questions:
 - How do we combine individual word meaning?
 - Does the word meaning remain static?

How to combine word meaning

- Assume that the word meaning is a vector or a tensor
- How can you calculate the meaning of a phrase?
 - Addition/aggregation
 - Complex (hierarchical) operations
 - Via a deep neural network

Vector addition

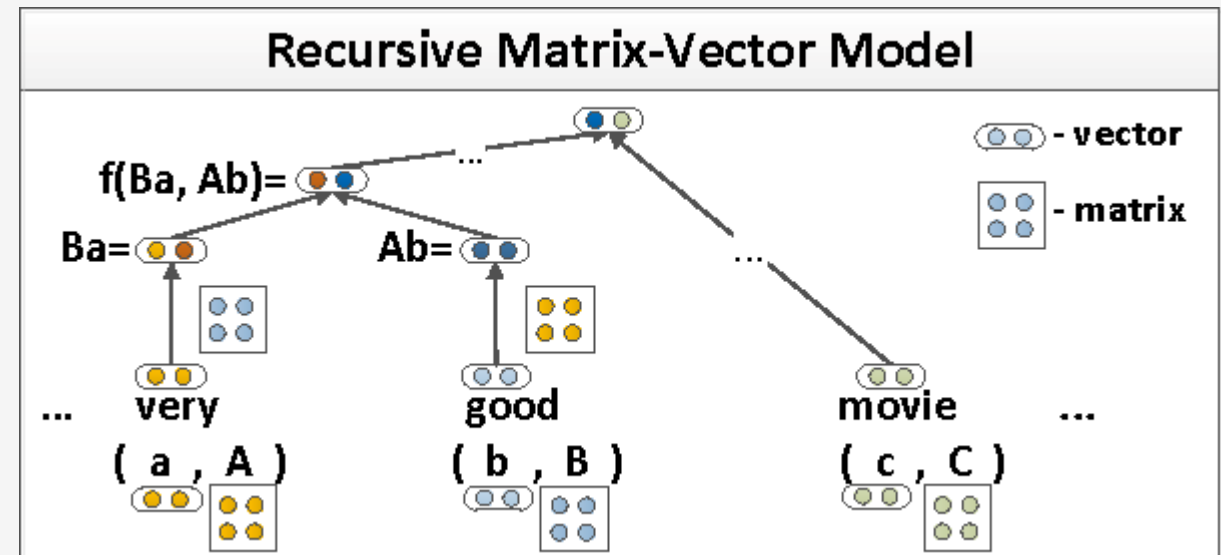
- The simplest form of compositionality
 - "The cat sat on a mat" = "The" + "cat" + "sat" + "on" + "a" + "mat"
- Advantages
 - Easy to calculate
 - Fixed vector length, regardless of text length
- Disadvantages
 - Loses word order
 - Lower impact of individual words (e.g., "not")

Vector concatenation

- Alternative to vector addition
 - Instead of adding dimensions, we concatenate the vectors
- Can you identify advantages and disadvantages of this approach?

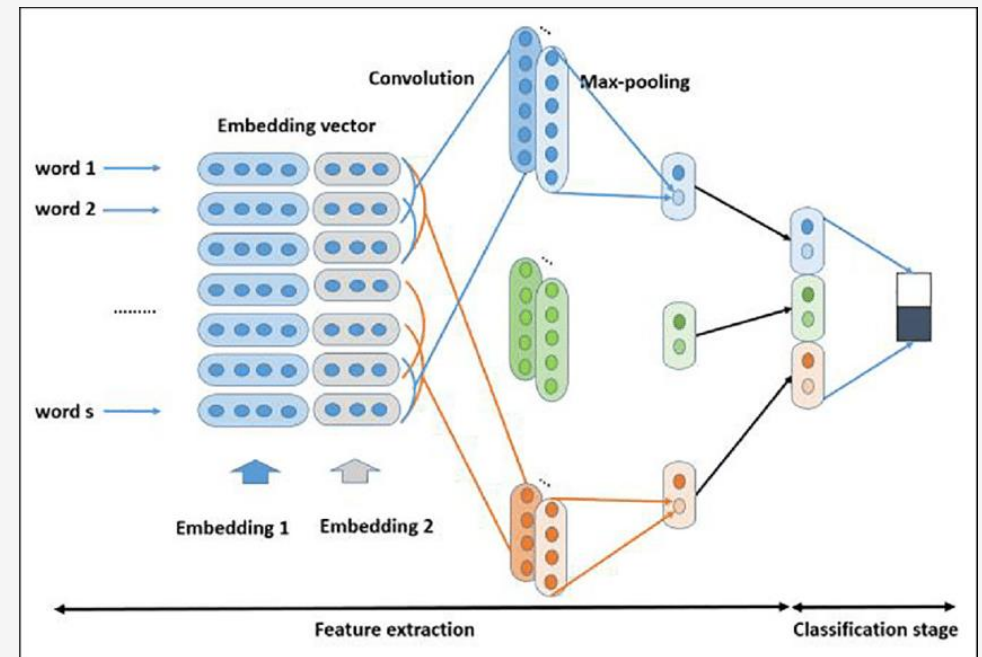
Recursive compositionality

- Baroni and Zampareli (2010), Socher et al. (2012, 2013)
- Meaning is not just a vector, but can be a vector + matrix
- Compositionality is a recursive vector-matrix operation
- Follow the syntactic structure



Compositionality via deep neural networks

- The current state-of-the-art
- Input the embeddings into a neural network
- Let the network handle the interactions
- The network architecture determines the compositionality



Task specific embeddings

- General purpose word embeddings
- Polysemy ("blue cat", "cat myfile.sh", "CAT scan")
- Retraining embeddings
 - Domain: news, medical, social media (Major et al. 2018, Soares, 2019)
 - Task: sentiment, NER (Siencnik, 2015)
 - Languages

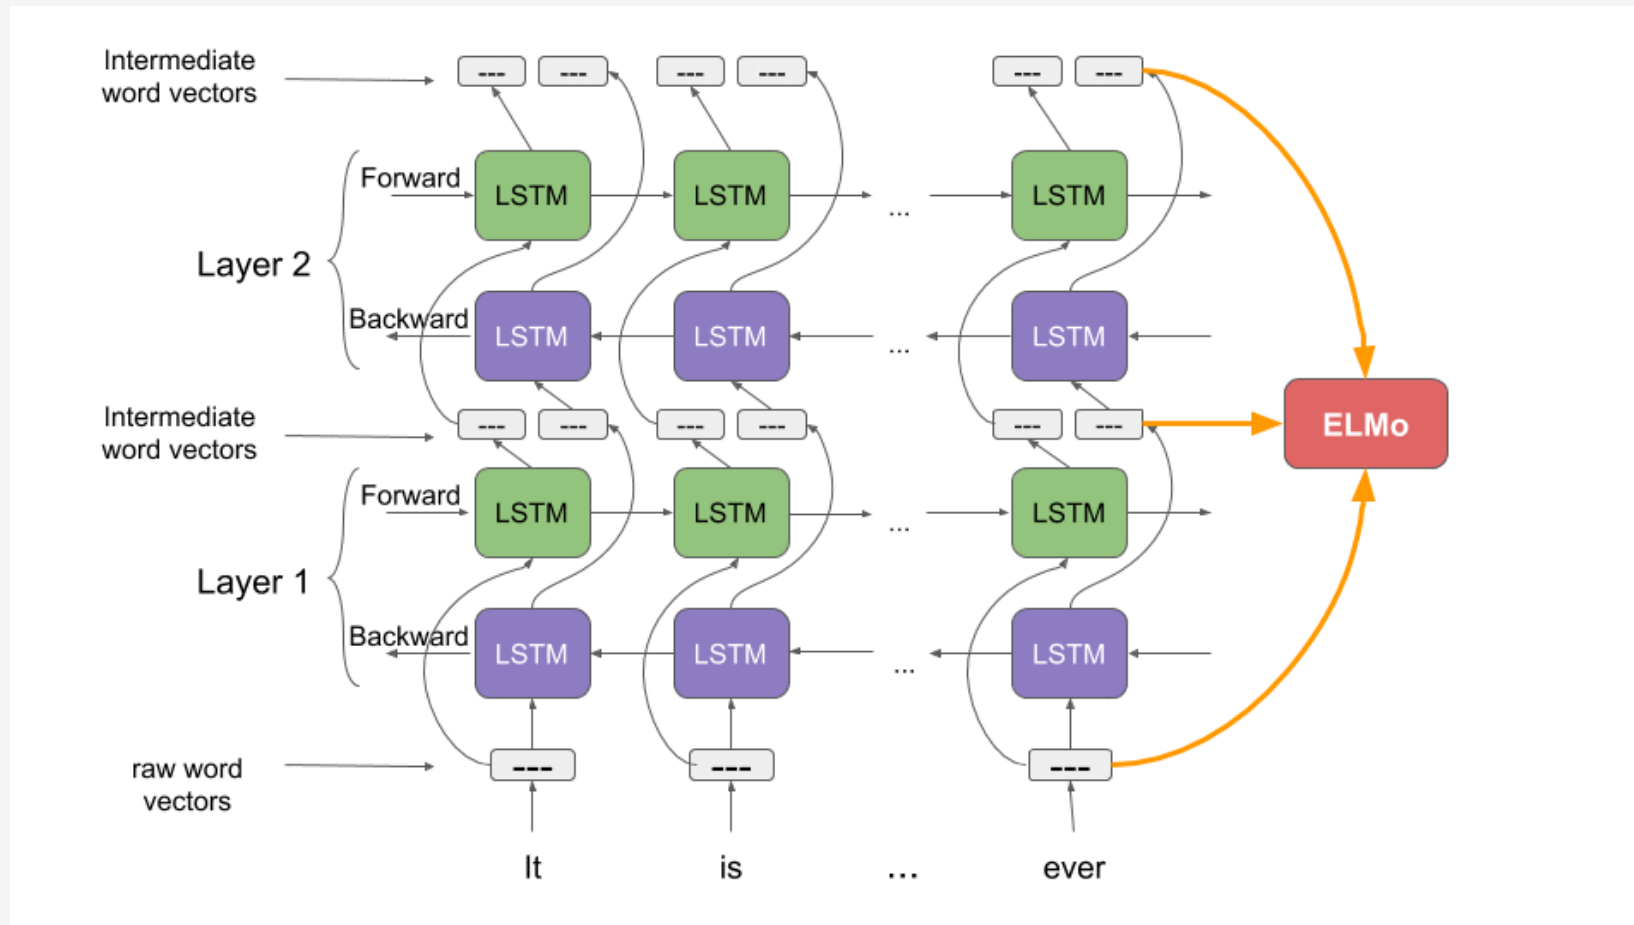
Contextual word embeddings

- The problem of polysemy
 - "The **cat** sat on a mat"
 - "You can **cat** this text file"
 - "I just got the results from my **cat** scanner"
- Are these the same word?
- Should they have the same embedding?
- Task specific embeddings may solve the problem, but can we do better?

ELMO – Deep contextualized word representations

- Key idea: generate a dynamic embedding, based on the context a word appears
 - “The cat sat on a mat” -> W2V -> the vector of “cat” depends only on “cat”
 - “The cat sat on a mat” -> ELMO -> the vector of “cat” depends on all words
- How? Bi-directional (LSTM) language model
- Concatenation of different layers
- Task-specific weights

Bi-directional language model



Bi-directional language model

- Forward language model:

$$p(t_1, t_2, \dots, t_N) = \prod_{k=1}^N p(t_k \mid t_1, t_2, \dots, t_{k-1}).$$

- Backward language model:

$$p(t_1, t_2, \dots, t_N) = \prod_{k=1}^N p(t_k \mid t_{k+1}, t_{k+2}, \dots, t_N).$$

- Bi-directional LM:

$$\sum_{k=1}^N (\log p(t_k \mid t_1, \dots, t_{k-1}; \Theta_x, \vec{\Theta}_{LSTM}, \Theta_s) \\ + \log p(t_k \mid t_{k+1}, \dots, t_N; \Theta_x, \overleftarrow{\Theta}_{LSTM}, \Theta_s)).$$

- Predict a word given its left and right context; keep input and output weights shared;

Deep representations

- For each token k , an L -layer bi-directional LM obtains $2L + 1$ representations

$$R_k = \{\mathbf{x}_k^{LM}, \vec{\mathbf{h}}_{k,j}^{LM}, \overleftarrow{\mathbf{h}}_{k,j}^{LM} \mid j = 1, \dots, L\} = \{\mathbf{h}_{k,j}^{LM} \mid j = 0, \dots, L\}$$

- Initial (static) representation \mathbf{x}
- For each layer – hidden representation of forward and backward

- ELMO learns a task-specific linear combination:

$$\mathbf{ELMO}_k^{task} = E(R_k; \Theta^{task}) = \gamma^{task} \sum_{j=0}^L s_j^{task} \mathbf{h}_{k,j}^{LM}. \quad (1)$$

- The representations R_k depend only on the context
- The embedding \mathbf{ELMO}_k depends on the context and the task

The impact and importance of ELMO

- ELMO (Peters et al. 2017) significantly improves NLP model performance
- Not the first contextual representation
- The first deep contextual representation (prior work only takes last layer)
- The first task-specific representation
- Short lived success due to the appearance of transformers and BERT
- Many of the concepts in ELMO are adopted in BERT

Using embeddings in downstream tasks and student projects

- What kind of vectors to use?
- Use pre-trained or re-train?
- Use as feature vectors or use to learn features?
 - E.g., can you “learn” which vectors are positive?
- Can you combine with other features?
- What would be the classifier?
- Size and scale of vectors?

End to end neural models

Accumulation of errors in a pipeline

- Assume a classifier working at 95%, but...
 - 95% accuracy tokenizer
 - 95% accuracy POS tagger
 - 95% accuracy dependency parser
 - 95% relevant feature mapping
 - 95% accurate classifier
 - What would be the total performance?
 - 77.3%

Why not go end to end?

- Do we need full pipelines?
- Embeddings make it possible to “feed” text directly into models
- Is it possible to go fully end-to-end and eliminate
 - Accumulation of errors
 - Human labor and supervision
 - (In)compatibility issues between elements?

High level idea

- Embeddings represent individual words
- Neural architecture combines embeddings (“words”) together
- Deeper hidden layers represent interactions between words (phrases, clauses, sentences)
- Backpropagation updates the way we combine words, but can also update embeddings

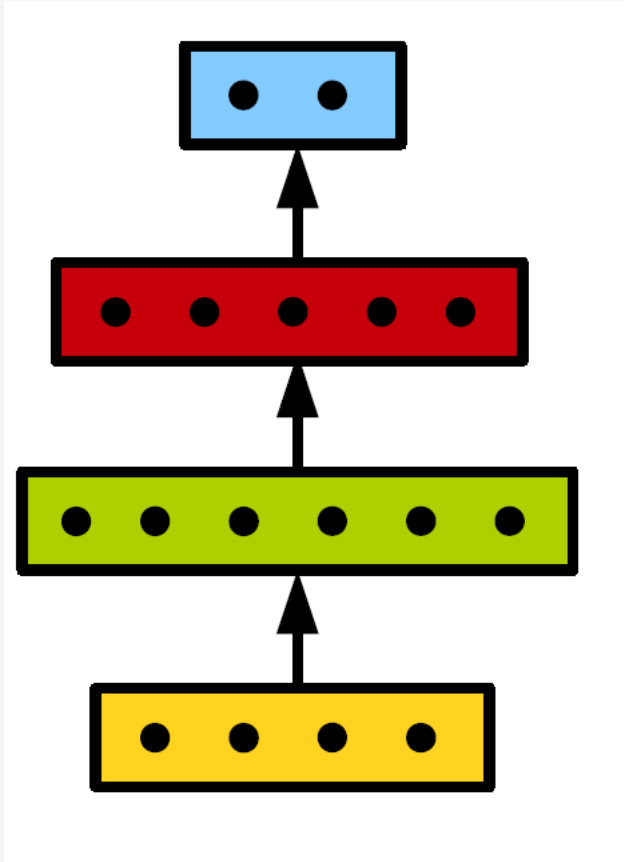
Neural architectures

- Feed forward neural networks
- Recurrent neural networks
 - LSTM, BiLSTM, GRU
- Convolutional neural networks

Feed-forward networks

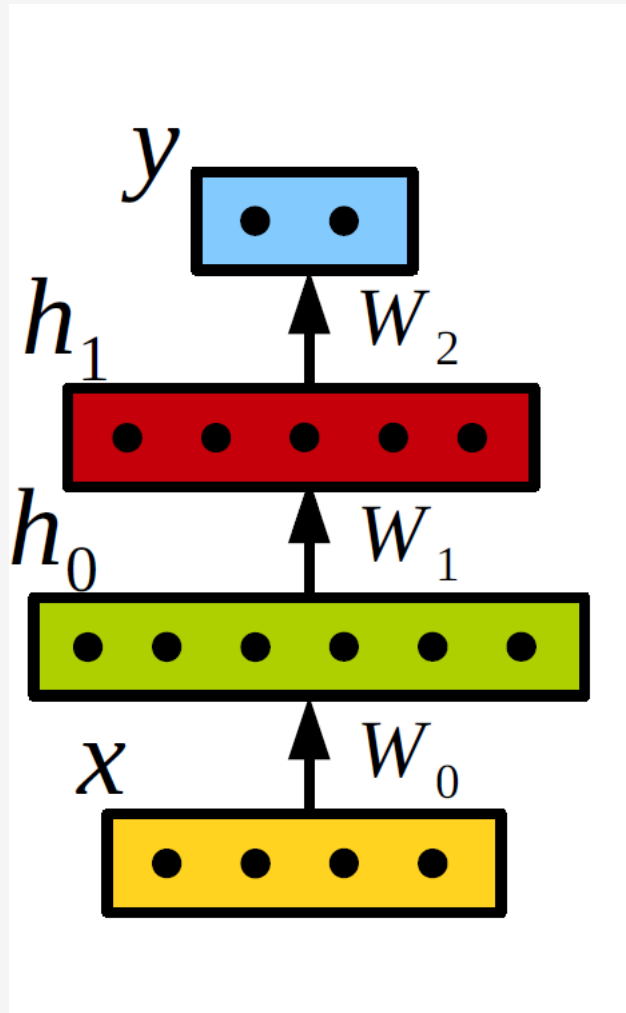
Multi-layer perceptron

Multi-layer perceptron



- An input layer (embedding)
- One or more hidden layers, each computed from the previous layer
- An output layer
 - Based on the last hidden layer
 - Class probabilities (spam / ham)

Multi-layer perceptron



- $y = \text{softmax}(h_1 \cdot W_2 + b_2)$
- $h_1 = f(h_0 \cdot W_1 + b_1)$
- $h_0 = f(x \cdot W_0 + b_0)$
- Non-linear functions f :
 - Sigmoid: $\sigma(x) = \frac{1}{\exp(-x)}$
 - Hyperbolic: $\tanh(x) = \frac{1 - \exp(-2x)}{1 + \exp(-2x)}$
 - ReLU: $\text{rect}(x) = \max(0, x)$

What is the input X?

- "The cat sat on a mat"
- We can convert all words into vectors
 - $w2v(\text{"the"})$, $w2v(\text{"cat"})$, $w2v(\text{"sat"})$, $w2v(\text{"on"})$, $w2v(\text{"a"})$, $w2v(\text{"mat"})$
- We obtain 6 300 dimensional vectors
- How do we add the input in the MLP?
 - Single fixed vector
 - A matrix of vectors
 - What are some advantages and disadvantages of each approach?

Training MLP

- Using Stochastic Gradient Descent
- Start with pretrained embeddings and random weights and biases
- Each epoch
 - Shuffle training data and select a training batch (set of k examples)
 - Compute the "forward" and the loss function
 - Compute the gradient of the loss function (backward)
 - Update the parameters (learning rate η): $W = W - \eta \frac{1}{K} \sum_{i=0}^{K-1} \nabla J_i(W)$
- Repeat until convergence

Training MLP

- Using Stochastic Gradient Descent
- Start with pretrained embeddings and random weights and biases
- Each epoch
 - Shuffle training data and select a training batch (set of k examples)
 - Compute the "forward" and the loss function
 - Compute the gradient of the loss function (backward)
 - Update the parameters (learning rate η): $W = W - \eta \frac{1}{K} \sum_{i=0}^{K-1} \nabla J_i(W)$
- Repeat until convergence

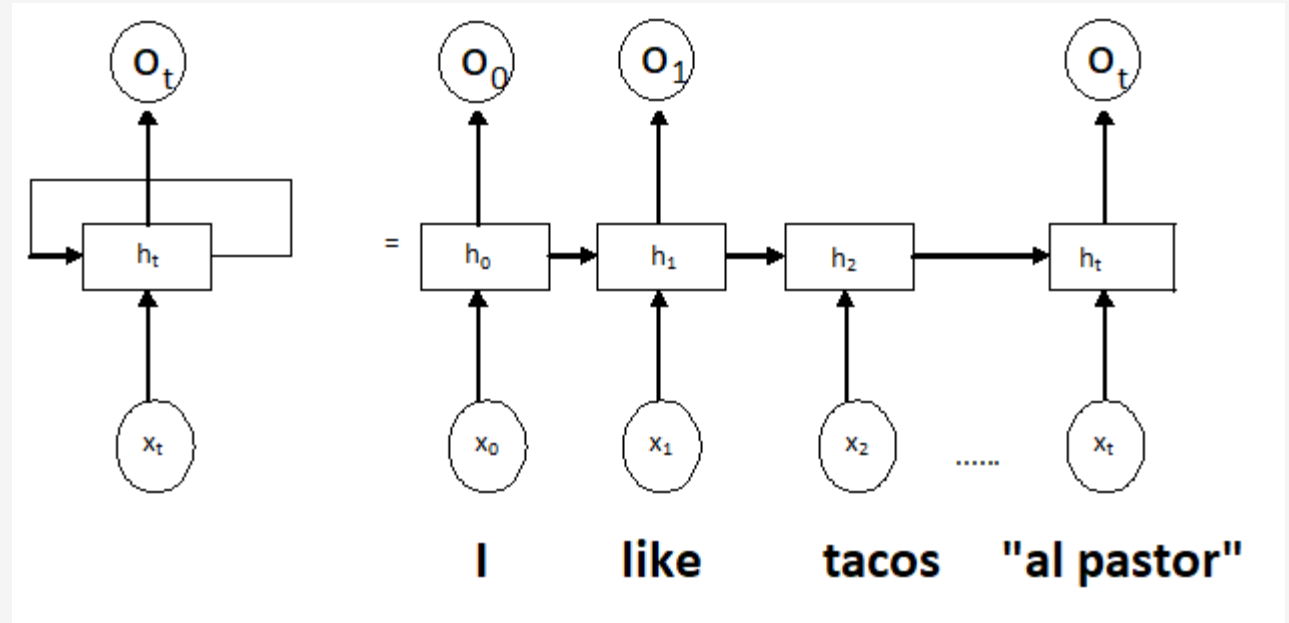
Training MLP

- SGD requires calculating the gradient of the loss function J , ∇J
 - For each example x_i
 - For each parameter $w \in W$
- Do we also update embeddings?
- How do we do that if we want to?

RNN, LSTM, and BiLSTM Neural Language Models

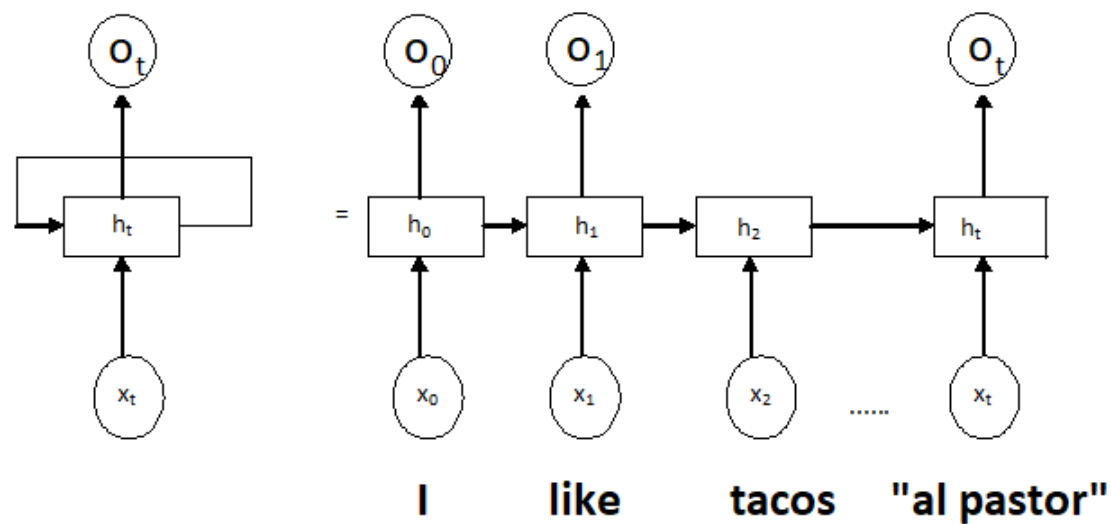
Recurrent neural networks

- How to represent text of varying length?
- Copy the network for each word
- At each timestep t , input is x_t and $h_{(t-1)}$
- I + like + tacos + "al pastor"
- Left to right, combining words one at a time
- Sequence classification
- Sequence to sequence



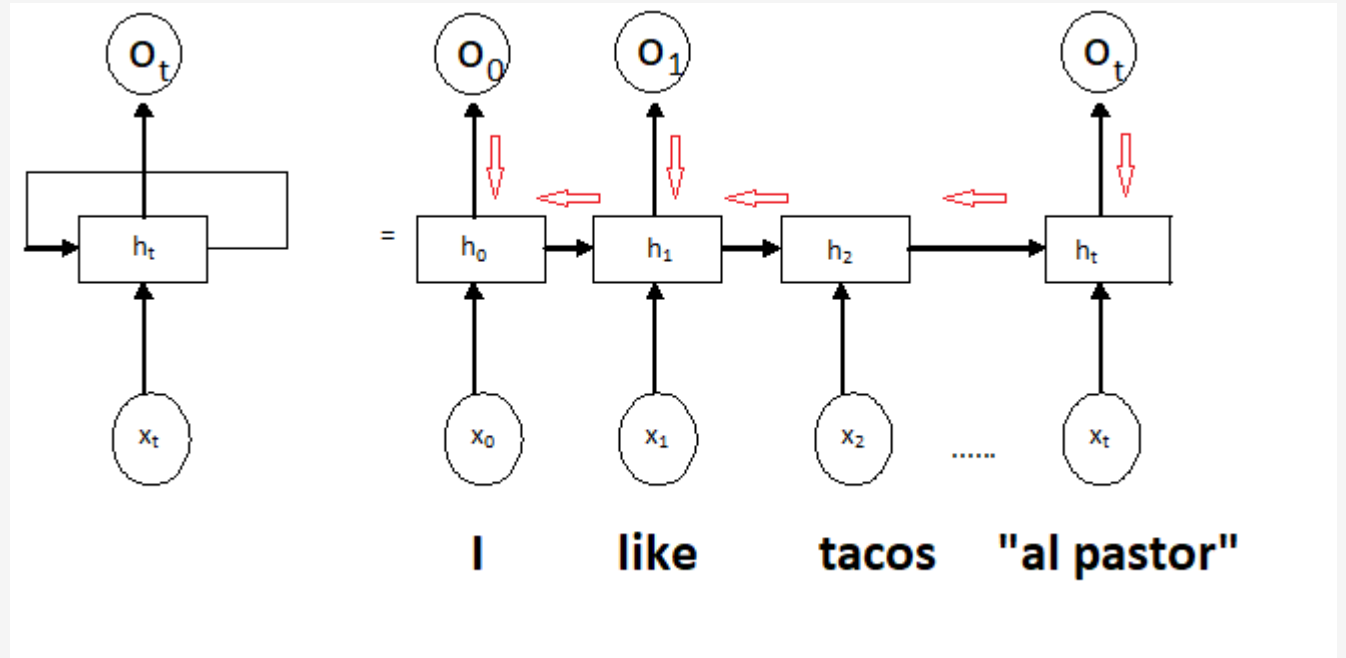
RNNs (formally)

- At each timestep
 - Input x_t and previous hidden state $h_{(t-1)}$
 - Three sets of weights:
 - input (W_x), hidden (W_h) , and output (W_o)
 - $h_t = f_h(W_x x_t + W_h h_{(t-1)})$
 - $O_t = f_o(W_o h_t)$
 - $J_t = f(O_t, y_t)$
- Pop quiz: are W_x , W_h , and W_o at time t_0 different than at time t_1 ?



Backpropagation through time

- Gradients flow backwards
- Propagate the gradient through each step, starting from the last
- Update the three weight matrices
- What would be a potential problem?
 - Exploding/vanishing gradients
 - Saturating non-linearity



Language modeling

- Recall: a language model assigns probability to a sequence of words:
 - Traditional Markov model: $p(w_1, \dots, w_m) = \prod_{t=1}^m p(w_t | w_1, \dots, w_{(t-1)}) \approx \prod_{t=1}^m p(w_t | w_{(t-n)}, \dots, w_{(t-1)})$
- Estimate the probability of a word given a history of a predefined length
- Estimate the probability of a sequence by multiplying the probabilities of all words
- Pop-quiz: is word2vec a traditional language model?

Neural language modeling

- If we use an RNN based model, we can drop Markov assumption
- Hidden state at time t keeps information about everything seen before
- At each step, we can add a classifier on top of the hidden state and predict the next word
- Add special symbols to mark the start and end of the sentence ($\langle s \rangle$ $\langle /s \rangle$)

Neural language modeling with RNNs

- More formally:
 - $h_t = \tanh(W_x x_t + W_h h_{(t-1)})$
 - $\hat{y}_t = \text{softmax}(W_y h_t)$
 - $J_t = -\log \hat{y}_{t, \text{correct}} \quad (-\log \text{prob the word at } t+1)$
-
- $J_{\text{sent}} = \frac{1}{T} \sum_{t=1}^T -\log \hat{y}_{t, \text{correct}}$

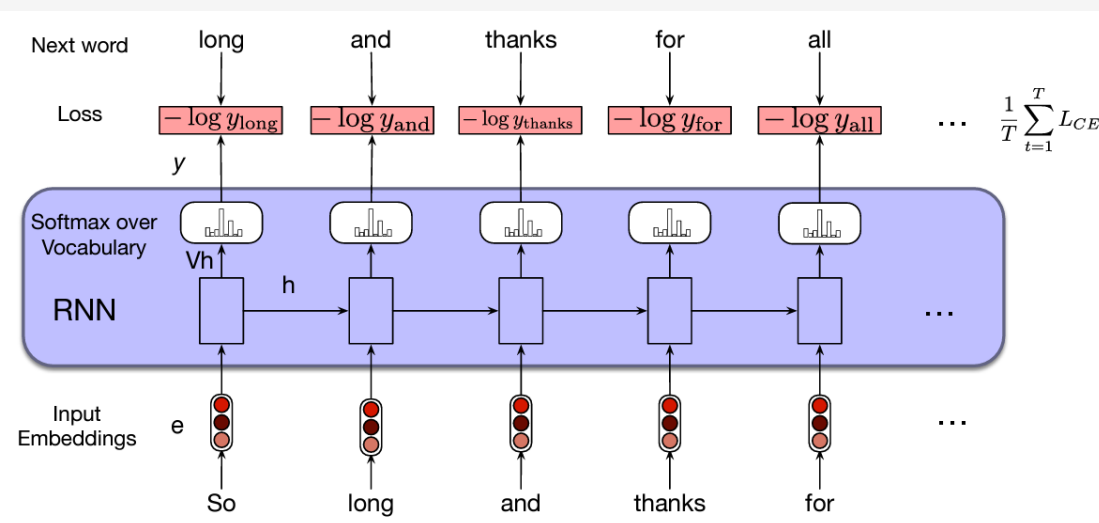


Figure 9.6 from SLP, chapter 9

Tying the weights

- Pop quiz: What are the dimensions of W_y ?
- What are the dimensions of the embedding matrix E ?
- We can use a quick “trick” and set the dimension of h to be same as the size of the vectors
- Then we can use E^T instead of W_y
 - Easier to train and performing better

The output layer in RNNs

- The output layer “knows” the history until time step t
- If the output layer predicts words, this the RNN becomes a language model (!)
- However, the output layer does not have to predict words
 - RNN can be used for classification, sequence tagging, generation

RNN for sequence tagging

- RNN can be trained to predict POS tags
- At each step, use the hidden state to predict POS
- Softmax over tags instead of words

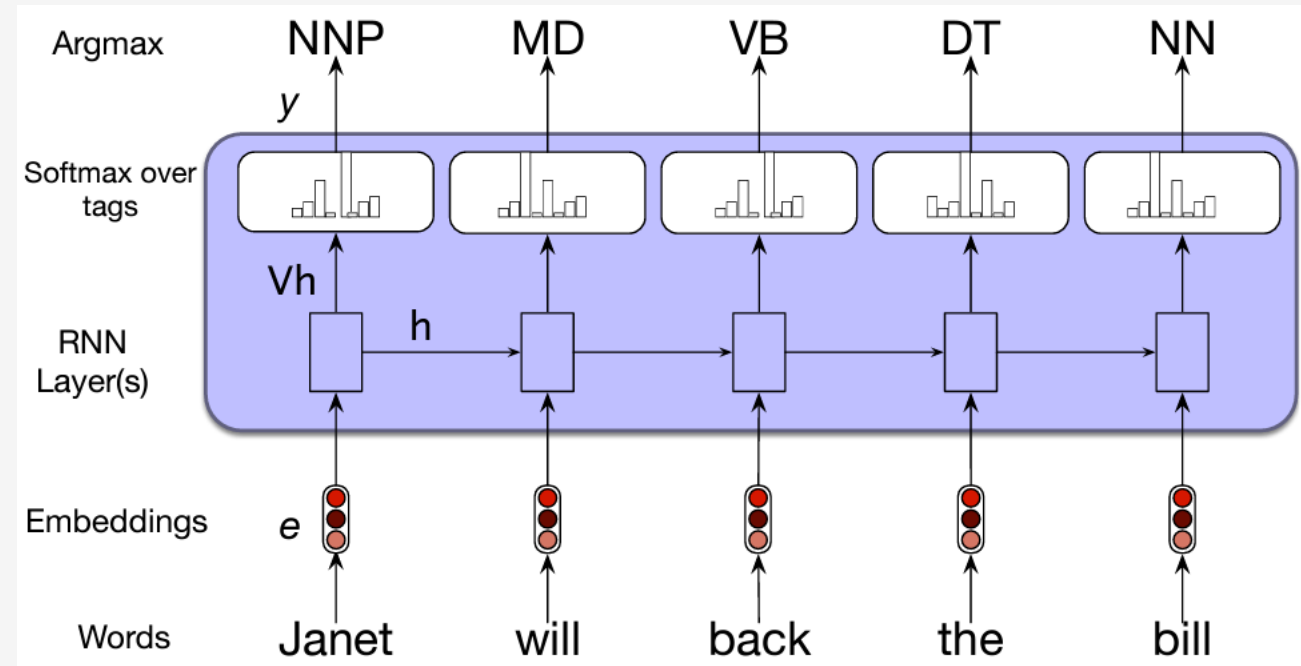


Figure 9.7 from SLP, chapter 9

RNN for text classification

- RNN can be used directly for classification
- Only use the last hidden layer
- Either predict directly, or add a FFN
- Predict intermediate sentiment?

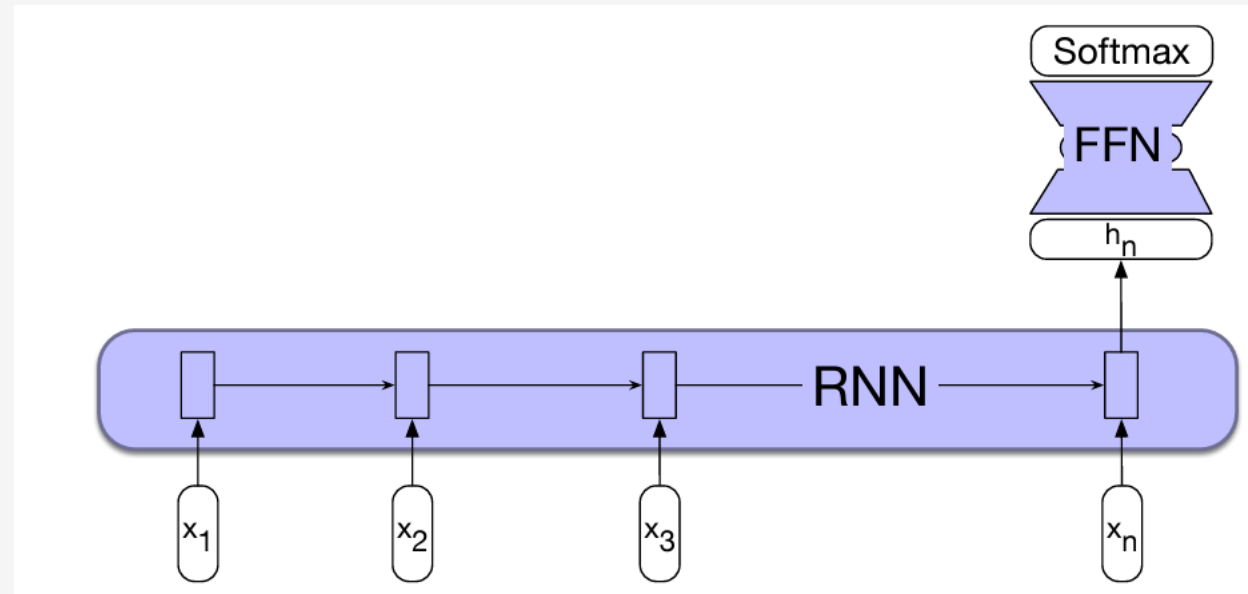


Figure 9.8 from SLP, chapter 9

Using RNN to generate text

- Generative AI
- How do you generate text using a Markov model?
- Generating using RNN
 - Sampling
 - Updating expectations
 - Teacher forcing

Stacked and bidirectional RNNs

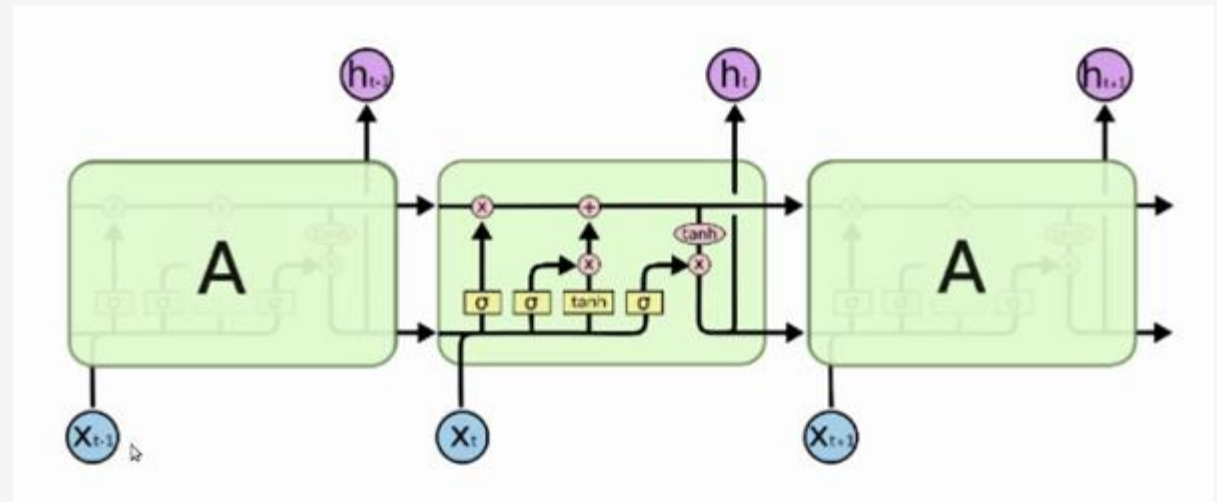
- The basic RNN is a powerful tool
- We can go deeper
- Stacking RNNs
- Bidirectional RNNs
 - Pop quiz: in a BiRNN for classification, which hidden states do we use from forward and backwards?

LSTM

- Popular variation of RNN that address native limitations
- Same recurrent concept (copy the network)
- Three different "gates":
 - Forget gate
 - Input gate
 - Output gate
- Gates manipulate the flow of information through time
 - Addressing conflicting objectives

LSTM. Some Notations

- x_t – input at time t
- h_t – hidden state at t after output gate
- “short term memory”
- C_t – state at t before output gate
- “long term memory”
- O_t – output at t
- W and U – weights with respect to x and h



Understanding the gates

- All gates have the same format:
 - A feedforward layer followed by a sigmoid
- The gates are filtering out information at a certain part of the network
- Each gate has two important aspects:
 - How to calculate the filter
 - What is the input that the filter is applied to

LSTM Gates: The forget gate

- Goal: determine how much to keep from the previous information (the long-term memory)
- Formula: $f_t = \sigma(U_f h_{(t-1)} + W_f x_t)$
 - Input at the current state (x_t), weighted by W_f
 - Hidden state at state $t-1$, weighted by U_f
 - Sigmoid activation (0 – forget everything, 1 – keep everything)
- “How much information to keep from the Long-term memory, given the current input and short-term memory”
- f_t is applied to $C_{(t-1)}$

LSTM Gates: The new information

- Goal: determine how surprising the new information is
- Formula: $g_t = \tanh(U_g h_{(t-1)} + W_g x_t)$
- This is our “normal” recurrence in RNN
- Tanh activation: what values does it take?
 - Tanh values are in $[-1, 1]$
- If the value of g_t is negative, the new information is subtracted from the memory, rather than added
- Note: this is not a gate, this is just your typical RNN hidden state

LSTM Gates: The input gate

- Goal: determine how much to use from the new information
- Formula: $i_t = \sigma(U_i h_{(t-1)} + W_i x_t)$
 - Similar to the forget gate, but with its own set of weights
- i_t is applied to the value of new information N_t
 - The input at state t is filtered twice: through N_t and i_t
- If we add the information that passes through the input and forget gates, we get the current long-term memory
 - $C_t = f_t * C_{(t-1)} + i_t * g_t$

LSTM Gates: The output gate

- Goal: determine what part of the long-term memory is necessary right now
- Consider machine translation: what is needed to translate current word vs what is needed to translate everything
- Formula: $o_t = \sigma(U_o h_{(t-1)} + W_o x_t)$
 - Very similar to forget and input gates, with its own set of weights
- o_t is applied to the value of C_t after passing it through another tanh
 - $h_t = o_t * \tanh(C_t)$
- h_t is the “traditional” hidden state – it can be used to make predictions at time step t

LSTM Results and improvements

- LSTM quickly became dominant paradigm in NLP
- Used in a variety of tasks:
 - Sentiment analysis (Wang et al. 2016)
 - Sequence labeling (Miwa et al. 2016)
 - Question Answering (Wang et al. 2015)
 - AMR parsing (Foland et al. 2017)
 - NLI (Chen et al. 2017)

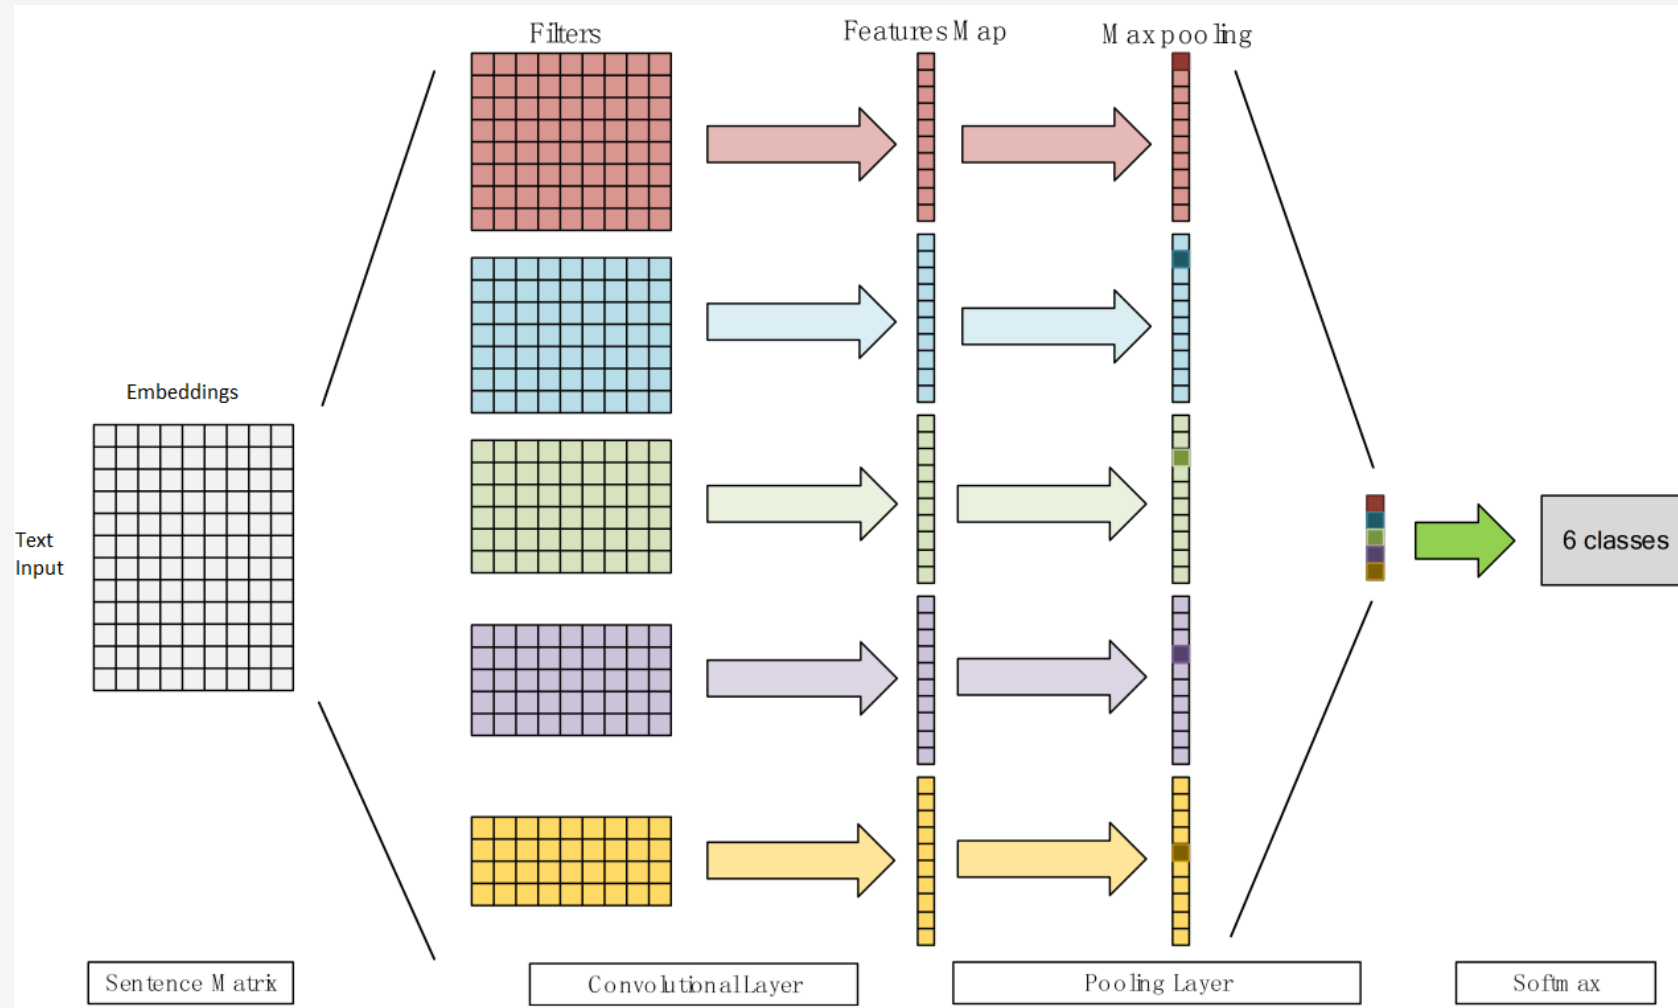
BiLSTM, stacked LSTMs, GRU

- Similar to RNN
- Deep LSTM
- BiLSTM – the de-facto norm of using LSTMs
- GRU – a different variation of gated RNN, slightly simpler than LSTM

Convolutional neural networks (CNN)

- Another popular architecture for NLP
 - Deals with text of various size
- Inspired by computer vision success
- Applying filters of different size to the input (2,3,4) + pooling
 - Analogous to n-grams

Convolutional neural networks (CNN)



Multimodality and multilinguality

- Embeddings and multilinguality
 - Words of any language can be mapped to vectors
 - Words of different languages can be mapped to the same space
 - Mapping and similarity across languages
- Embeddings and multimodality
 - Images, sound, and other modalities can also be mapped to vectors (e.g., using CNNs)
 - Shared multimodal spaces (vision + language)

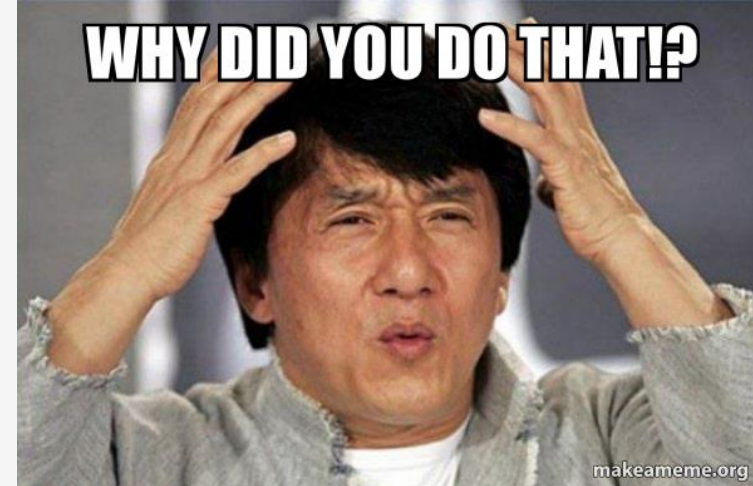
The first “should I worry about my job”

- Rapid change in technologies can be stressful
- Three “major” milestones in NLP in the past 10 years
 - Word2Vec and end-to-end models (BiLSTM, CNN) – 2013
 - BERT (and the transformer family) – 2018
 - Large generative language models (GPT3, ChatGPT, Bard) – 2022 - 2023
- There are still many unsolved problems

Some concerns

Explainability and Interpretability

- Interpreting feature-based models
 - Feature values ("v1agra") + weights = prediction ("spam")
- Interpreting end-to-end neural networks
 - Feature values (300d dense vector)
 - weights (input, forget, output gates)
 - different types of nonlinearity



Explainability and Interpretability

- Why did you do that?
- How did you do that? What makes you say that?
- How can I verify your results?

Bias, Guarantees, and Robustness

- Does the algorithm discriminate?
- Does the algorithm contain bias with respect to race, gender, religion, sexual orientation?
- Does the algorithm guarantee consistent and robust performance?
- Is the algorithm secure from adversarial attacks?

New fields of study in NLP in the are of deep learning

- Explainability of neural networks
- Algorithmic fairness
- Evaluation, unit testing, and adversarial attacks for NLP
- Data centric AI

Conclusions

Embeddings in NLP

- Embeddings changed the way we do NLP
- Valuable stand-alone resource
 - Can be used to query lexical information
 - Can be used as automatically extracted features
 - Can be used for a simple text representation
 - Static and dynamic embeddings, polisemy
- Enable end-to-end neural models

End-to-end neural models

- Minimizing human interaction and supervision
 - Remove the need of feature engineering
 - Only require labeled data for classification tasks
- Improving efficiency and removing accumulation of errors in pipeline
- Enabling full training without depending on external resources

End-to-end neural models

- Introducing new challenges and problems
 - Increased computational complexity
 - Need for more data
 - Error accumulation becomes internal
 - Difficult to interpret and debug
 - Potentially containing biases
 - Ultimately re-invent many of the pipeline parts during training