

# Encoder-Decoder Models

## Attention. “The” transformer model.

Venelin Kovatchev

Lecturer in Computer Science

[v.o.kovatchev@bham.ac.uk](mailto:v.o.kovatchev@bham.ac.uk)

# Outline

---

- Quick recap
- Encoder-decoder networks
- Attention
- Original transformer

# Quick recap: End-to-end neural networks

# What are end-to-end models


---

- Task specific models
- Map directly from input to output
- No feature engineering
- Trained via backpropagation
  - Data and compute expensive

# What are some advantages of end-to-end

---

- Better performance
- Simpler pipeline
- Changing the problem formulation
  - The task is defined by the data and the metrics
- Making NLP more accessible
  - Plug and play



# Challenges with going end to end


---

- My take on key challenges
  - Computational and data cost
  - Dependency on data and task formulation
  - Explainability and Interpretability
  - Bias, guarantees, and robustness

# Explainability and Interpretability

---

- Interpreting feature-based models
  - Feature values ("v1agra") + weights = prediction ("spam")
- Interpreting end-to-end neural networks
  - Feature values (300d dense vector)
  - weights (input, forget, output gates)
  - different types of nonlinearity



# Explainability and Interpretability

---


- Provide a (valid) justification for the model behavior
- Provide a faithful explanation of the model behavior
- Provide an explanation that is useful for a human
  - To assess the model
  - To learn how to perform the task
  - In a Human-Computer collaboration



# Bias, Guarantees, and Robustness

---

- An end-to-end neural network finds the (mathematically) optimal solution to a formally defined problem
- Sometimes the optimal solution can lead to undesired behavior
  - bias with respect to race, gender, religion, sexual orientation
  - “shortcuts” to solving tasks
- How do we guarantee the model is consistent and bias-free?
  - Evaluation and algorithmic fairness
- How do we know if the algorithm is safe from adversarial attacks?



# What networks do we know so far

---

- Feed forward networks (FFN)
- Recurrent neural networks (RNN) (+ LSTM, GRU)
- Convolutional neural networks (CNN)
- Pop quiz: are these networks for supervised or unsupervised NLP?

# Encoder-Decoder Models


# Input and output in NLP tasks

---

- What is the input and output of the following tasks
  - Sentiment analysis
  - Automated fact checking
  - Clustering documents based on topic
  - Machine translation
- How many possible outputs does each of those tasks have?


# Sequence labeling vs sequence-to-sequence

- Consider the following two tasks



- What is similar between them?

- What is different?



- How would you approach each of them?

# Key differences


---

- Same length vs different length
- One-to-one alignment vs no one-to-one alignment
- Local dependencies vs long-distance dependencies
  - Within the output
  - Between the input and the output

# Different task formulations

- Which of the following images corresponds to:

- FFN
- RNN for text classification
- Machine translation
- Image captioning
- Sequence labeling



The Unreasonable Effectiveness of Recurrent Neural Networks by Andrej Karpathy

# Encoder decoder

---

- We use a model family called encoder-decoder
- Simple idea
  - Encoder “represents” the source (e.g., English)
  - Decoder “generates” the target (e.g., German)
- Can you suggest tasks that can use encoder-decoder?

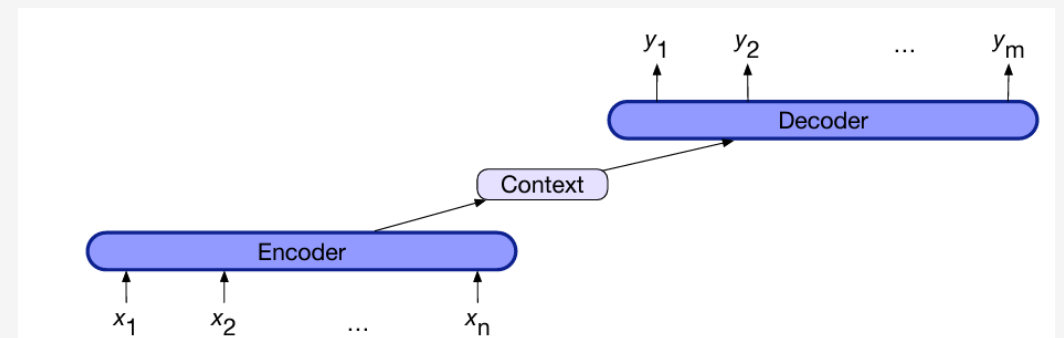


Fig 9.16



# Usage of encoder decoder

---

- General usage of encoder decoder
  - Mapping between data of different format, size, and structure
  - Encoder-decoder vs sequence-to-sequence
- Examples for tasks that can use encoder-decoder:
  - Machine translation
  - Text summarization
  - Question answering
  - Image captioning



# How do we implement encoder-decoder?

---

- Can you propose a way to implement enc-dec model with what we know so far?
- How do we encode the input?
- How do we decode the output?

# Single RNN as encoder decoder

---

- Let's consider a single RNN for the task
- Add a separator between the two texts:
  - [sentence] [in] [English] [SEP] [sentence] [in] [German]
- The hidden state at SEP will contain all the information about the first sentence

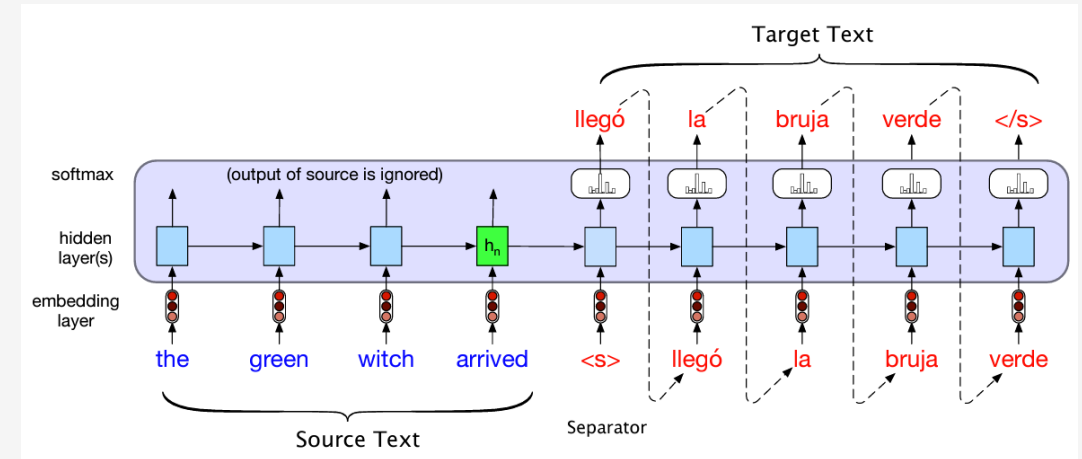
# Conditional generation

---

- How does a traditional language model generate text?
- How does an encoder-decoder RNN generate text?
- Does that concept look familiar?

# A single RNN as encoder-decoder

- Consider using the following model
  - We use "English" as a "prompt"
  - Hidden state at  $\langle s \rangle$  "encodes" the text
  - We generate Spanish step by step from  $x$  and  $h$




- What would be some problems with this model?
  - What if the task was text captioning?

# Using separate RNNs for encoder and decoder

- Train two models
- Pass the context at every step

$$\mathbf{h}_t^d = g(\hat{y}_{t-1}, \mathbf{h}_{t-1}^d, \mathbf{c})$$

- Can you point a potential problem?
- What could improve this architecture?
- What is the purpose of the encoder?
- Should it be able to generate?



# Formal representation of RNN based decoder


---

- The context is the last  $h$  of the encoder
- The hidden stage at step 0 is just the context
- For every step after 0, we use both  $h$  and  $c$
- We use the hidden state to predict  $y$  at time  $t$
- Why is there a "y" at the calculation of the hidden state  $h_t^d$ ?

$$\begin{aligned}\mathbf{c} &= \mathbf{h}_n^e \\ \mathbf{h}_0^d &= \mathbf{c} \\ \mathbf{h}_t^d &= g(\hat{y}_{t-1}, \mathbf{h}_{t-1}^d, \mathbf{c}) \\ \mathbf{z}_t &= f(\mathbf{h}_t^d) \\ y_t &= \text{softmax}(\mathbf{z}_t)\end{aligned}$$

# Training encoder-decoder models

- Models are trained end-to-end
- Encoder is trained through hidden layers
- Decoder is trained through teacher forcing
  - Remember “teacher forcing”?





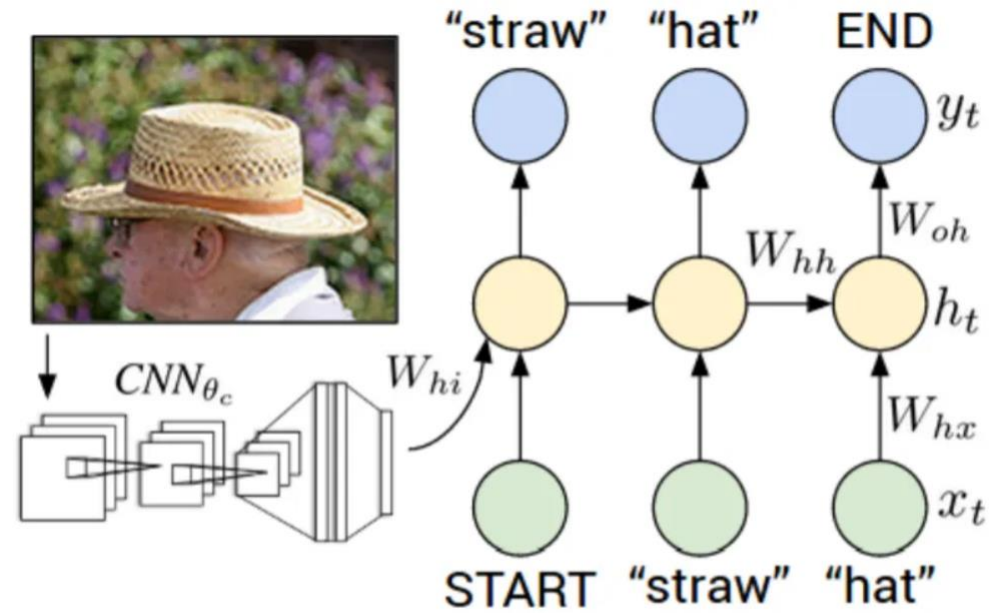
# Why are encoder-decoder models important?

---

- Instant improvement on machine translation
  - Google Translate switching to NMT
- Key concepts reused (and giving raise to) Attention and Transformers
- Bridging the gap between modalities

# Encoder decoder across modalities: image captioning

- The encoder and decoder “talk” via the context
- They don’t have to be the same type of model
- The modalities don’t have to match
  - Speech to text
  - Image to text

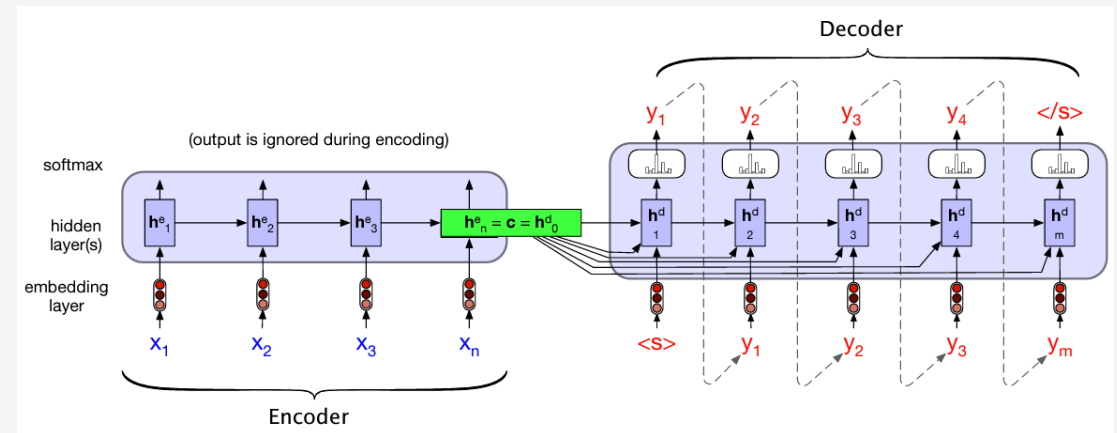
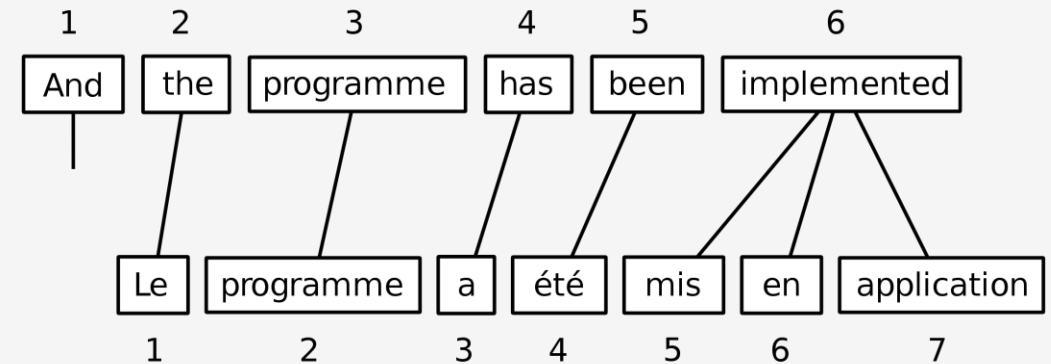


Deep Visual-Semantic Alignments for Generating Image Descriptions

Attention

# A bottleneck of RNNs

- Consider the problem of MT and an encoder-decoder solution
- The “context” is the information from the input that we need to generate the target
- To generate word  $y_t$ , we use the prior information for  $y_1 - y_{(t-1)}$  and the same  $c$
- Should  $c$  be the same for every word?



# Attention – intuition and restrictions


---

- Intuition: each token in the target should use a “personalized” context
- Access all the hidden states in the encoder
  - Still needs to have a fixed length, regardless of variable input length
- Any ideas how we can do that?

# Attention – basic implementation

---

- Weighted sum of all encoder hidden states
  - Calculated separately at each decoder step
  - Using the hidden state at (t-1)
- 
- Dot product attention
  - Calculate the similarity between  $h_{(t-1)}$  and each encoder state  $h^e$
  - Use the similarity scores to calculate the weighted sum



# Dot product attention (formally)

---

- Scoring function:

$$\text{score}(\mathbf{h}_{i-1}^d, \mathbf{h}_j^e) = \mathbf{h}_{i-1}^d \cdot \mathbf{h}_j^e$$

- Weight vector:

$$\begin{aligned} \alpha_{ij} &= \text{softmax}(\text{score}(\mathbf{h}_{i-1}^d, \mathbf{h}_j^e)) \\ &= \frac{\exp(\text{score}(\mathbf{h}_{i-1}^d, \mathbf{h}_j^e))}{\sum_k \exp(\text{score}(\mathbf{h}_{i-1}^d, \mathbf{h}_k^e))} \end{aligned}$$

- Personalized context:


$$\mathbf{c}_i = \sum_j \alpha_{ij} \mathbf{h}_j^e$$

- More complex scoring functions:

$$\text{score}(\mathbf{h}_{i-1}^d, \mathbf{h}_j^e) = \mathbf{h}_{i-1}^d \mathbf{W}_s \mathbf{h}_j^e$$

# Visualization of RNN with attention


- RNN with attention
- Attention is learned via a simple FFN





# Visualizing attention

- Linear weights are interpretable
- We can see which word is more important
- Can we use attention for explainability?



Attention is all you need  
The original Transformer

# Training vs Finetuning

---

- Simple end-to-end models are trained for a single task
- Word embeddings can be reused, compositionality is learned
- Transfer learning has limited capabilities
  - From similarity to inference
  - From emotion to sentiment

# Need for powerful transfer learning models

---

- Generic representation framework
  - Represent (contextual) word meaning
  - Represent interactions between words
  - Capture different types of meaning and interactions
- Easy to adapt to new tasks with minimal adjustment
- Looks familiar?
  - Many of the problems and RQs remain the same, just the context changes


# Look back at ELMO

---

- ELMO embeddings meet most of those expectations
  - In-context meaning
  - Interactions between words
  - Deep representation capturing different relations
  - Task specific weight learning
- Pop quiz: how did ELMO embeddings work?

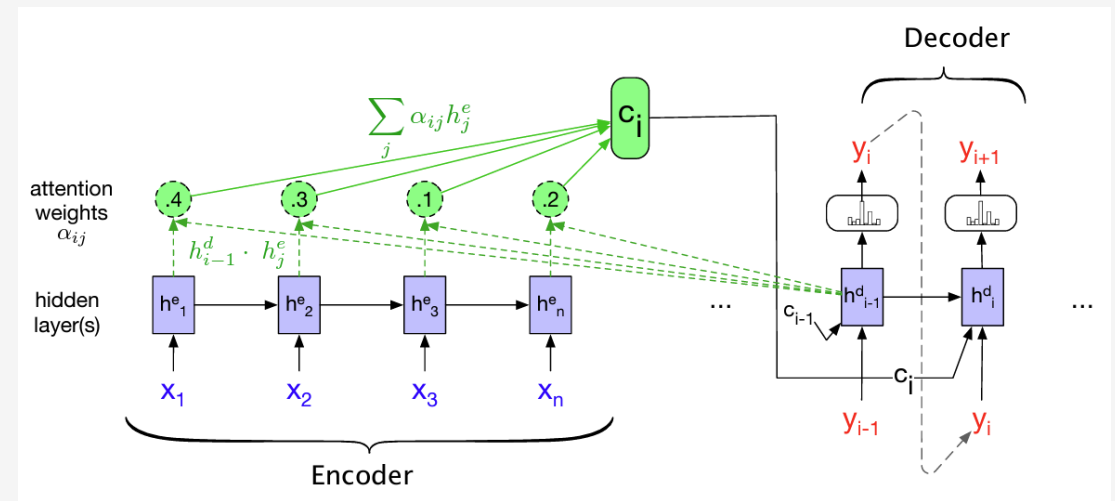
# Elmo architecture

- How can we improve over that?




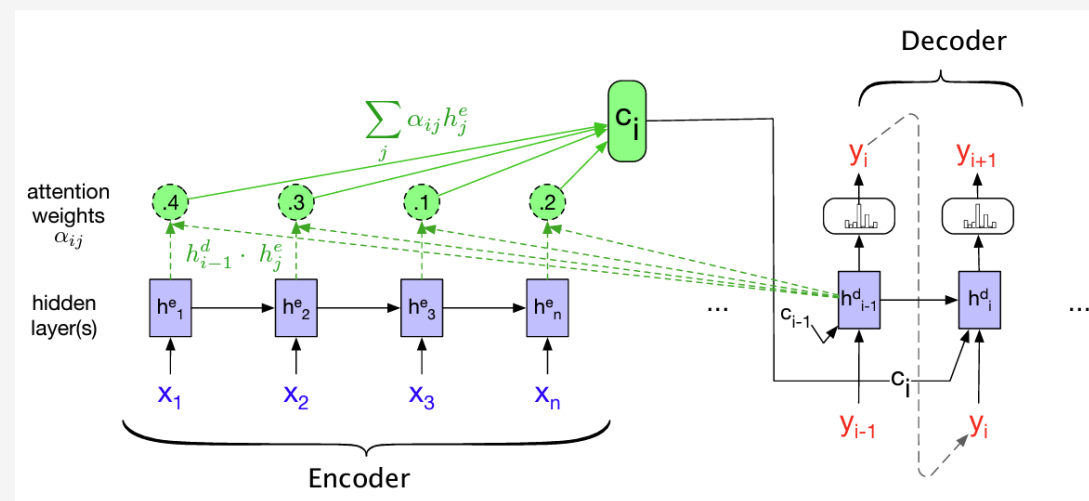
# Self attention

- Attention works better than RNN/LSTM for encoder-decoder models
- Can we use attention for a standalone network?



# Self attention (2)

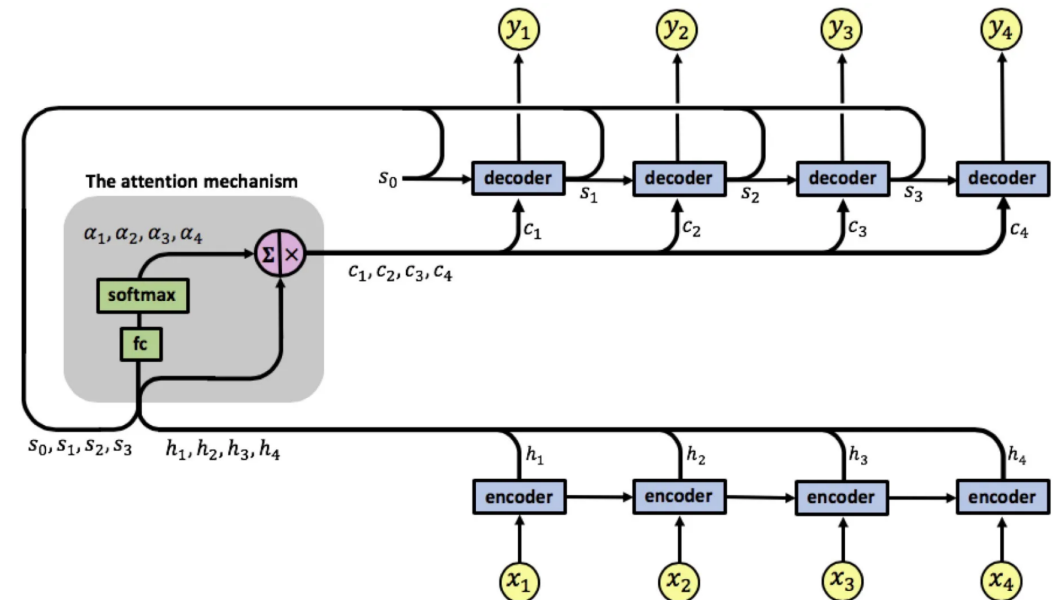
- Self attention is a key concept in building transformers
- It applies the same approach as attention in encoder-decoder, but on itself





# Causal attention vs bidirectional attention

- In encoder-decoder attention, the attention is the weighted sum of all hidden states of the encoder
- Which hidden states do we use in self attention?
- Why?



# Causal self attention

---

- Causal self-attention is used in models like GPT
- Two key properties
  - Only calculated using words in one direction (left for european languages)
  - Each representation at a layer L is calculated independently of the others
- How does this compare to RNNs and LSTMs?
  - Why are these two properties important?

# Pop quiz

---

- Can a transformer model process infinite input?
- Can an RNN be (natively) parallelized?

# Causal self attention (intuition)


- Similar to RNNs, we have a 1:1 input-output mapping
- Same basic approach as original attention
- Dot product + softmax + weighted sum

$$\text{score}(\mathbf{x}_i, \mathbf{x}_j) = \mathbf{x}_i \cdot \mathbf{x}_j$$

$$\begin{aligned}\alpha_{ij} &= \text{softmax}(\text{score}(\mathbf{x}_i, \mathbf{x}_j)) \quad \forall j \leq i \\ &= \frac{\exp(\text{score}(\mathbf{x}_i, \mathbf{x}_j))}{\sum_{k=1}^i \exp(\text{score}(\mathbf{x}_i, \mathbf{x}_k))} \quad \forall j \leq i\end{aligned}$$

$$\mathbf{a}_i = \sum_{j \leq i} \alpha_{ij} \mathbf{x}_j$$

- Which is the most similar token to  $\mathbf{x}_3$ ? What is the input to the first hidden layer?



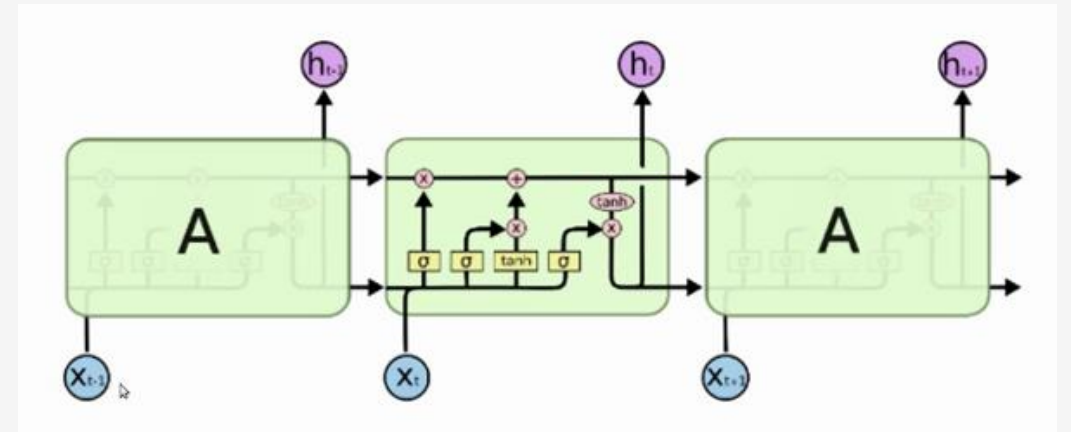
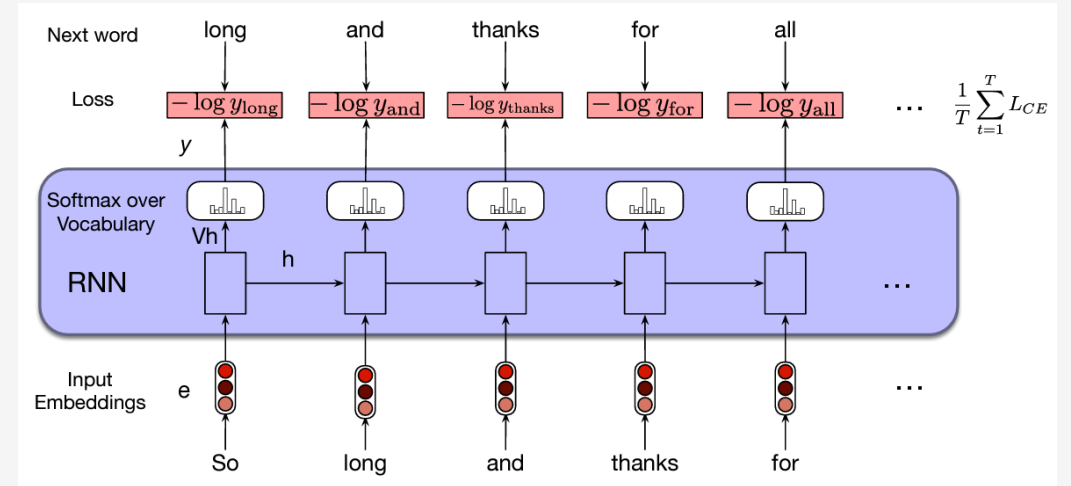
# Decomposing input vectors

---

- We can use simple attention
- Transformers introduce query, key, value
- What are they, why do we need them and how do we use them?
  - The “dictionary” analogy
  - A semantic explanation, grounded in NLP

# The LSTM and RNN

- What is the difference between LSTM and RNN?
- Why do we need that “evolution”?
- Break a single hidden state into two + gates
  - Filtering and specialization



# Compositionality of meaning

---

- Consider the following phrases
  - "A black dog"
  - "A house for my dog"
- What is the meaning of the dog in each phrase?
- Where is the dog in the second picture?



# Compositionality of meaning (2)

---

- What about "A house for my black dog"?
  - Does "dog" change the meaning of "house"?
  - Does "house" change the meaning of "dog"?
  - Does "dog" change the meaning of "black"?
  - Does "black" change the meaning of "dog"?
- Meaning compositionality can be asymmetrical!



# Different aspects of meaning compositionality

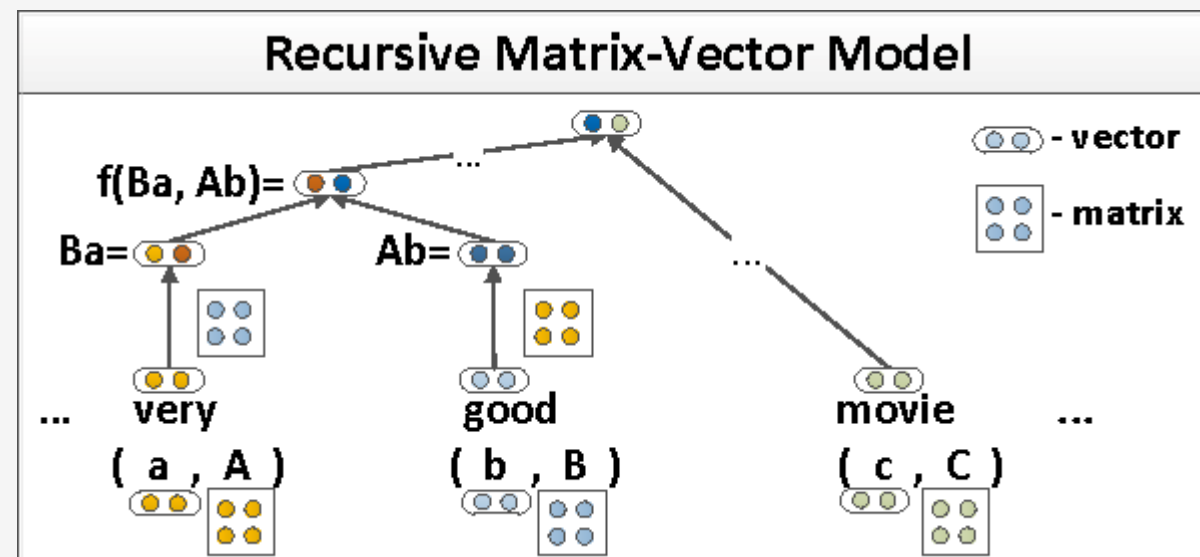
- Meaning compositionality is not a simple addition

- Words "behave" differently in different context

- Socher's Vector-Matrix representation

- Vector for the head, matrix for the complement


- Pop quiz: what would be the vector and what would be the matrix in "black dog"?



# How to model asymmetric compositionality in attention?

---

- Self attention (that we have seen) has 1:1 correspondence
- Dot product attention is commutative
  - $a \cdot b = b \cdot a$
  - $\text{score}(\text{"black"}, \text{"dog"}) = \text{score}(\text{"dog"}, \text{"black"})$



- Pop quiz: would "black" have the same importance on "dog" as "dog" would have on "black"?

# The query, key, value

---

- We project the input vector  $x$  to three vectors that serve different purpose: "query", "key", and "value"
- Two vector operations in the original attention:
  - "Score": for indexes  $i$  and  $j$ , calculate how important is  $x_j$  for  $x_i$ :  $\text{score}(x_i, x_j)$
  - "Scale": for index  $i$ , calculate the hidden state  $h_i$  as a weighted sum of  $x_1 \dots x_i$ :  $h_i = \sum_{j \leq i} \alpha_{ij} x_j$
- Each input vector  $x$  can have three different roles
  - Argument 1 in  $\text{score}()$  ["dog" in  $\text{score}(\text{"dog"}, \text{"black"})$ ] -> **query**
  - Argument 2 in  $\text{score}()$  ["black" in  $\text{score}(\text{"black"}, \text{"dog"})$ ] -> **key**
  - The **value** used in scale to calculate the hidden state

# Query, Key, Value (formally)

---

- We learn three different matrices ( $W^Q, W^K, W^V$ )
- Every input vector  $x_i$  is projected to three different representations
  - $q_i = x_i W^Q ; k_i = x_i W^K ; v_i = x_i W^V$
- The new formula for score:  $\text{score}(\mathbf{x}_i, \mathbf{x}_j) = \mathbf{q}_i \cdot \mathbf{k}_j$
- The new formula for calculating weights:  $\mathbf{a}_i = \sum_{j \leq i} \alpha_{ij} \mathbf{v}_j$
- Pop quiz: which token will have the most impact on  $x_3$ ?


# The transformer self attention

1.  $\mathbf{q}_i = \mathbf{x}_i \mathbf{W}^Q; \mathbf{k}_i = \mathbf{x}_i \mathbf{W}^K; \mathbf{v}_i = \mathbf{x}_i \mathbf{W}^V$

2. and 3.  $\text{score}(\mathbf{x}_i, \mathbf{x}_j) = \frac{\mathbf{q}_i \cdot \mathbf{k}_j}{\sqrt{d_k}}$

4.  $\alpha_{ij} = \text{softmax}(\text{score}(\mathbf{x}_i, \mathbf{x}_j)) \quad \forall j \leq i$

5. and 6.  $\mathbf{a}_i = \sum_{j \leq i} \alpha_{ij} \mathbf{v}_j$



# Parallelizing and masking the future

- Calculating hidden state  $h_t$  is independent of  $h_{(t-1)}$
- We can compute all hidden states in a single operation
- $Q = XW^Q$ ;  $K = XW^K$ ;  $V = XW^V$
- $A = \text{SelfAttention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$
- Can you see a problem for causal self attention?
- Pop quiz: What is the complexity of the self-attention w.r.t. length of the input?

N	q1•k1	−∞	−∞	−∞	−∞
	q2•k1	q2•k2	−∞	−∞	−∞
	q3•k1	q3•k2	q3•k3	−∞	−∞
	q4•k1	q4•k2	q4•k3	q4•k4	−∞
	q5•k1	q5•k2	q5•k3	q5•k4	q5•k5
N					

# Multiheaded self-attention


---

- Instead of using a single self attention, we can use multiple
  - Each “head” has its own weights  $W^Q, W^K, W^V$
  - The outputs of all heads are concatenated and projected to input dimensions
- Formally:

$$\begin{aligned}\mathbf{Q} &= \mathbf{XW}_i^Q ; \mathbf{K} = \mathbf{XW}_i^K ; \mathbf{V} = \mathbf{XW}_i^V \\ \mathbf{head}_i &= \text{SelfAttention}(\mathbf{Q}, \mathbf{K}, \mathbf{V}) \\ \mathbf{A} &= \text{MultiHeadAttention}(\mathbf{X}) = (\mathbf{head}_1 \oplus \mathbf{head}_2 \dots \oplus \mathbf{head}_h) \mathbf{W}^O\end{aligned}$$

# The transformer block

- Residual connection
  - Copy the input of a layer to its output
- Layer normalize
  - Rescale each  $x$  vector to 0-mean with  $STD=1$
- Feedforward
  - Apply the same fully connected FFN to each  $x$






# The transformer block (formally)

---

- Simplified representation
  - $O = \text{LayerNorm}(X + \text{MultiHeadAttention}(X))$
  - $H = \text{LayerNorm}(O + \text{FFN}(O))$
- You can change the order of operations in some implementations


# Does transformer consider word order?

- Consider an autoregressive transformer
- Does it handle long-distance dependencies?
- Does it handle word order?
- Does the position of  $x_1$  and  $x_2$  matter?



# Encoding the Input. Positional Embeddings.

- Semantic embeddings
  - One-hot encoding maps to a row in a matrix
- Positional embeddings
  - One embedding for each position
  - Learnable; Same dimension as semantic
- Add semantic and positional embeddings



- Alternative techniques: use functions (sine/cosine); calculate relative positional embeddings


# Classification layer: The “Head” of the model

---

- How did we train word2vec?
  - How did we reduce the computational cost?
- The concept of transfer learning
  - Train on one objective
  - Reuse the model on another task
  - We keep the stacked transformer blocks, change the “head”

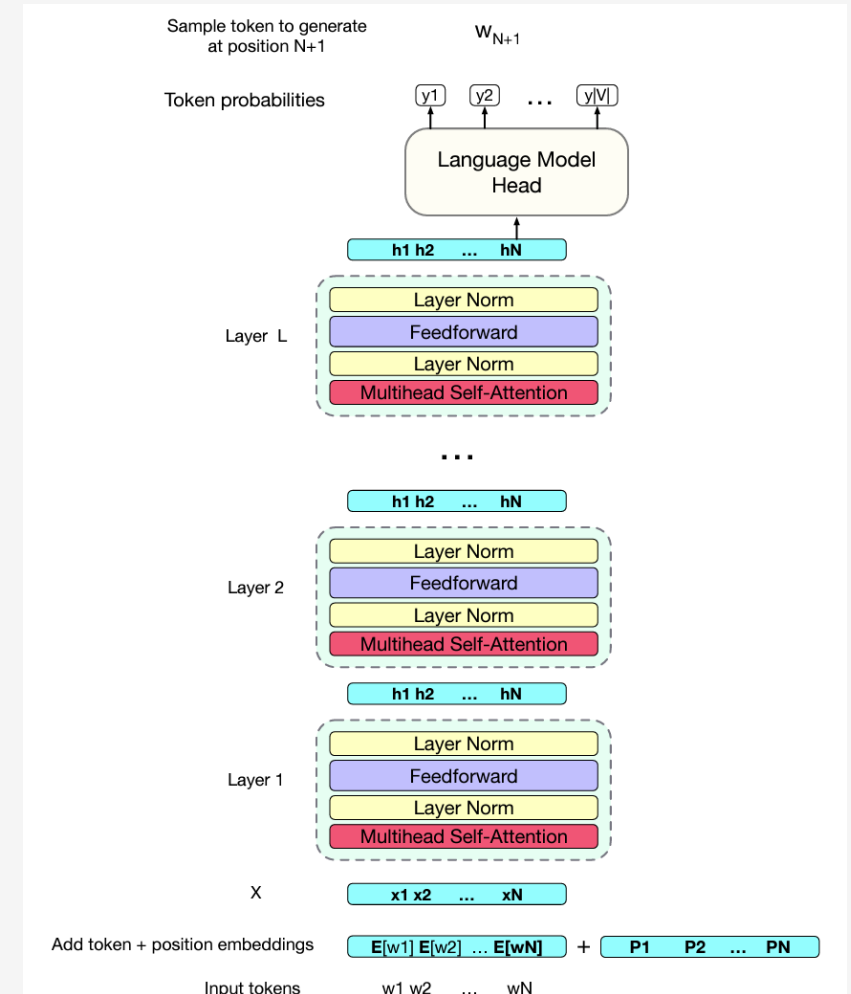
# Language modeling head

- Language modeling
  - Efficient for learning representations
  - Self-supervised
- Project  $h_N$  to vocabulary size
  - Do we know any computational tricks for that?
  - What would  $h_N^L$  look like?



# A final transformer representation for LM

- Token + positional embedding
- Multiple stacked transformer blocks
- A classification head
- Language modeling with weight tying and sampling



# Conclusions

# Encoder decoder models

---

- Dealing with tasks where input and output are mismatched
  - Different length
  - No 1 to 1 alignment
- We use one model to encode the input (image, text in English)
- We use another model to generate text in target language
- Simple encoder decoder is based on RNNs/LSTMs



# Attention

---

- RNNs have problems with long-distance dependencies
- Decoding from a single hidden state is restricted
- Attention uses all hidden states and compares current decoder state
  - Dot product attention

# Transformers

---

- Self attention builds upon the attention from encoder-decoder
- Query, Key, Value project the input based on its function
- Multiheaded attention stacks multiple self attentions

# Transfer learning

---

- The goal of transformers – learn contextual (and text) representations and reuse
- The head of the transformer determines the task
- Multiple problems can be framed as classification or generation