

# Revision week

## Neural NLP – from embeddings to LLMs

Venelin Kovatchev

Lecturer in Computer Science

[v.o.kovatchev@bham.ac.uk](mailto:v.o.kovatchev@bham.ac.uk)

# Outline

---

- Word embeddings
- Compositionality and end-to-end networks
- Encoder-decoder and attention
- Transformers
- Transfer learning and types of transformers
- LLMs and RLHF finetuning

# Vector representations and word embeddings

# The distributional hypothesis and word embeddings

---

- The distributional hypothesis in semantics
  - What does it state?
  - How does it affect NLP?
- Word embeddings
  - Encoding words as “semantic” vectors

# Count based vector representations

---

- Obtain a large corpus in the language/domain of interest
- Define a context of co-occurrence
  - What contexts can you think of?
- Count and fill in a co-occurrence matrix
- Apply transformations (which?)

# Word2Vec

---


- Learning embeddings directly from text

- A simple neural architecture


- Two algorithms

- What is the difference between them?

CBoW



Skip-gram



# Negative Sampling

---

- Global objective: maximize the log probability of the dataset of size  $T$  with a context size  $c$

$$\frac{1}{T} \sum_{t=1}^T \sum_{-c \leq j \leq c, j \neq 0} \log p(w_{t+j} | w_t)$$

- Calculate the probability of every word given context (or the other way around)

$$p(w_O | w_I) = \frac{\exp(v'_{w_O} v_{w_I})}{\sum_{w=1}^W \exp(v'_w v_{w_I})}$$

- Training with a softmax over the whole vocabulary is expensive

- Convert the task into "classifying the correct objective" using a logistic

$$P(+|w, c) = \sigma(\mathbf{c} \cdot \mathbf{w}) = \frac{1}{1 + \exp(-\mathbf{c} \cdot \mathbf{w})}$$

# Calculating similarity

---

- How do we calculate similarity between vectors?
- Dot product

$$\text{dot product}(\mathbf{v}, \mathbf{w}) = \mathbf{v} \cdot \mathbf{w} = \sum_{i=1}^N v_i w_i = v_1 w_1 + v_2 w_2 + \dots + v_N w_N$$

- Cosine

$$\text{cosine}(\mathbf{v}, \mathbf{w}) = \frac{\mathbf{v} \cdot \mathbf{w}}{|\mathbf{v}| |\mathbf{w}|} = \frac{\sum_{i=1}^N v_i w_i}{\sqrt{\sum_{i=1}^N v_i^2} \sqrt{\sum_{i=1}^N w_i^2}}$$




# Post-training word embeddings

---

- We train embeddings using a logistic classification (with negative sampling)
- What happens with the logistic after the training?
- Can you think of another algorithm that uses similar approach?

# ELMO

- What is the most important difference between ELMO and word2vec?
- What makes ELMO representations “deep”?
- What is the training objective behind ELMO?



# Compositionality and End-to-end

# Feature engineering

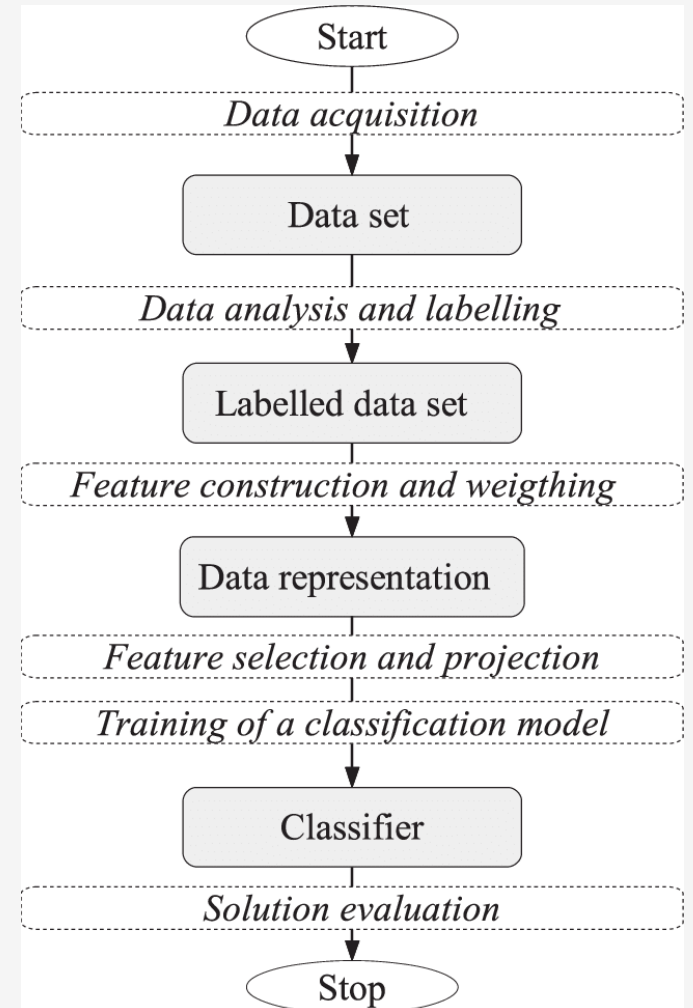
---

- Analyze the problem, the input, and the desired outcome
- Explore existing resources and processing techniques
- Select the most relevant features and feature-extraction methods
- Empirically test what works best

# Text classification using features

---

- Step by step process
- Involves active human engagement
- Feature selection and extraction
- Data is fed into a classifier (Logistic, NB, SVM)
- Iteratively improve feature selection and model (hyper) parameters



# Why not go end to end?

---

- Do we need full pipelines?
- Embeddings make it possible to “feed” text directly into models
- Is it possible to go fully end-to-end and eliminate
  - Accumulation of errors
  - Human labor and supervision
  - (In)compatibility issues between elements?

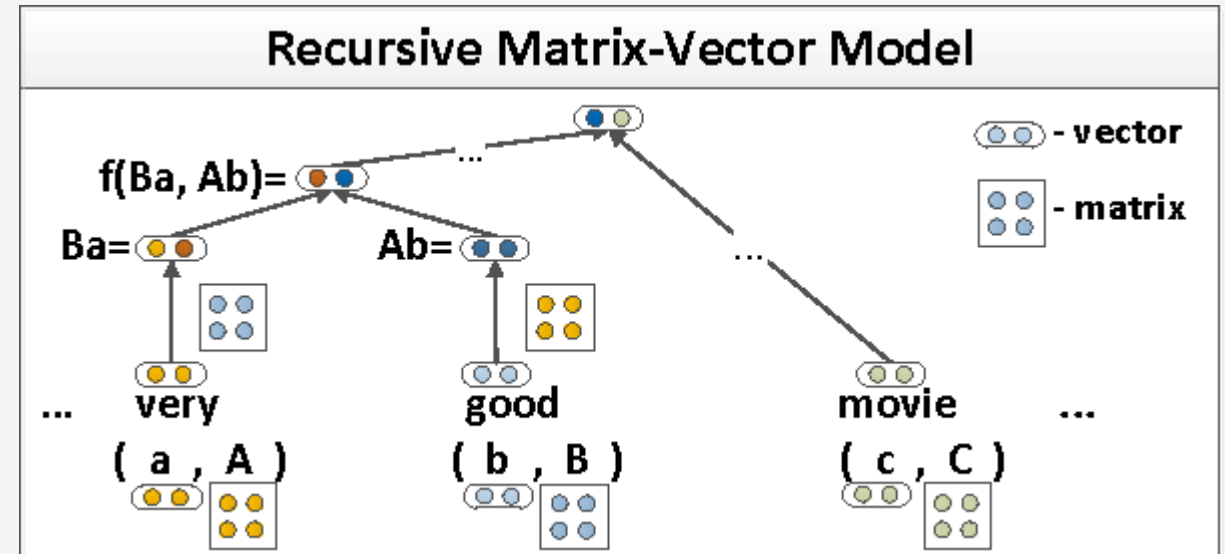
# Embeddings and the problem of compositionality

---

- Embeddings represent individual words
- NLP deals with processing texts
- How to go from word representations to text representations?

# Composing word meaning


- Vector operations
  - Vector addition
  - Pointwise vector multiplication
  - Vector concatenation
  - Complex matrix-vector operations
- Which of these operations consider text structure?
- Let a neural network do the compositionality





# Multi-layer perceptron


---



- $y = \text{softmax}(h_1 \cdot W_2 + b_2)$
- $h_1 = f(h_0 \cdot W_1 + b_1)$
- $h_0 = f(x \cdot W_0 + b_0)$
- Non-linear functions  $f$ :
  - Sigmoid:  $\sigma(x) = \frac{1}{\exp(-x)}$
  - Hyperbolic:  $\tanh(x) = \frac{1 - \exp(-2x)}{1 + \exp(-2x)}$
  - ReLU:  $\text{rect}(x) = \max(0, x)$


# Recurrent neural networks

- How to represent text of varying length?
- Copy the network for each word
- At each timestep  $t$ , input is  $x_t$  and  $h_{(t-1)}$
- I + like + tacos + "al pastor"
- Left to right, combining words one at a time
- Sequence classification
- Sequence to sequence



# RNNs (formally)

- At each timestep
  - Input  $x_t$  and previous hidden state  $h_{(t-1)}$
  - Three sets of weights:
    - input ( $W_x$ ), hidden ( $W_h$ ) , and output ( $W_o$ )
  - $h_t = f_h(W_x x_t + W_h h_{(t-1)})$
  - $O_t = f_o(W_o h_t)$
  - $J_t = f(O_t, y_t)$
- Pop-quiz: What are we predicting at O?



# Neural language modeling with RNNs

- More formally:
- $h_t = \tanh(W_x x_t + W_h h_{(t-1)})$
- $\hat{y}_t = \text{softmax}(W_y h_t)$
- $J_t = -\log \hat{y}_{t, \text{correct}}$  (- log prob the word at t+1)

- $J_{\text{sent}} = \frac{1}{T} \sum_{t=1}^T -\log \hat{y}_{t, \text{correct}}$

- Pop quiz: What is “weight tying” in RNNs?




Figure 9.6 from SLP, chapter 9

# Stacked and bidirectional RNNs

---

- The basic RNN is a powerful tool
- We can go deeper
- Stacking RNNs
- Bidirectional RNNs

# LSTM


---

- Popular variation of RNN that address native limitations
- Same recurrent concept (copy the network)
- Three different "gates":
  - Forget gate
  - Input gate
  - Output gate
- Gates manipulate the flow of information through time
  - Addressing conflicting objectives

# Understanding the gates

---

- All gates have the same format:
  - A feedforward layer followed by a sigmoid
- The gates are filtering out information at a certain part of the network
- Each gate has two important aspects:
  - How to calculate the filter
  - What is the input that the filter is applied to




# Convolutional neural networks (CNN)

---

- Another popular architecture for NLP
  - Deals with text of various size
- Inspired by computer vision success
- Applying filters of different size to the input (2,3,4) + pooling
  - Analogous to n-grams



# Convolutional neural networks (CNN)



# Compositionality


---

- How each network “composes” meaning
  - Feed-forward network
    - Linear combination of words (no order) + activation function
  - RNN/LSTM
    - Recursively, word by word (linear order)
  - CNN
    - Locally, similar to n-gram (proximity window, limited order importance)

# Encoder-decoder and attention

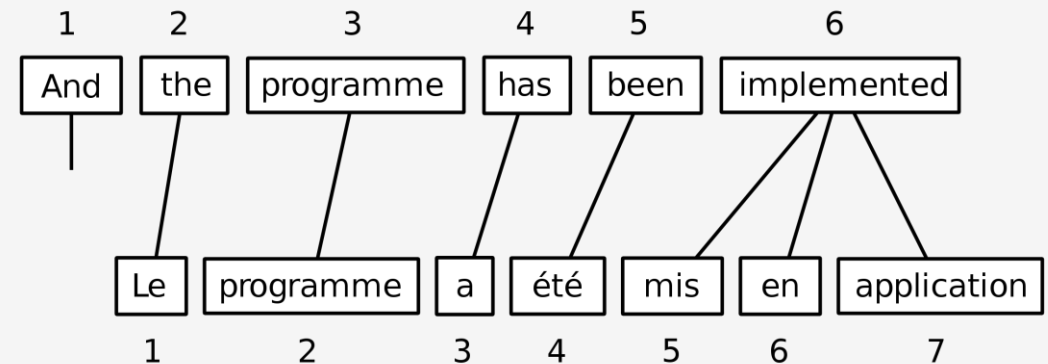
# Sequence labeling vs sequence-to-sequence problems

- Consider the following two tasks



- What is similar between them?

- What is different?



- How would you approach each of them?

# Key differences

---

- Same length vs different length
- One-to-one alignment vs no one-to-one alignment
- Local dependencies vs long-distance dependencies
  - Within the output
  - Between the input and the output

# Encoder decoder

---

- We use a model family called encoder-decoder
- Simple idea
  - Encoder “represents” the source (e.g., English)
  - Decoder “generates” the target (e.g., German)
- Can you suggest tasks that can use encoder-decoder?





Fig 9.16

# Using separate RNNs for encoder and decoder

- Train two models
- Pass the context at every step

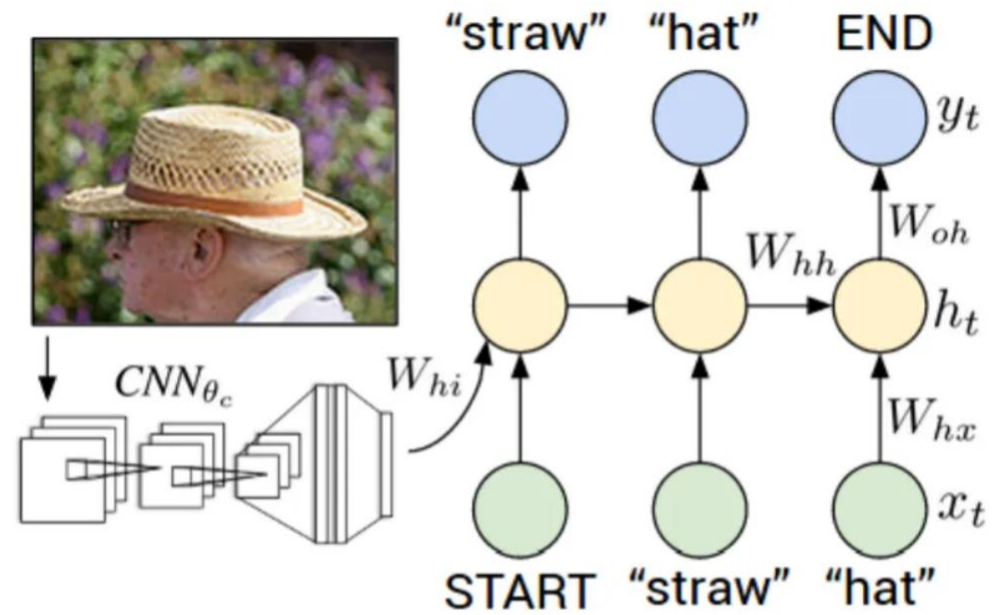
$$\mathbf{h}_t^d = g(\hat{y}_{t-1}, \mathbf{h}_{t-1}^d, \mathbf{c})$$

- Can you point a potential problem?
- What could improve this architecture?
- What is the purpose of the encoder?
- Should it be able to generate?



# Encoder decoder across modalities: image captioning

- The encoder and decoder “talk” via the context
- They don’t have to be the same type of model
- The modalities don’t have to match
  - Speech to text
  - Image to text




Deep Visual-Semantic Alignments for Generating Image Descriptions



# Attention – basic implementation

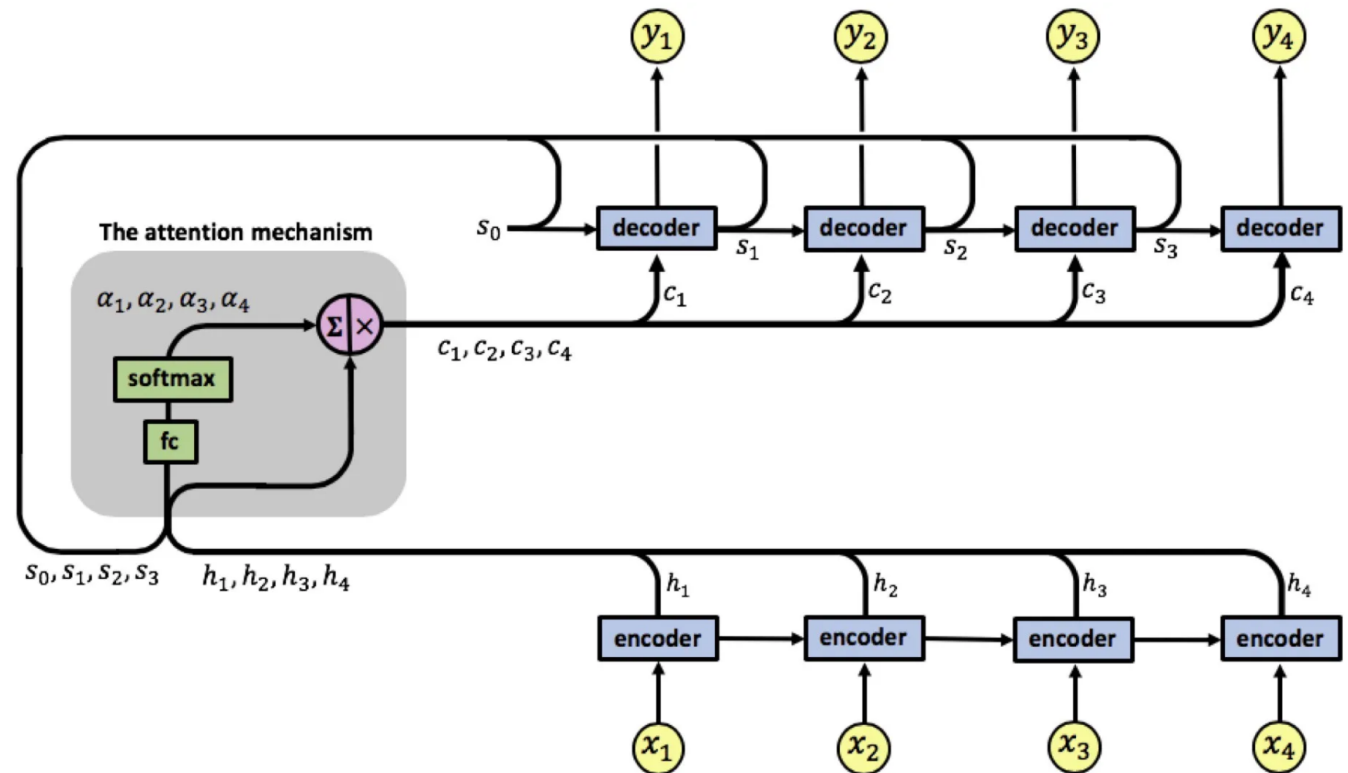
---

- Intuition: each token in the target should use a “personalized” context
  - Should have access to all hidden states in the encoder
  - The context should have a fixed (vector) length
- Weighted sum of all encoder hidden states
  - Calculated separately at each decoder step
  - Using the hidden state at (t-1)
- Dot product attention
  - Calculate the similarity between  $h_{(t-1)}$  and each encoder state  $h^e$
  - Use the similarity scores to calculate the weighted sum




# Visualization of RNN with attention

- RNN with attention
- Attention is learned via a simple FFN



# Visualizing attention

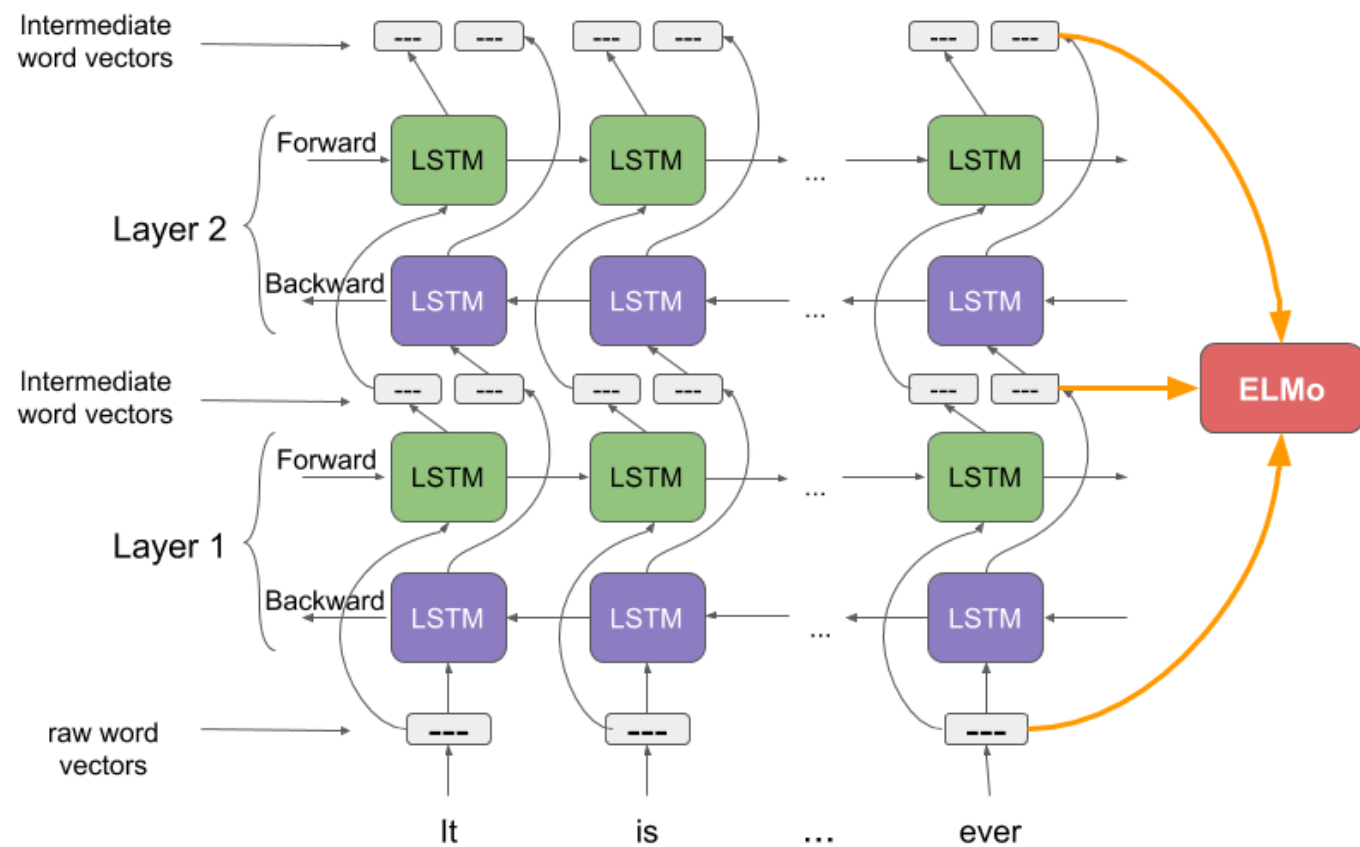
- Linear weights are interpretable
- We can see which word is more important
- Can we use attention for explainability?



# Transformers


# Elmo architecture

- How can we improve over that?




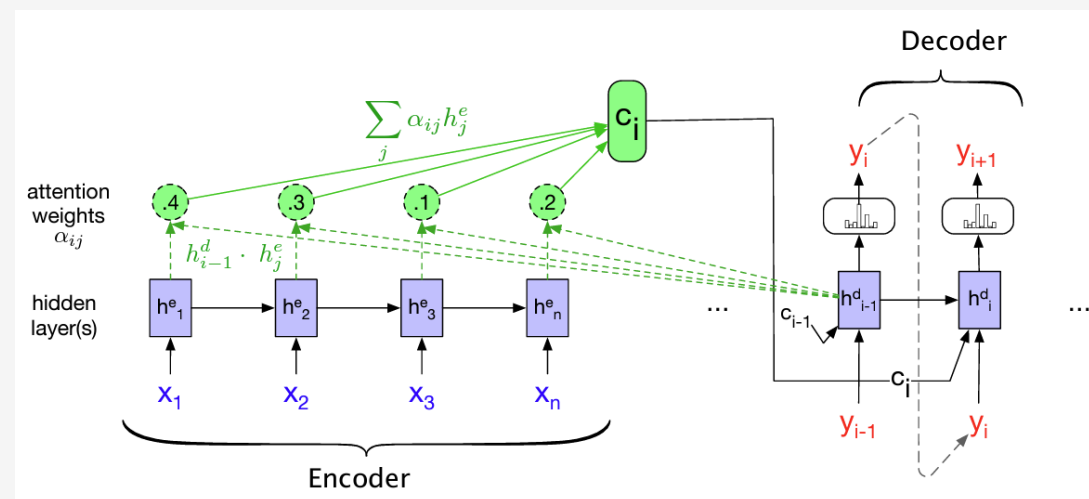
# Self attention

- Attention works better than RNN/LSTM for encoder-decoder models
- Can we use attention for a standalone network?



## Self attention (2)

- Self attention is a key concept in building transformers
- It applies the same approach as attention in encoder-decoder, but on itself



# Causal self attention (intuition)


- Similar to RNNs, we have a 1:1 input-output mapping
- Same basic approach as original attention
- Dot product + softmax + weighted sum

$$\text{score}(\mathbf{x}_i, \mathbf{x}_j) = \mathbf{x}_i \cdot \mathbf{x}_j$$

$$\begin{aligned}\alpha_{ij} &= \text{softmax}(\text{score}(\mathbf{x}_i, \mathbf{x}_j)) \quad \forall j \leq i \\ &= \frac{\exp(\text{score}(\mathbf{x}_i, \mathbf{x}_j))}{\sum_{k=1}^i \exp(\text{score}(\mathbf{x}_i, \mathbf{x}_k))} \quad \forall j \leq i\end{aligned}$$

$$\mathbf{a}_i = \sum_{j \leq i} \alpha_{ij} \mathbf{x}_j$$

- Which is the most similar token to  $\mathbf{x}_3$ ? What is the input to the first hidden layer?





# Decomposing input vectors

---

- We can use simple attention and it works
- Transformers introduce query, key, value
- What are they, why do we need them and how do we use them?
  - The “dictionary” analogy
  - A semantic explanation, grounded in NLP

# How to model asymmetric compositionality in attention?


---

- Classical attention (that we have seen) has 1:1 correspondence

- Dot product attention is commutative

- $a \cdot b = b \cdot a$

- $\text{score}(\text{"black"}, \text{"dog"}) = \text{score}(\text{"dog"}, \text{"black"})$



- Pop quiz: would "black" have the same importance on "dog" as "dog" would have on "black"?

# The query, key, value

---

- We project the input vector  $x$  to three vectors that serve different purpose: "query", "key", and "value"
- Two vector operations in the original attention:
  - "Score": for indexes  $i$  and  $j$ , calculate how important is  $x_j$  for  $x_i$ :  $\text{score}(x_i, x_j)$
  - "Scale": for index  $i$ , calculate the hidden state  $h_i$  as a weighted sum of  $x_1 \dots x_i$ :  $h_i = \sum_{j \leq i} \alpha_{ij} x_j$
- Each input vector  $x$  can have three different roles
  - Argument 1 in  $\text{score}()$  ["dog" in  $\text{score}(\text{"dog"}, \text{"black"})$ ] -> **query**
  - Argument 2 in  $\text{score}()$  ["dog" in  $\text{score}(\text{"black"}, \text{"dog"})$ ] -> **key**
  - The **value** used in scale to calculate the hidden state

# Query, Key, Value (formally)

---

- We learn three different matrices ( $W^Q, W^K, W^V$ )
- Every input vector  $x_i$  is projected to three different representations
  - $q_i = x_i W^Q ; k_i = x_i W^K ; v_i = x_i W^V$
- The new formula for score:  $\text{score}(\mathbf{x}_i, \mathbf{x}_j) = \mathbf{q}_i \cdot \mathbf{k}_j$
- The new formula for calculating weights:  $\mathbf{a}_i = \sum_{j \leq i} \alpha_{ij} \mathbf{v}_j$
- Pop quiz: which token will have the most impact on  $x_3$ ?


# The transformer self attention

1.  $\mathbf{q}_i = \mathbf{x}_i \mathbf{W}^Q; \mathbf{k}_i = \mathbf{x}_i \mathbf{W}^K; \mathbf{v}_i = \mathbf{x}_i \mathbf{W}^V$

2. and 3.  $\text{score}(\mathbf{x}_i, \mathbf{x}_j) = \frac{\mathbf{q}_i \cdot \mathbf{k}_j}{\sqrt{d_k}}$

4.  $\alpha_{ij} = \text{softmax}(\text{score}(\mathbf{x}_i, \mathbf{x}_j)) \quad \forall j \leq i$

5. and 6.  $\mathbf{a}_i = \sum_{j \leq i} \alpha_{ij} \mathbf{v}_j$



# Multiheaded self-attention

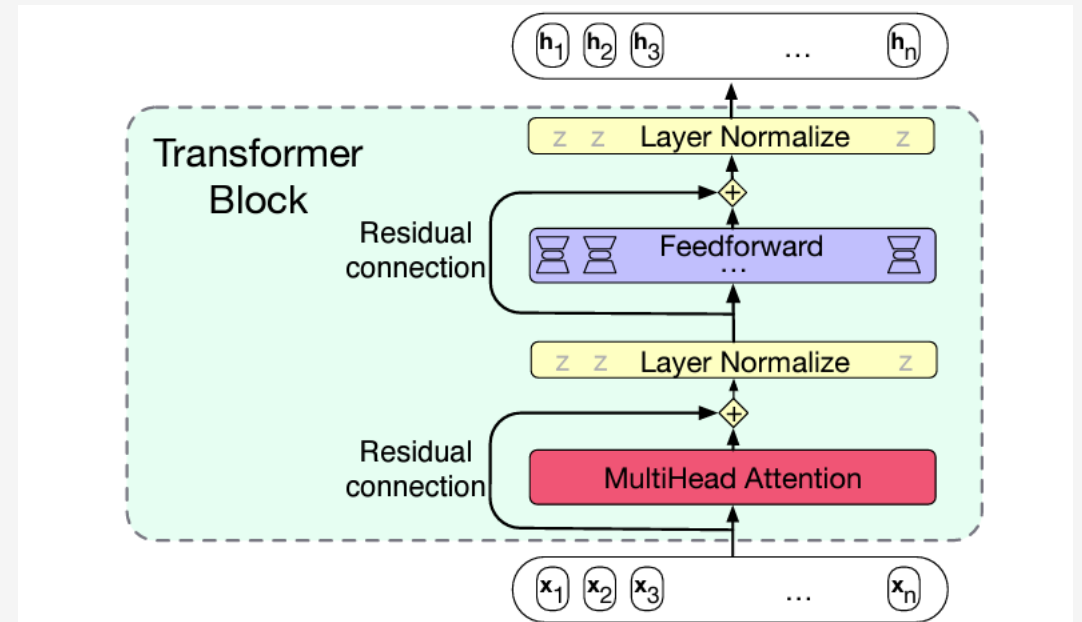
---

- Instead of using a single self attention, we can use multiple
  - Each "head" has its own weights  $W^Q, W^K, W^V$
  - The outputs of all heads are concatenated and projected to input dimensions
  - You can also think of multiheaded attention as "breaking" one big attention into specialized subsets
- Formally:

$$\begin{aligned}\mathbf{Q} &= \mathbf{XW}_i^Q ; \mathbf{K} = \mathbf{XW}_i^K ; \mathbf{V} = \mathbf{XW}_i^V \\ \mathbf{head}_i &= \text{SelfAttention}(\mathbf{Q}, \mathbf{K}, \mathbf{V}) \\ \mathbf{A} &= \text{MultiHeadAttention}(\mathbf{X}) = (\mathbf{head}_1 \oplus \mathbf{head}_2 \dots \oplus \mathbf{head}_h) \mathbf{W}^O\end{aligned}$$

# The transformer block

- Residual connection
  - Copy the input of a layer to its output
- Layer normalize
  - Rescale each  $x$  vector to 0-mean with  $STD=1$
- Positional feedforward
  - Apply the same fully connected FFN to each  $x$



# The transformer block (formally)


---

- Simplified representation
  - $O = \text{LayerNorm}(X + \text{MultiHeadAttention}(X))$
  - $H = \text{LayerNorm}(O + \text{FFN}(O))$
- You can change the order of operations in some implementations



# Encoding the Input. Positional Embeddings.

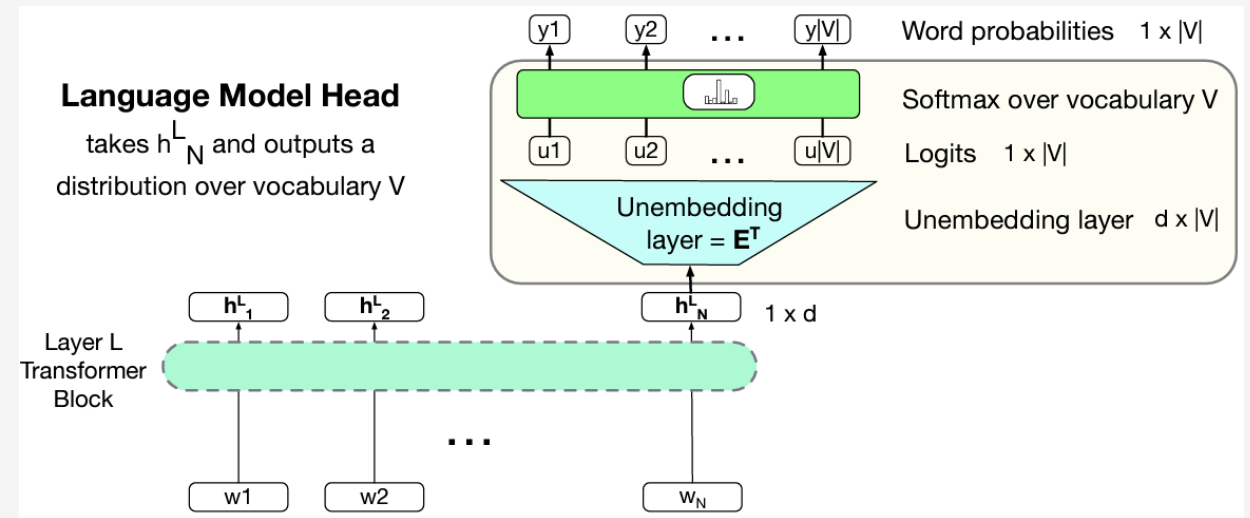
- Semantic embeddings
  - One-hot encoding maps to a row in a matrix
- Positional embeddings
  - One embedding for each position
  - Learnable; Same dimension as semantic
- Add semantic and positional embeddings



- Alternative techniques: use functions (sine/cosine); calculate relative positional embeddings

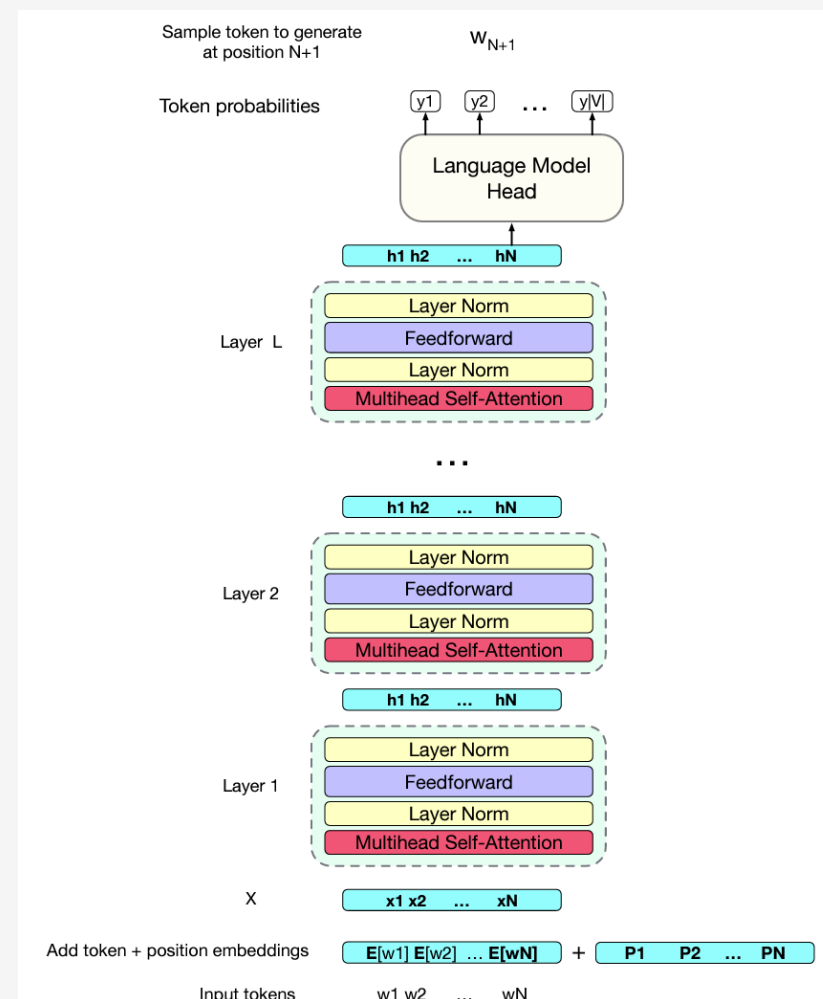
# Language modeling head

- Language modeling
  - Efficient for learning representations
  - Self-supervised
- Project  $h_N$  to vocabulary size
  - Do we know any computational tricks for that?
  - What would  $h_N^L$  look like?



# A final transformer representation for LM

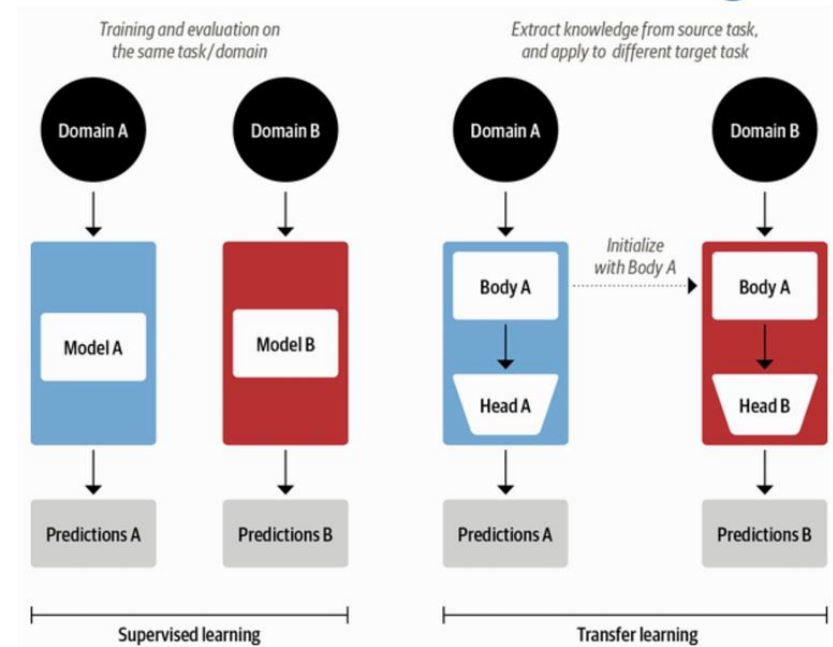
- Token + positional embedding
- Multiple stacked transformer blocks
- A classification head
- Language modeling with weight tying and sampling



# Transfer Learning and Types of Transformers

# Supervised learning vs Transfer learning

- What are the goals and benefits of transfer learning?
- What are some potential issues or risks?
- Are there any other paradigms that you can think of in that area?



Source: Lewis Tunstall, Leandro von Werra, and Thomas Wolf (2022), Natural Language Processing with Transformers: Building Language Applications with Hugging Face, O'Reilly Media.

# The decoder transformer: GPT

---

- GPT1 combines different concepts we know so far
  - The standard transformer block
  - Neural Language Modeling
  - Transfer learning capabilities
- Pop quiz: What kind of attention does it use?
- Intuition:
  - Generative pre-training
  - Discriminative finetuning

# Finetuning GPT

- After pretraining, use the hidden state at last layer
- Add a last linear layer with m neurons (m = number of classes)
- Predict the target class:


$$P(y|x^1, \dots, x^m) = \text{softmax}(h_l^m W_y).$$

- Maximize the probability of the correct labels (need labeled data)

$$L_2(\mathcal{C}) = \sum_{(x,y)} \log P(y|x^1, \dots, x^m).$$

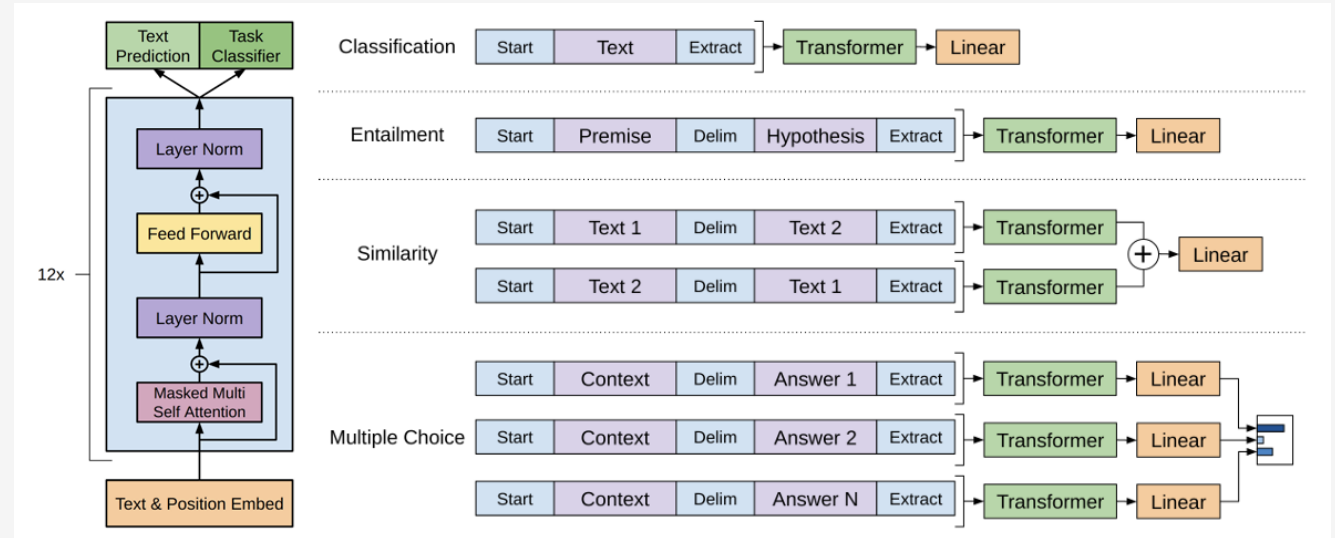
- Combining both losses together

$$L_3(\mathcal{C}) = L_2(\mathcal{C}) + \lambda * L_1(\mathcal{C})$$



# Task specific input transformations

- Task reformulating
- Using special tokens (sep, start/end)
- Comparing separate "streams"
- Task design is a non-trivial task
  - Task formulation; Data format; Metrics and Evaluation





# The encoder transformer: BERT


---

- The original encoder-only transformer
- An English-only sub-word vocabulary consisting of 30,000 tokens
  - Most of the modern algorithms use subwords tokenizers and embeddings
- 768 hidden size
- 12 layers, 12 heads in each multi-head attention
- 100M parameters
- Trained on two tasks: Masked Language Modeling and Next Sentence Prediction

# The encoder transformer

---

- The encoder in “attention is all you need”
- Same architecture as the decoder
- Bi-directional self-attention
  - All key/query values, no masking
- Better for encoding source information



# Masked language modeling objective

---

- Based on “cloze” tasks:
  - “Can I have a \_\_\_\_ of water, please?”
  - Does that remind you of something?
- Masked Language Modeling (MLM)
  - Randomly sample tokens from the text and perform alternations
  - Predict the original inputs for each position

# Next Sentence Prediction

---

- MLM predicts relationships between words
- Transformers want to also process sentences
- Next sentence prediction task
  - Given two sentences, predict whether they are a pair of adjacent sentences



# Next sentence prediction. The CLS token.

---

- Next sentence prediction
  - 50 % true adjacent pairs
  - A special [CLS] token added at the beginning
  - A special [SEP] token added between texts
  - Special "sentence position" (first/second) are added to input
- When predicting the sentence relation, we use the CLS as an input to softmax

# Adapting other tasks to work with BERT

- How would we perform paraphrase identification?
- What is the input/output/classification process?
- Performing other tasks:
  - Extractive QA
  - Sequence labeling



# The encoder-decoder model

- Encoder
  - Bi-directional attention can “see” all tokens
  - Follows the architecture we have seen last week
- Decoder
  - Causal attention
  - Additional Multiheaded Attention
    - Why do we need it?
    - What would be the Q, K, V used by it?

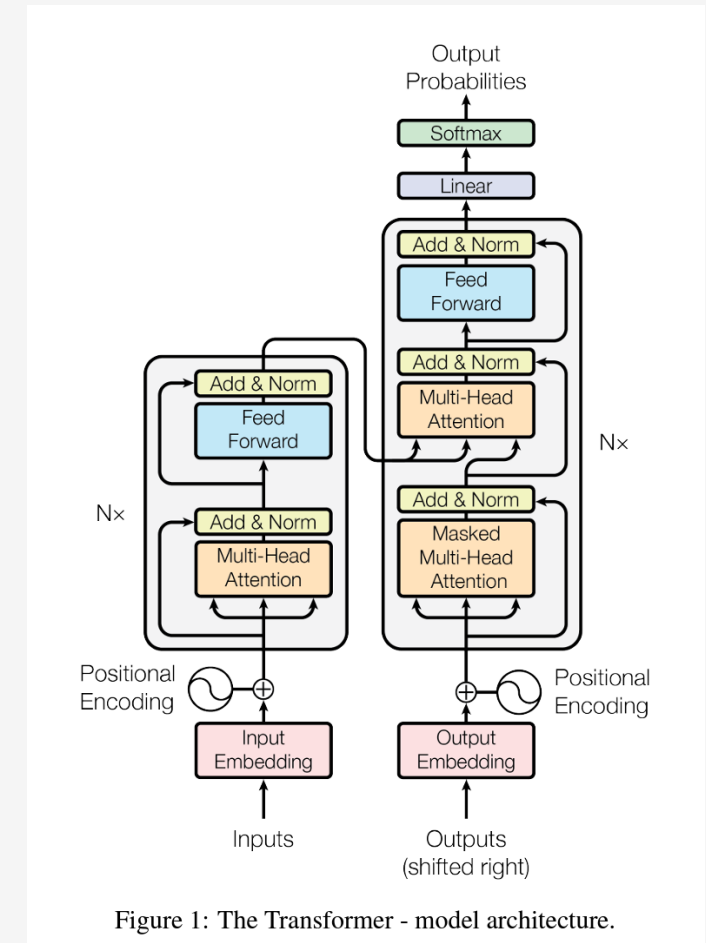
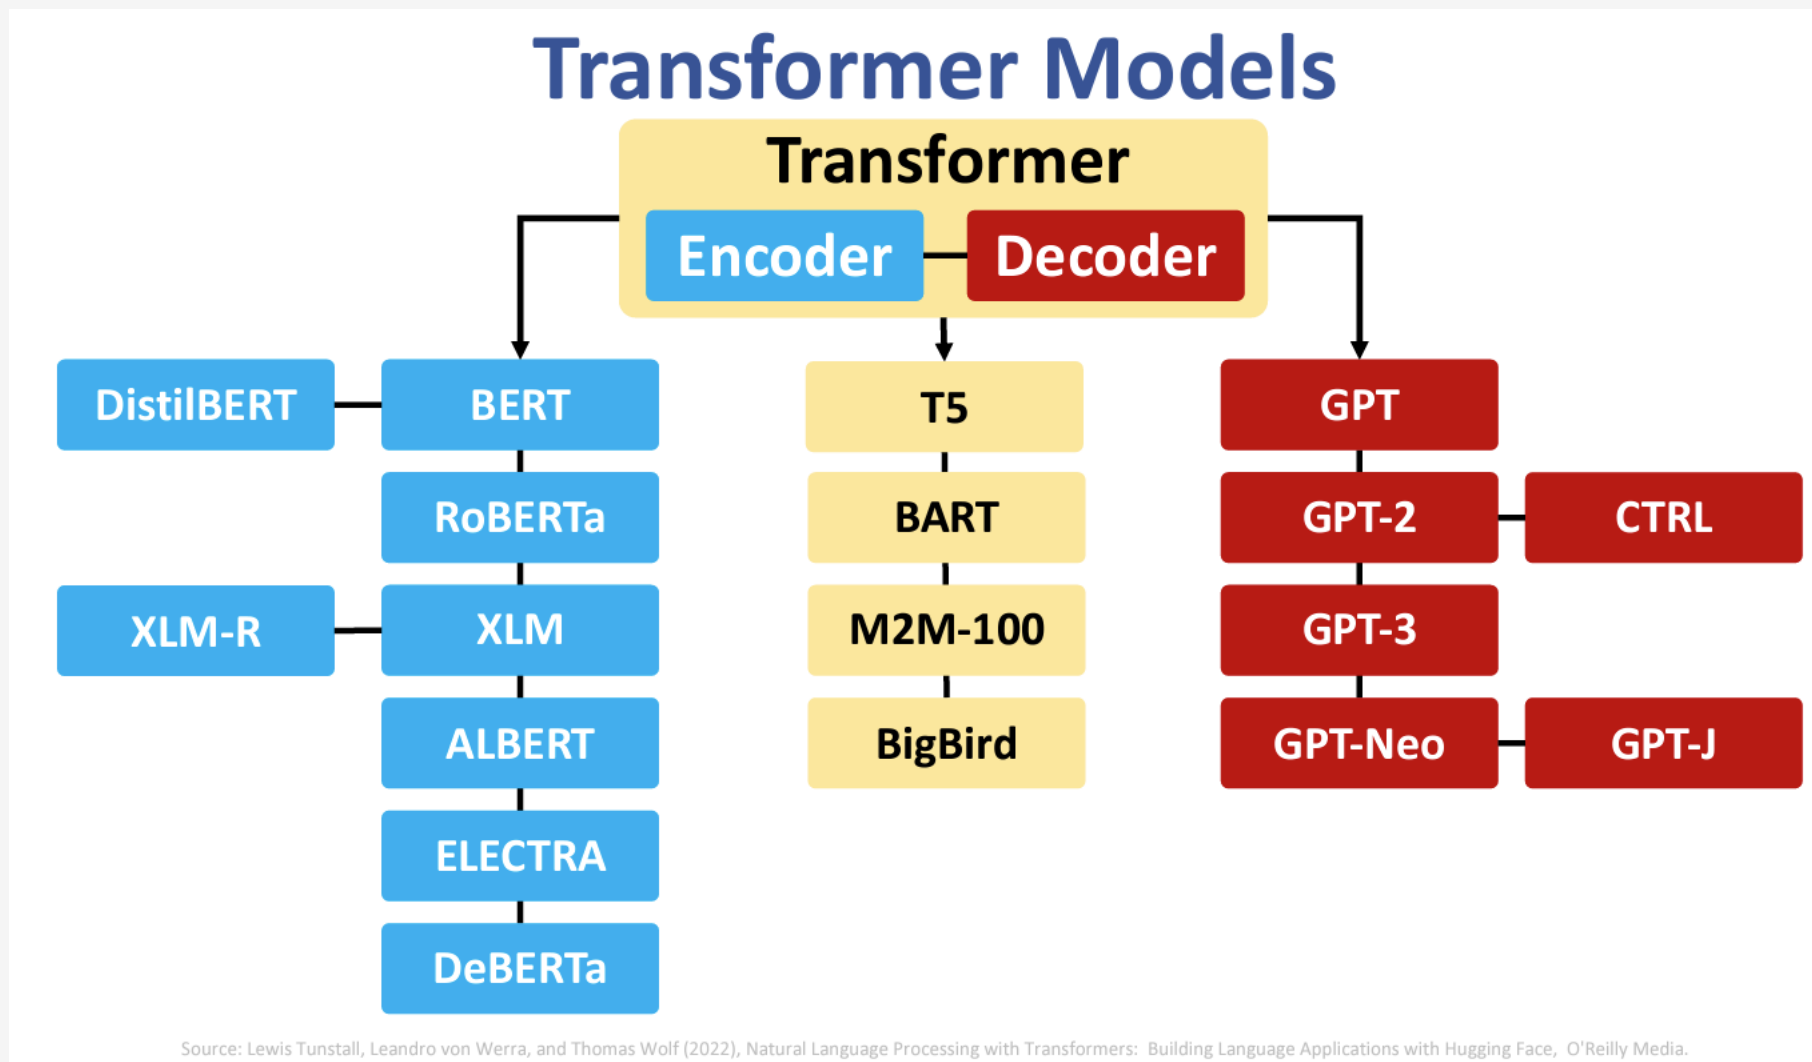


Figure 1: The Transformer - model architecture.

# The transformer family tree





# In-context learning and RLHF finetuning

# In-context learning

---

- Taking transfer learning to the extreme
- Using the input to specify the task
  - "What is the sentiment of the following text: I like this movie, it's the best in the Avengers series!"
  - "Do those sentences contradict each other: I bike to work every day. <SEP> I drive to work every day."
- Emerging property
  - A by-product of scaling the model above a certain size

# Zero- One- and Few-shot learning

- Three different experimental conditions
- No gradient update or finetuning
- The only difference – number of examples

The three settings we explore for in-context learning

## Zero-shot

The model predicts the answer given only a natural language description of the task. No gradient updates are performed.

```
1 Translate English to French: ← task description
2 cheese => ..... ← prompt
```

## One-shot

In addition to the task description, the model sees a single example of the task. No gradient updates are performed.

```
1 Translate English to French: ← task description
2 sea otter => loutre de mer ← example
3 cheese => ..... ← prompt
```

## Few-shot

In addition to the task description, the model sees a few examples of the task. No gradient updates are performed.

```
1 Translate English to French: ← task description
2 sea otter => loutre de mer ← examples
3 peppermint => menthe poivrée ←
4 plush girafe => girafe peluche ←
5 cheese => ..... ← prompt
```

# InstructGPT – training models to follow instructions

---

- Scale is not everything
  - Hallucinations
  - Toxicity
  - Lack of helpfulness
- Improve the training (and evaluation) procedures
- “Align” models with their users

# The LM training objective

---

- Language modeling is not “following instructions”
- Language modeling does not take (individual) preferences
- Every training sequence is equally important
- Preference in output (in language modeling) depends on
  - The observed frequency in training data
  - The sampling strategy
- Few- and Zero-shot learning are an “emerging” side effect, not an intentionally defined goal

# The training process of InstructGPT


## Step 1

**Collect demonstration data,  
and train a supervised policy.**


A prompt is  
sampled from our  
prompt dataset.



A labeler  
demonstrates the  
desired output  
behavior.




This data is used  
to fine-tune GPT-3  
with supervised  
learning.



## Step 2

**Collect comparison data,  
and train a reward model.**


A prompt and  
several model  
outputs are  
sampled.



A labeler ranks  
the outputs from  
best to worst.



This data is used  
to train our  
reward model.



## Step 3

**Optimize a policy against  
the reward model using  
reinforcement learning.**


A new prompt  
is sampled from  
the dataset.



The policy  
generates an output.



The reward model  
calculates a  
reward for the output.



The reward is  
used to update  
the policy  
using PPO.

