

# Week 2 Note

---

## 2.1 Gradient descent(GD)

---

### 2.1.1 Introduction

- The gradient vector at a point  $x$  points in the direction of greatest increase of the function  $f$ : each element of the gradient shows how fast  $f(x)$  is changing

- Example:

$$f(\vec{x}) = f(x_1, x_2) = 6x_1^2 + 4x_2^2 - 4x_1x_2$$

- The gradient vector is

$$\nabla f(x_1, x_2) = \begin{pmatrix} 12x_1 - 4x_2 \\ 8x_2 - 4x_1 \end{pmatrix}$$

### 2.1.2 Optimisation

- Optimisation algorithm
  1. Start with a point  $w$ (initial guess)
  2. Find a direction  $d$  to move on
  3. Determine how far ( $\eta$ ) to move along  $d$
  4. Update:  $w = w + \eta d$

### 2.1.3 Minimisation

- it is an iterative algorithm, starting from  $\vec{w}^{(0)}$  and producing a new  $\vec{w}^{(t+1)}$  at each iteration as:

$$\vec{w}^{(t+1)} = \vec{w}^{(t)} - \eta_t \nabla C(\vec{w}^{(t)})$$

where  $t = 0, 1, \dots, T$

◼  $\eta_t > 0$  is the learning rate or step size

### 2.1.4 Choosing a step size

- Choosing a step size
  - If step size is too large - algorithm may never converge
  - If step size is too small - convergence may be very slow

### 2.1.5 GD for least squares regression

- Least squares regression
  - For least square regression, let's recall:

$$C(\vec{w}) = \frac{1}{2n} (\underbrace{\vec{w}^T X^T X \vec{w}}_{\text{quadratic}} - \underbrace{2\vec{w}^T X^T \vec{y}}_{\text{linear}} + \underbrace{\vec{y}^T \vec{y}}_{\text{constant}})$$

- The gradient is computed as:

$$\nabla C(\vec{w}) = \frac{1}{n} (X^T X \vec{w} - X^T \vec{y})$$

- GD updates  $\vec{w}^{(t)}$  by

$$\begin{aligned}\vec{w}^{(t+1)} &= \vec{w}^{(t)} - \eta \nabla C(\vec{w}^{(t)}) \\ \vec{w}^{(t+1)} &= \vec{w}^{(t)} - \frac{\eta}{n} (X^T X \vec{w}^{(t)} - X^T \vec{y})\end{aligned}$$

## 2.2 Stochastic gradient descent(SGD)

---

### 2.2.1 Introduction

- Replace the computationally expensive term  $\nabla C(\vec{w}^{(t)})$  by a stochastic gradient computed on a random example

### 2.2.2 Algorithm

- Algorithm
  1. Initialise the weights  $\vec{w}^{(0)}$
  2. For  $t = 0, 1, \dots, T$ 
    - Draw  $i_t$  from  $1, 2, \dots, n$  with equal probability
    - Compute stochastic gradient  $\nabla C_{i_t}(\vec{w}^{(t)})$  and update


$$\vec{w}^{(t+1)} = \vec{w}^{(t)} - \eta_t \nabla C_{i_t}(\vec{w}^{(t)})$$

### 2.2.3 SGD vs GD

- GD requires more computations per iteration but makes a good progress per iteration
  - It needs few iterations to get a good solution
- SGD requires less computations per iteration but makes less update per iteration
  - Therefore, it needs more iterations to get a good solution
- GD and SGD cannot always dominate the other.
  - If we want high accuracy and  $n$  is small, then **GD** is better
  - If we want moderate accuracy and  $n$  is large, then **SGD** is better

### 2.2.4 Effect of learning rates

- If we choose a low learning rate, then SGD would converge very slowly
- If we choose a large learning rate, then SGD would not go further as we run more and more iterations
- If we choose a huge learning rate, then SGD would become unstable
- A typical choice is  $\eta_t = \frac{c}{\sqrt{t}}$ , where  $c$  is a parameter needed to tune



## 2.3 Minibatch SGD

### 2.3.1 Introduction

- Randomly select a batch of indices:  $B_t \subseteq \{1, 2, \dots, n\}$  and update the model

$$\vec{w}^{(t)} = \vec{w}^{(t)} - \frac{\eta_t}{b} \sum_{i \in B_t} \nabla C_i(\vec{w}^{(t)})$$

where  $b$  is the batch size

■ If  $b = 1$ , it is clear that minibatch SGD is SGD

### 2.3.2 Algorithm


- Let  $\{\eta_t\}$  be a sequence of step sizes
- Algorithm
  1. Initialise the weights  $\vec{w}^{(0)}$
  2. For  $t = 0, 1, \dots, T$ 
    - Randomly select a batch  $B_t \subseteq \{1, 2, \dots, n\}$  of size  $b$
    - Compute stochastic gradient  $\nabla C_i(\vec{w}^{(t)})$  with  $i \in B_t$  and update

$$\vec{w}^{(t)} = \vec{w}^{(t)} - \frac{\eta_t}{b} \sum_{i \in B_t} \nabla C_i(\vec{w}^{(t)})$$

### 2.3.3 minibatch selection

- There are two ways to sample the minibatch  $B_t$ 
  - sampling with replacement
  - sampling without replacement

### 2.3.4 Minibatch SGD vs SGD vs GD



## 2.4 Linear classification

---

### 2.4.1 Introduction

- Suppose we have

$$D = \{(\vec{x}^1, y^1), (\vec{x}^2, y^2), \dots, (\vec{x}^n, y^n)\}$$

and

$$y^i \in \{-1, +1\}$$

- To build a linear model to separate posite examples from negative examples

### 2.4.2 0-1 loss

$$L(\hat{y}, y) = II[\hat{y} \neq y] = \begin{cases} 1 & \text{if } \hat{y} \neq y \\ 0 & \text{otherwise} \end{cases}$$

- The behaviour of a model on  $D$  can be measured by:

$$C(\vec{w}) = \frac{1}{n} \sum_{i=1}^n II[\text{sgn}(\vec{w}^T \vec{x}^i) \neq y^i]$$

The surrogate of  $C(\vec{w})$  which is easy to minimise

### 2.4.3 margin-based loss

- Margin
  - The margin of a model  $\vec{w}$  on an example  $(\vec{x}, y)$  is defined as  $y\vec{w}^T \vec{x}$
- A model with a positive margin means a correct prediction
- A model with a negative margin means an incorrect prediction

$$\begin{aligned} \hat{y} &= \text{sgn}(\vec{w}^T x) \\ \bar{y} &= \vec{w}^T x \end{aligned}$$

Margin(边距), 正确分类的情况下, 距离决策边界越远的数据预测的越准确

This further motivateds a model with large margin: a large margin means the model is robust in making a correct prediction

- Loss function of the form:

$$L(\hat{y}, y) = g(y\hat{y})$$

where  $g$  is decreasing

- minimising L means maximising the margin
  - Maximising the margin means getting a model with good performance

### 2.4.4 Surrogate loss

- We mainly consider

$$g(t) = \frac{1}{2}(\max(0, 1 - t))^2 = \begin{cases} \theta & \text{if } t \geq 1 \\ \frac{1}{2}(1 - t)^2 & \text{otherwise} \end{cases}$$

- The loss function becomes

$$L(\hat{y}, y) = \frac{1}{2}(\max\{0, 1 - y\hat{y}\})^2 = \frac{1}{2}(\max\{0, 1 - y\vec{w}^T \vec{x}\})^2$$

- The behaviour on  $D$  is quantified by


$$C(\vec{w}) = \frac{1}{2n} \sum_{i=1}^n (\max\{0, 1 - y^i \vec{w}^T \vec{x}^i\})^2$$

- We further get


$$\nabla C_i(\vec{w}) = \begin{cases} \theta & \text{if } y^i \vec{w}^T \vec{x}^i \geq 1 \\ (\vec{w}^T \vec{x}^i - y^i) \vec{x}^i & \text{otherwise} \end{cases}$$

- There are some other choices, including:


- $g(t) = \max\{0, 1 - t\}$ 
    - $\min_{\vec{w}} \frac{1}{n} \sum_{i=1}^n (\max\{0, 1 - y^i \vec{w}^T \vec{x}^i\})$
  - $g(t) = \log(1 + \exp(-t))$ 
    - $\min_{\vec{w}} \frac{1}{n} \sum_{i=1}^n \log(1 + \exp(-y^i \vec{w}^T \vec{x}^i))$



$g(t) = \max\{0, 1 - t\}$



$g(t) = \frac{1}{2}(\max\{0, 1 - t\})^2$



$g(t) = \log(1 + \exp(-t))$

## 2.4.5 SGD for linear classification

- Consider linear classification with the hinge loss

$$C_i(\vec{w}) = \max\{0, 1 - y^i \vec{w}^T \vec{x}^i\}$$

- The formula for SGD update:

$$\vec{w}^{(t+1)} = \begin{cases} \vec{w}^{(0)} & \text{if } y^{i_t} \vec{w}^{(t)T} \vec{x}^{i_t} \geq 1 \\ \vec{w}^{(0)} + \eta_t y^{i_t} \vec{x}^{i_t} & \text{otherwise} \end{cases}$$

- Consider regularisation in the loss function

$$C_i(\vec{w}) = \frac{1}{2}(\max\{0, 1 - y^i \vec{w}^T \vec{x}^i\})^2 + \frac{1}{2}\lambda \|\vec{w}\|_2^2$$

- The formula for SGD update:

$$\vec{w}^{(t+1)} = \begin{cases} \vec{w}^{(t)} - \lambda \vec{w}^{(t)} & \text{if } y^{i_t} \vec{w}^{(t)T} \vec{x}^{i_t} \geq 1 \\ \vec{w}^{(t)} - \lambda \vec{w}^{(t)} + \eta_t (1 - y^{i_t} \vec{w}^{(t)T} \vec{x}^{i_t}) y^{i_t} \vec{x}^{i_t} & \text{otherwise} \end{cases}$$