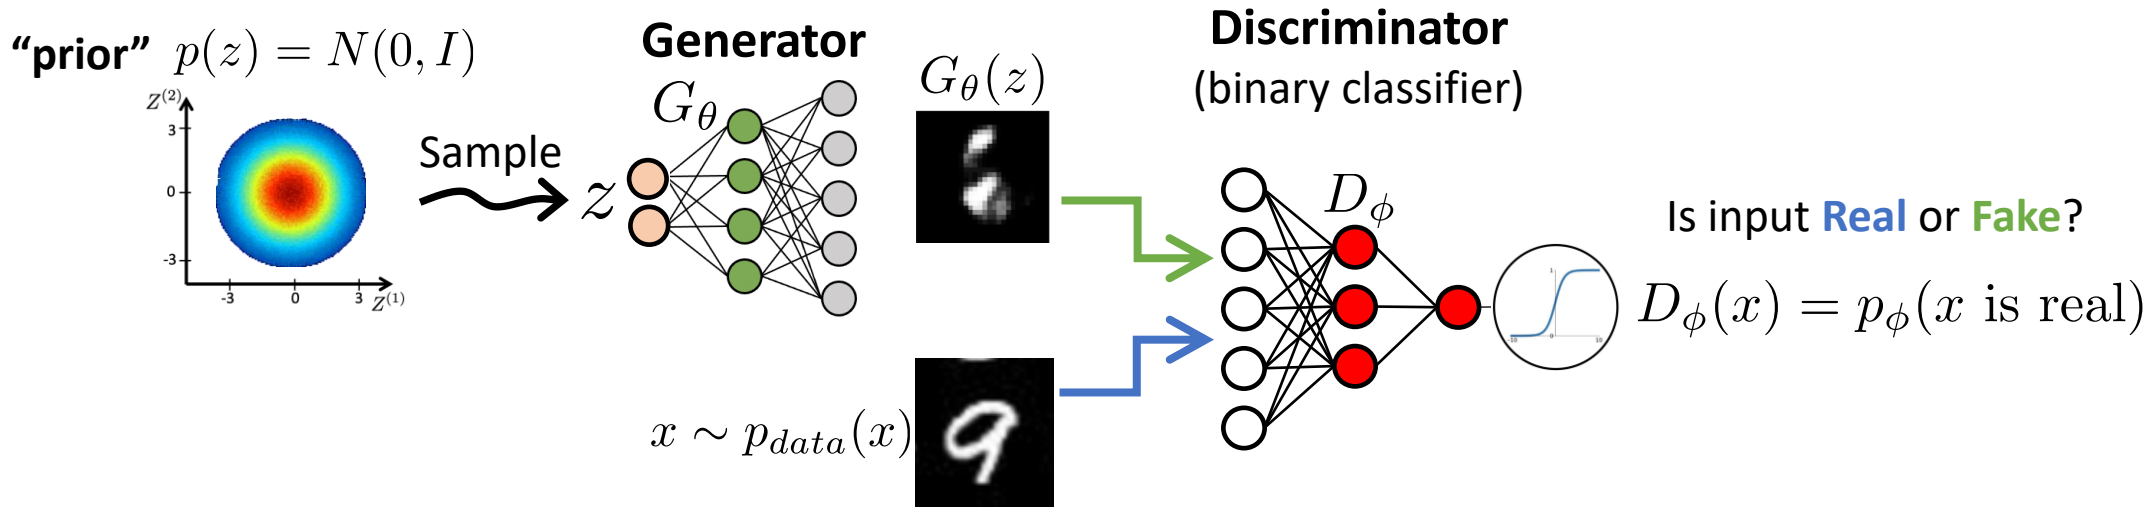



Limitations of basic GAN and Advanced Models

Generative Adversarial Networks



$$\{\theta', \phi'\} = \arg \min_{\theta} \max_{\phi} \mathbb{E}_{x \sim p_{data}(x)} [\log D_\phi(x)] + \mathbb{E}_{z \sim p(z)} [\log(1 - D_\phi(G_\theta(z)))]$$

Theoretically motivated
vs practical loss for G



Local Minima and Mode Collapse

Research on solutions:

- Better optimizers
- Better network architectures
- Better losses (Wasserstein, GAN-GP,...)
- ...

Two potential losses for G from previous video lecture, one offering better gradients.


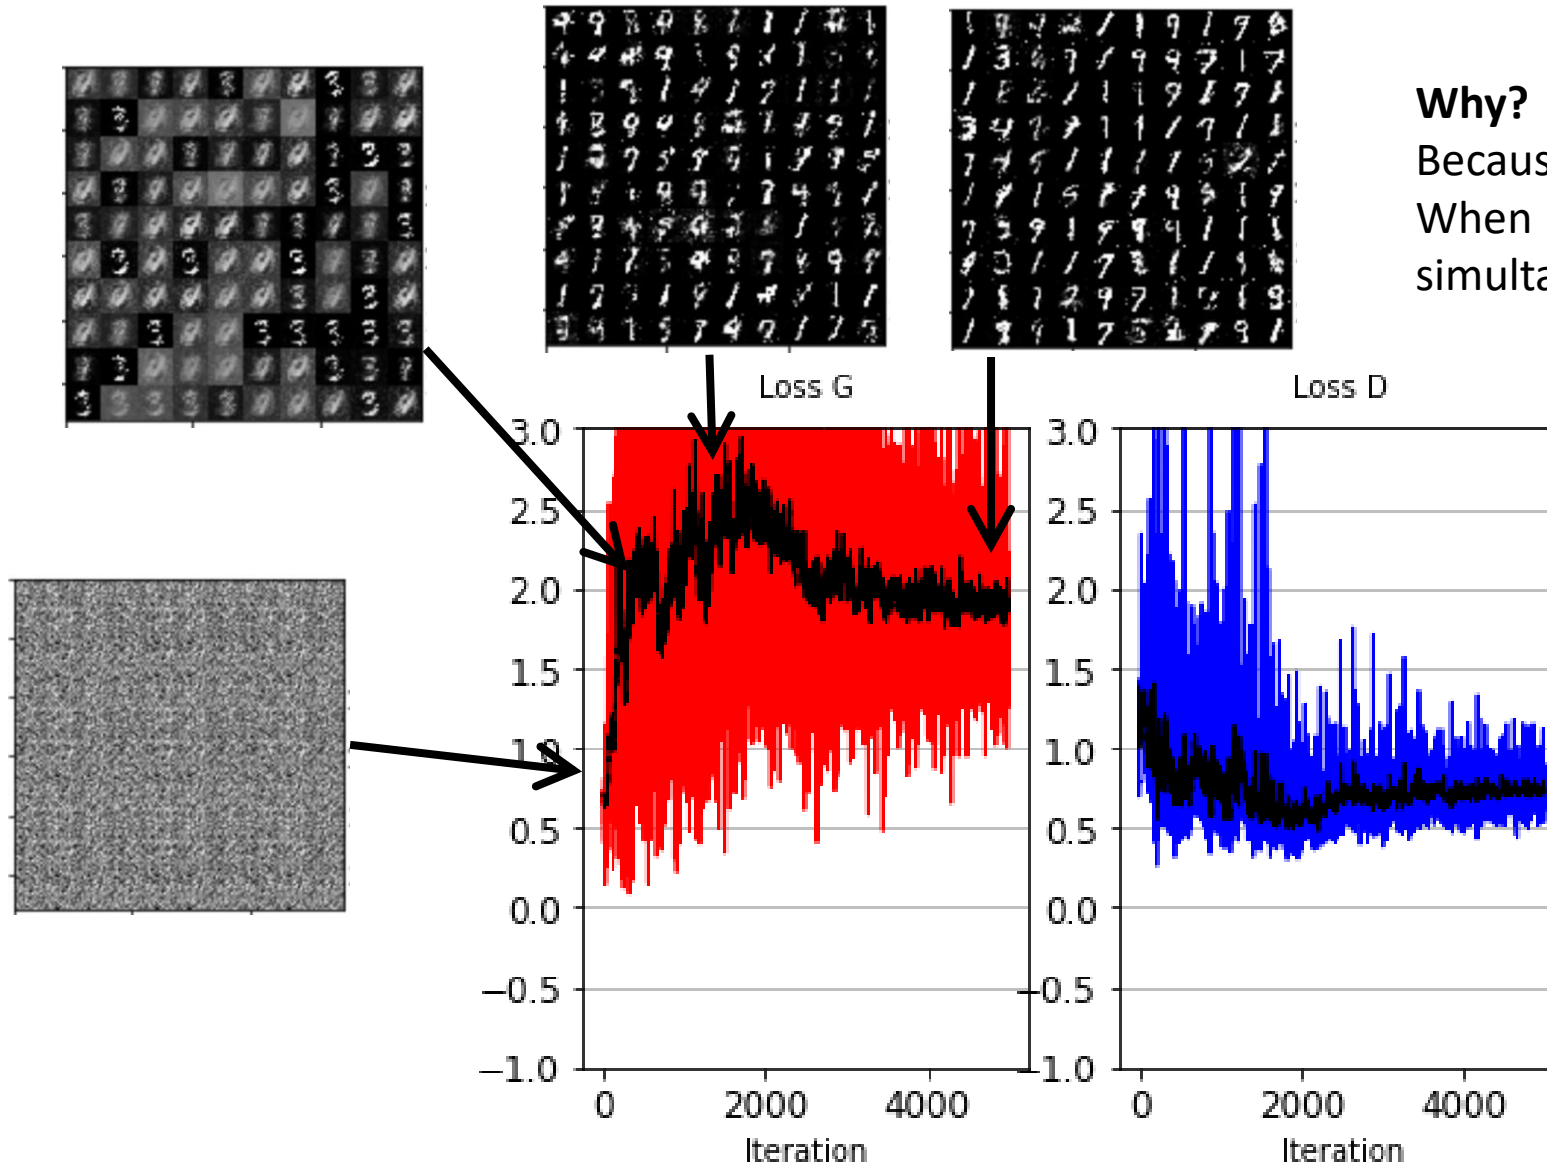


Image: Fedus et al, Many paths to equilibrium: GANS do not need to decrease a divergence at every step, 2018

Problem: Generator's Loss is not decreasing



Why?

Because D also improves during training. When D changes, if D reduces its loss, it simultaneously increases G's loss.

Training on MNIST
(from our Tutorial on GANs)

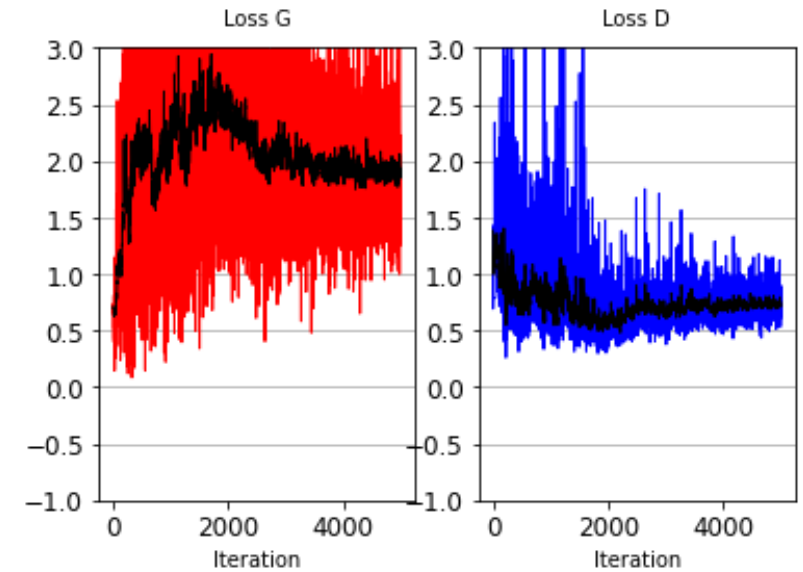
Discriminator's Accuracy as “Learned Loss function” for the Generator

Theoretical loss:

$$\mathcal{L}_G(z) = \log(1 - \underset{\text{red}}{D}_{\underset{\text{green}}{\phi}}(\underset{\text{green}}{G}_{\theta}(z)))$$

Practical loss:

$$\mathcal{L}_G(z) = -\log \underset{\text{red}}{D}_{\underset{\text{green}}{\phi}}(\underset{\text{green}}{G}_{\theta}(z))$$




Generator's loss is a function of the Discriminator and its parameters.

G's loss decreases when G improves during training and pushes $D(G(z))$ to 1, but also increases when D improves during training and pushes $D(G(z))$ to 0.

Therefore, even when results by G improve, we may see G's loss go up!

Problem: Generator's Loss is not decreasing



Why?

Because D also improves during training. When D changes, if D reduces its loss, it simultaneously increases G's loss.

Solution?

Open research on alternative GAN losses (e.g. Wasserstein GAN)

From: Arjovsky et al, Wasserstein GAN, 2017

Limited uses of basic GAN:

How to use the basic GAN for inferring latent variables (code) of a given sample?

How to use the basic GAN for altering a feature of a specific input sample?

Can we use basic GAN for compression?

How to use the basic GAN for pre-training a supervised classifier?

We cannot, because basic GAN has no encoder $x \rightarrow z$. Only generator $z \rightarrow x$

...but there are 1000s of extensions for solving these!

Advanced models for unsupervised learning

From: Makhzani et al, Adversarial Auto-Encoders, 2016

Adversarial Auto-Encoder

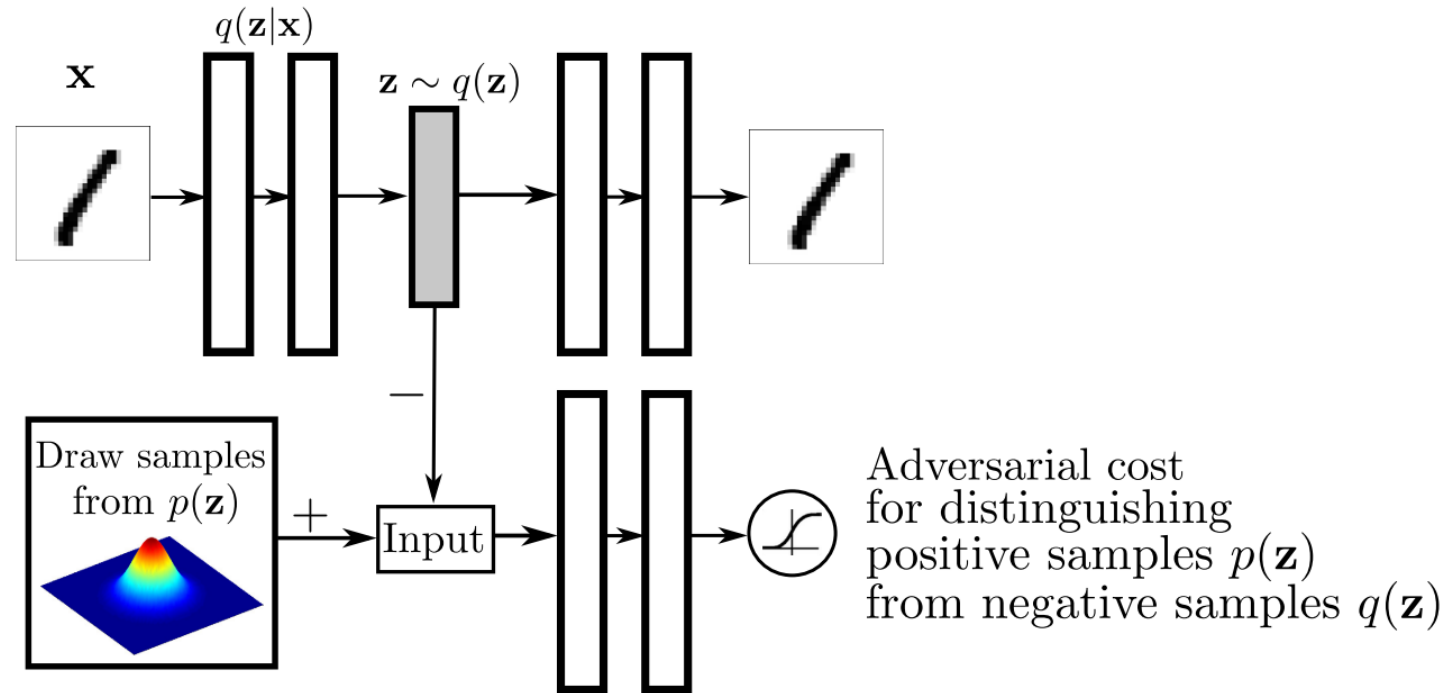


Figure 1: Architecture of an **adversarial autoencoder**. The top row is a standard autoencoder that reconstructs an image x from a latent code z . The bottom row diagrams a second network trained to discriminatively predict whether a sample arises from the hidden code of the autoencoder or from a sampled distribution specified by the user.

From: Makhzani et al, Adversarial Auto-Encoders, 2016

VAE-GAN


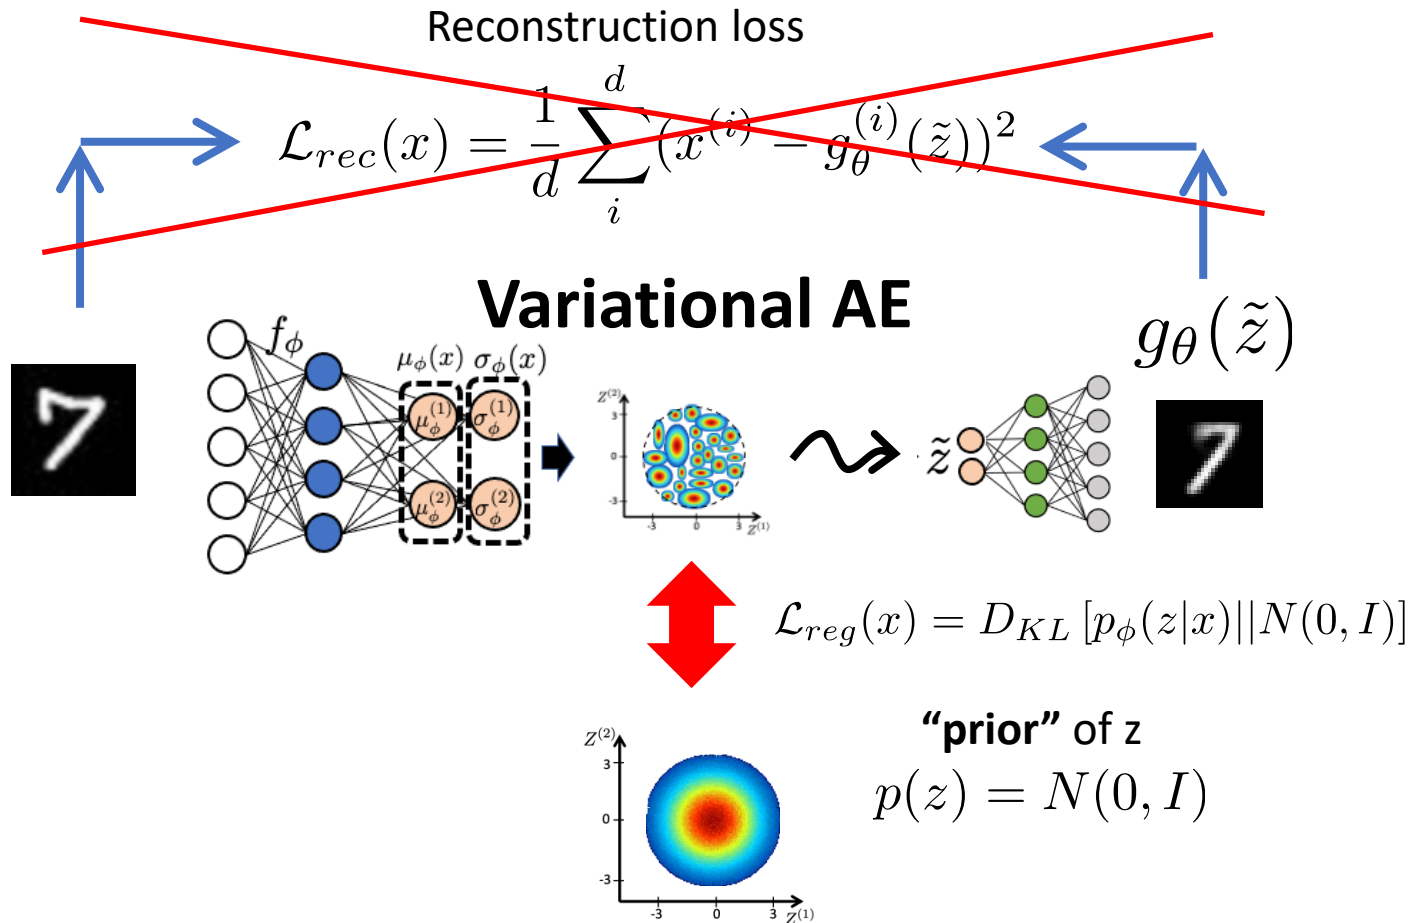


Figure 1. Overview of our network. We combine a VAE with a GAN by collapsing the decoder and the generator into one.

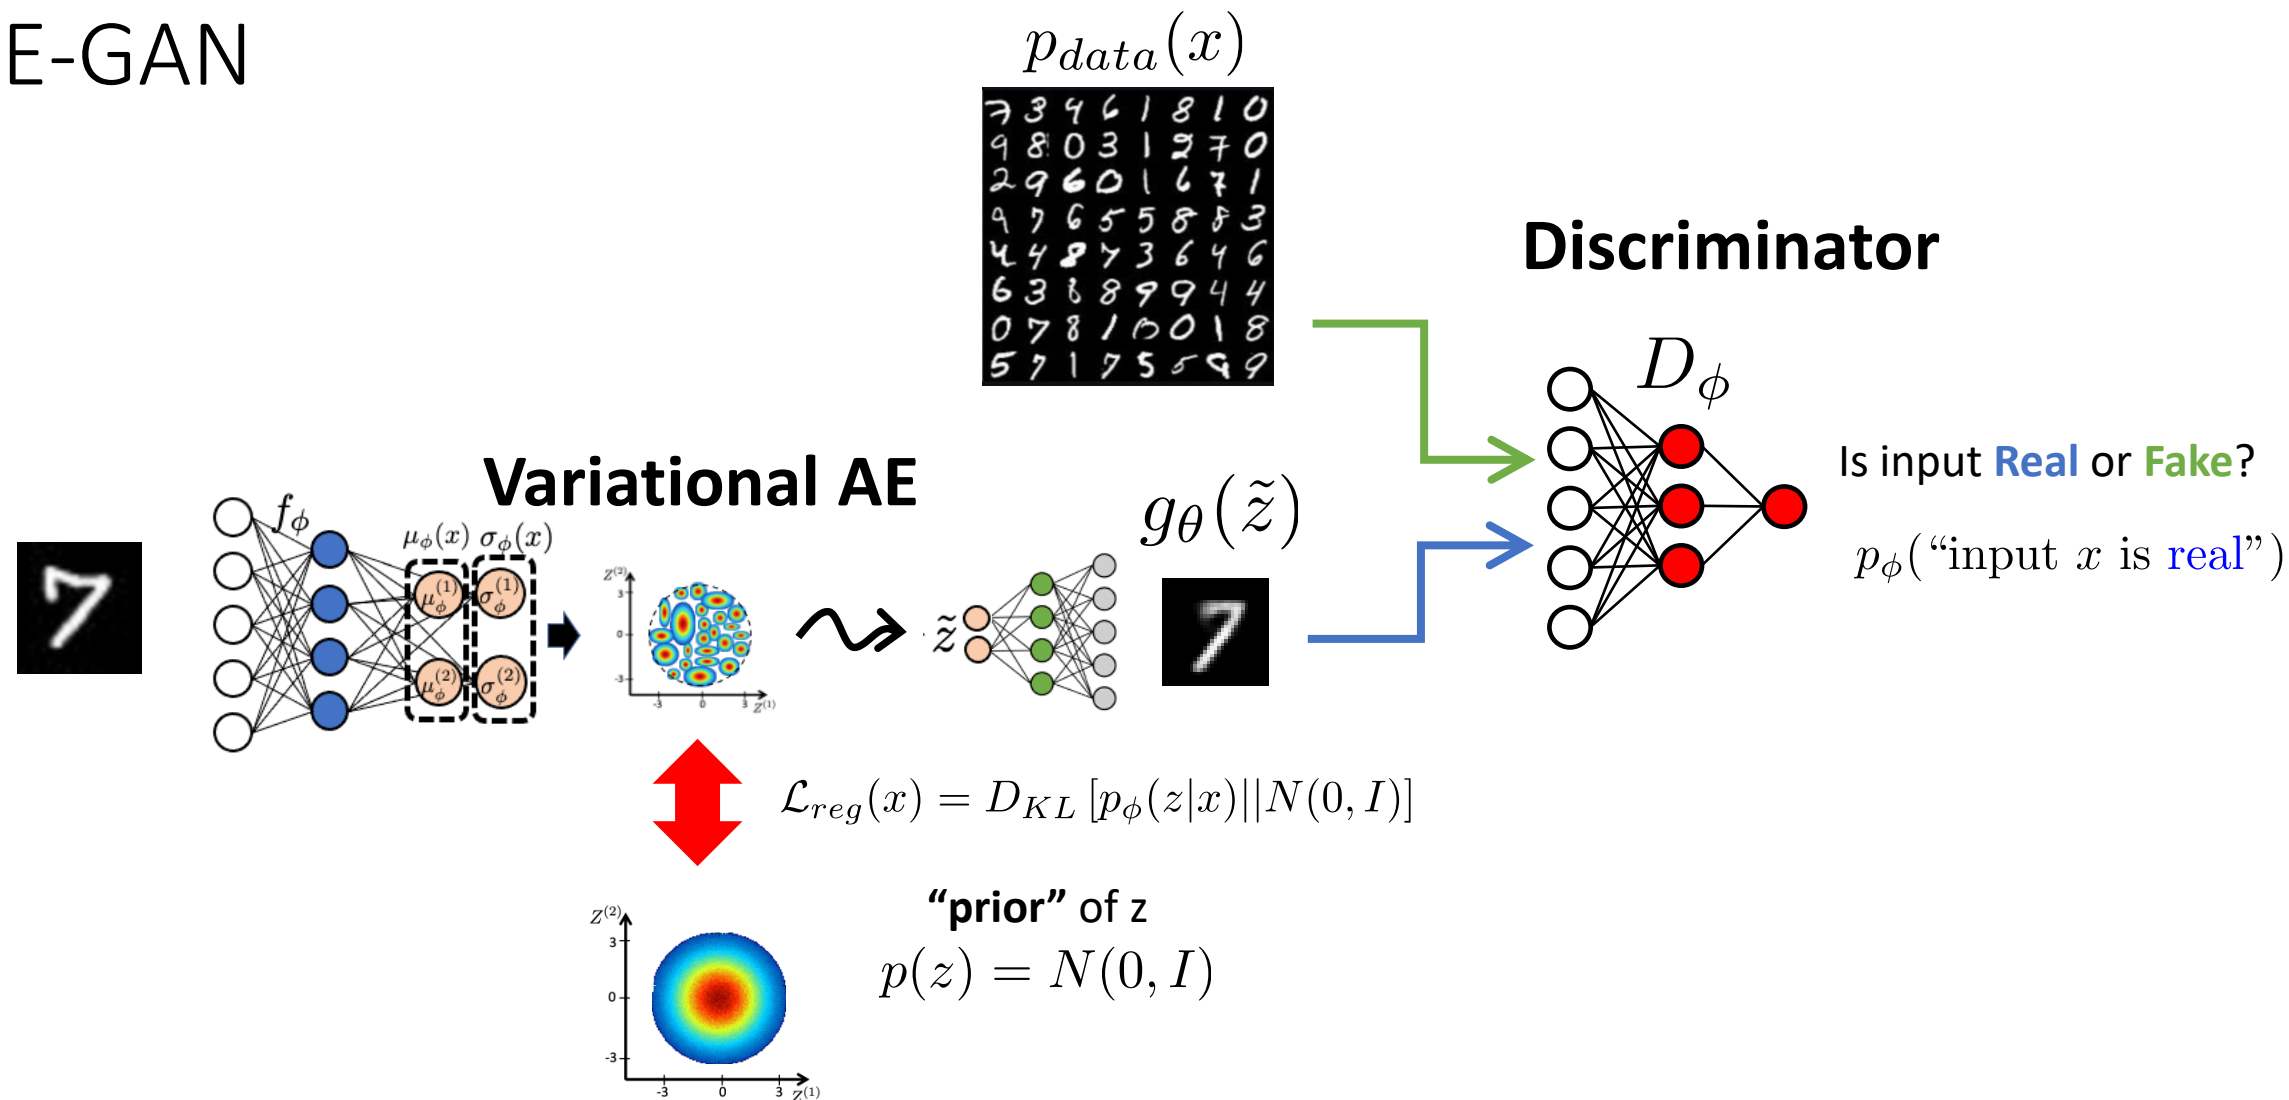
From: Larsen et al, Autoencoding beyond pixels using a learned similarity metric, 2016

VAE-GAN



From: Larsen et al, Autoencoding beyond pixels using a learned similarity metric, 2016

VAE-GAN



From: Larsen et al, Autoencoding beyond pixels using a learned similarity metric, 2016

VAE-GAN

Encoding => Reconstruction


Input



VAE



GAN



VAE/GAN



Generation of new datapoints

VAE



GAN



VAE/GAN



From: Larsen et al, Autoencoding beyond pixels using a learned similarity metric, 2016

Discriminator's Accuracy as “Learned Loss function” for the Generator

$$\mathcal{L}_G(z) = \log(1 - \underline{D}_{\underline{\phi}}(\underline{G}_{\theta}(z)))$$

Theoretical loss

$$\mathcal{L}_G(z) = -\log \underline{D}_{\underline{\phi}}(\underline{G}_{\theta}(z))$$

Practical loss

Other “**engineered**” losses were designed to measure a **specific type** of error. For example:

Reconstruction loss: intensity squared error

VAE regularizer: divergence from prior

Cross Entropy: entropy (difference) between predicted labels and real labels

Dice loss: Overlap of predicted segmentation with real segmentation

Discriminator's accuracy is a “Learned Loss” for the Generator:

D can learn to distinguish generated/real based on any type of features!


Real or Fake?

Day or Night?

Summer or Winter?

Cartoon or Photo?

Style Transfer

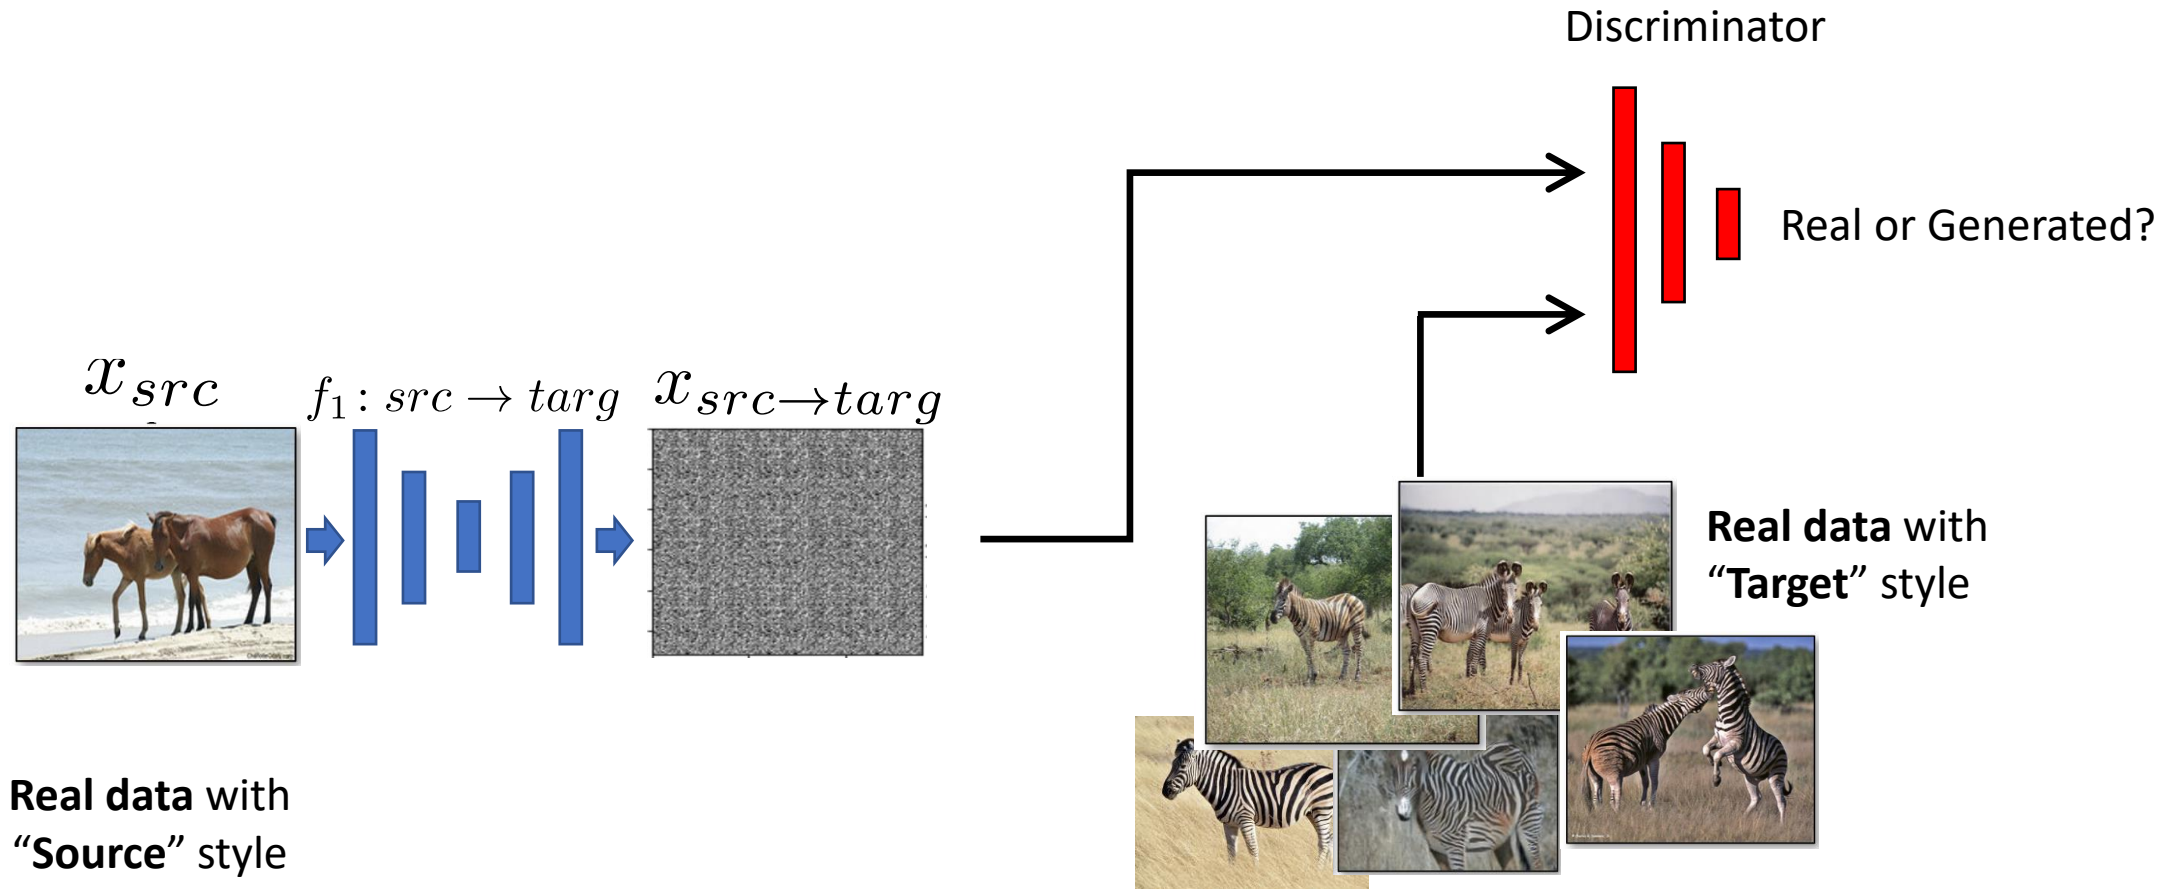


Horse => Zebra video:

<https://www.youtube.com/watch?v=9reHvktowLY>

Cycle GAN:

Converting “Style” of one data distribution to the “style” of another

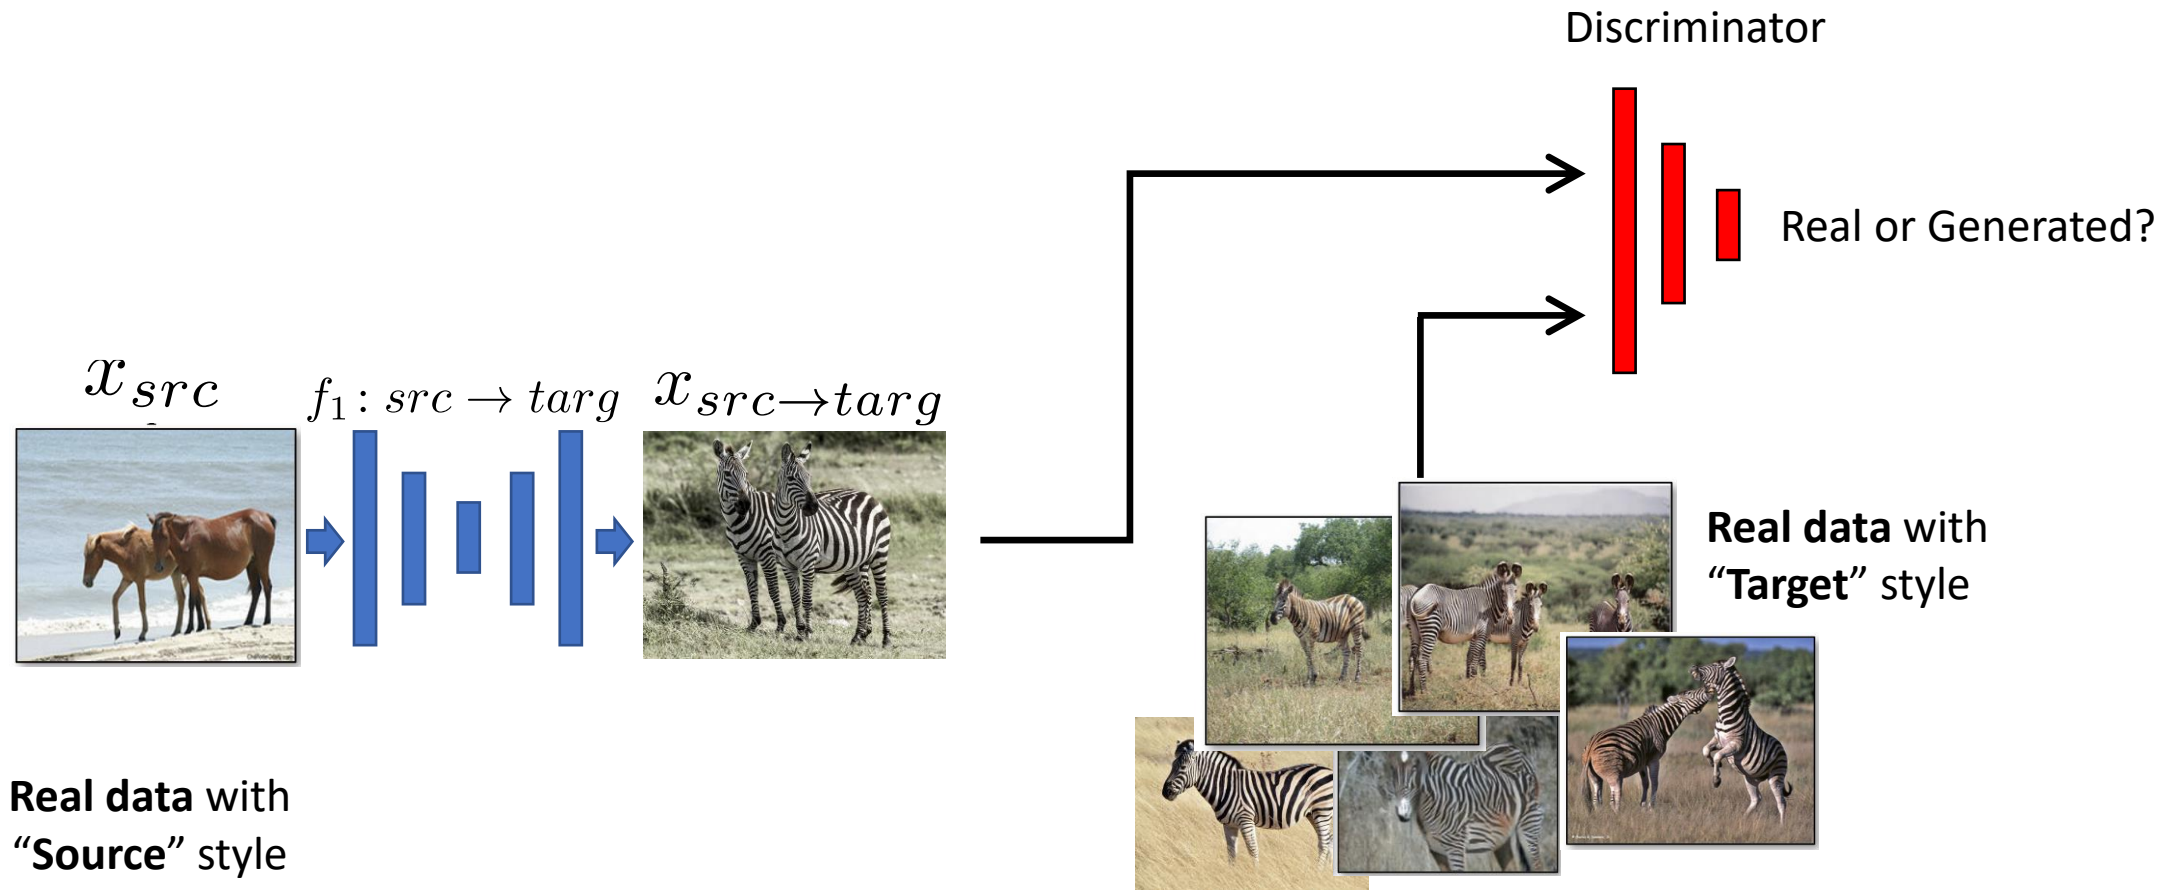


Zhu et al, Unpaired Image-to-Image Translation using Cycle-Consistent Adversarial Networks, 2017

Neural Computation – Konstantinos Kamnitsas

Cycle GAN:

Converting “Style” of one data distribution to the “style” of another

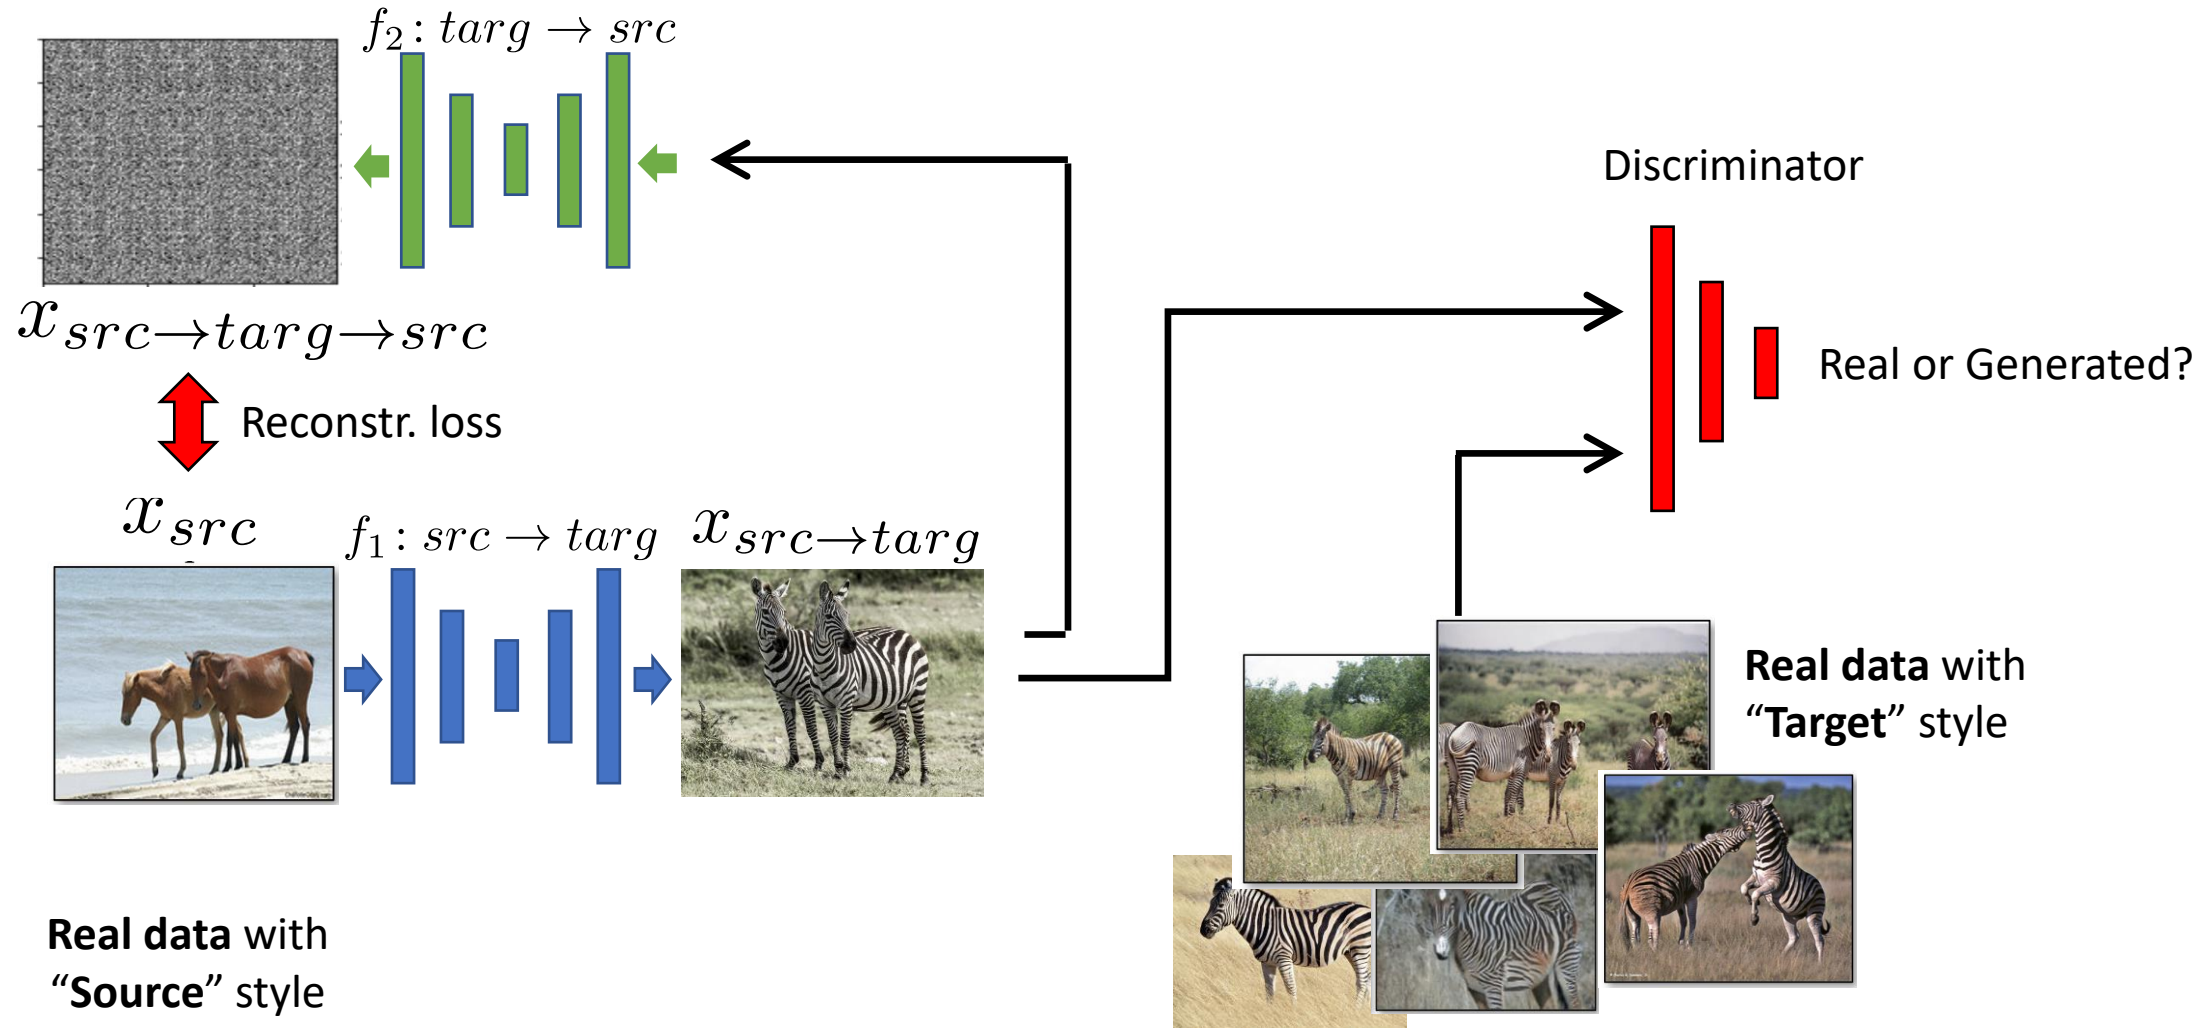


Zhu et al, Unpaired Image-to-Image Translation using Cycle-Consistent Adversarial Networks, 2017

Neural Computation – Konstantinos Kamnitsas

Cycle GAN:

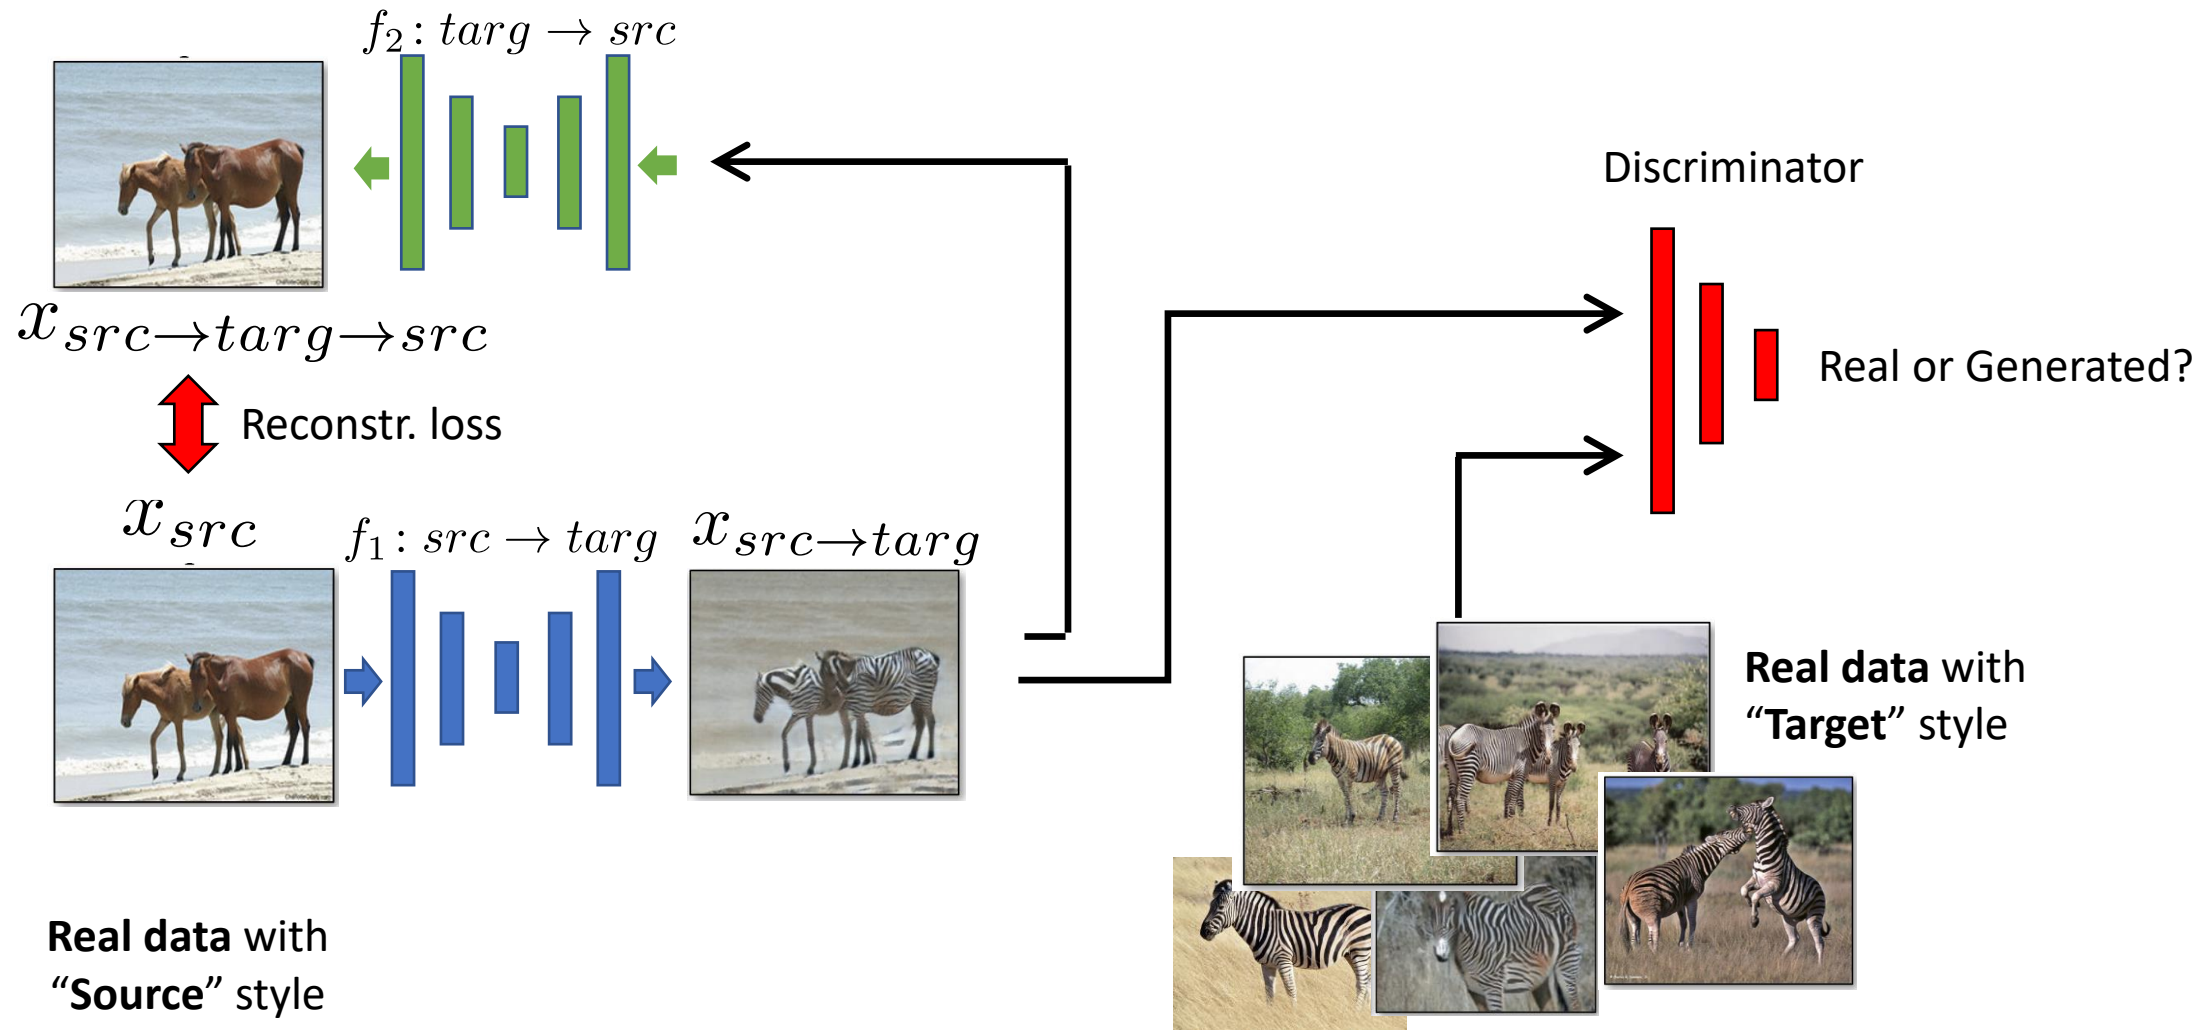
Converting “Style” of one data distribution to the “style” of another



Zhu et al, Unpaired Image-to-Image Translation using Cycle-Consistent Adversarial Networks, 2017

Cycle GAN:

Converting “Style” of one data distribution to the “style” of another



Zhu et al, Unpaired Image-to-Image Translation using Cycle-Consistent Adversarial Networks, 2017

Cycle-GAN: Results

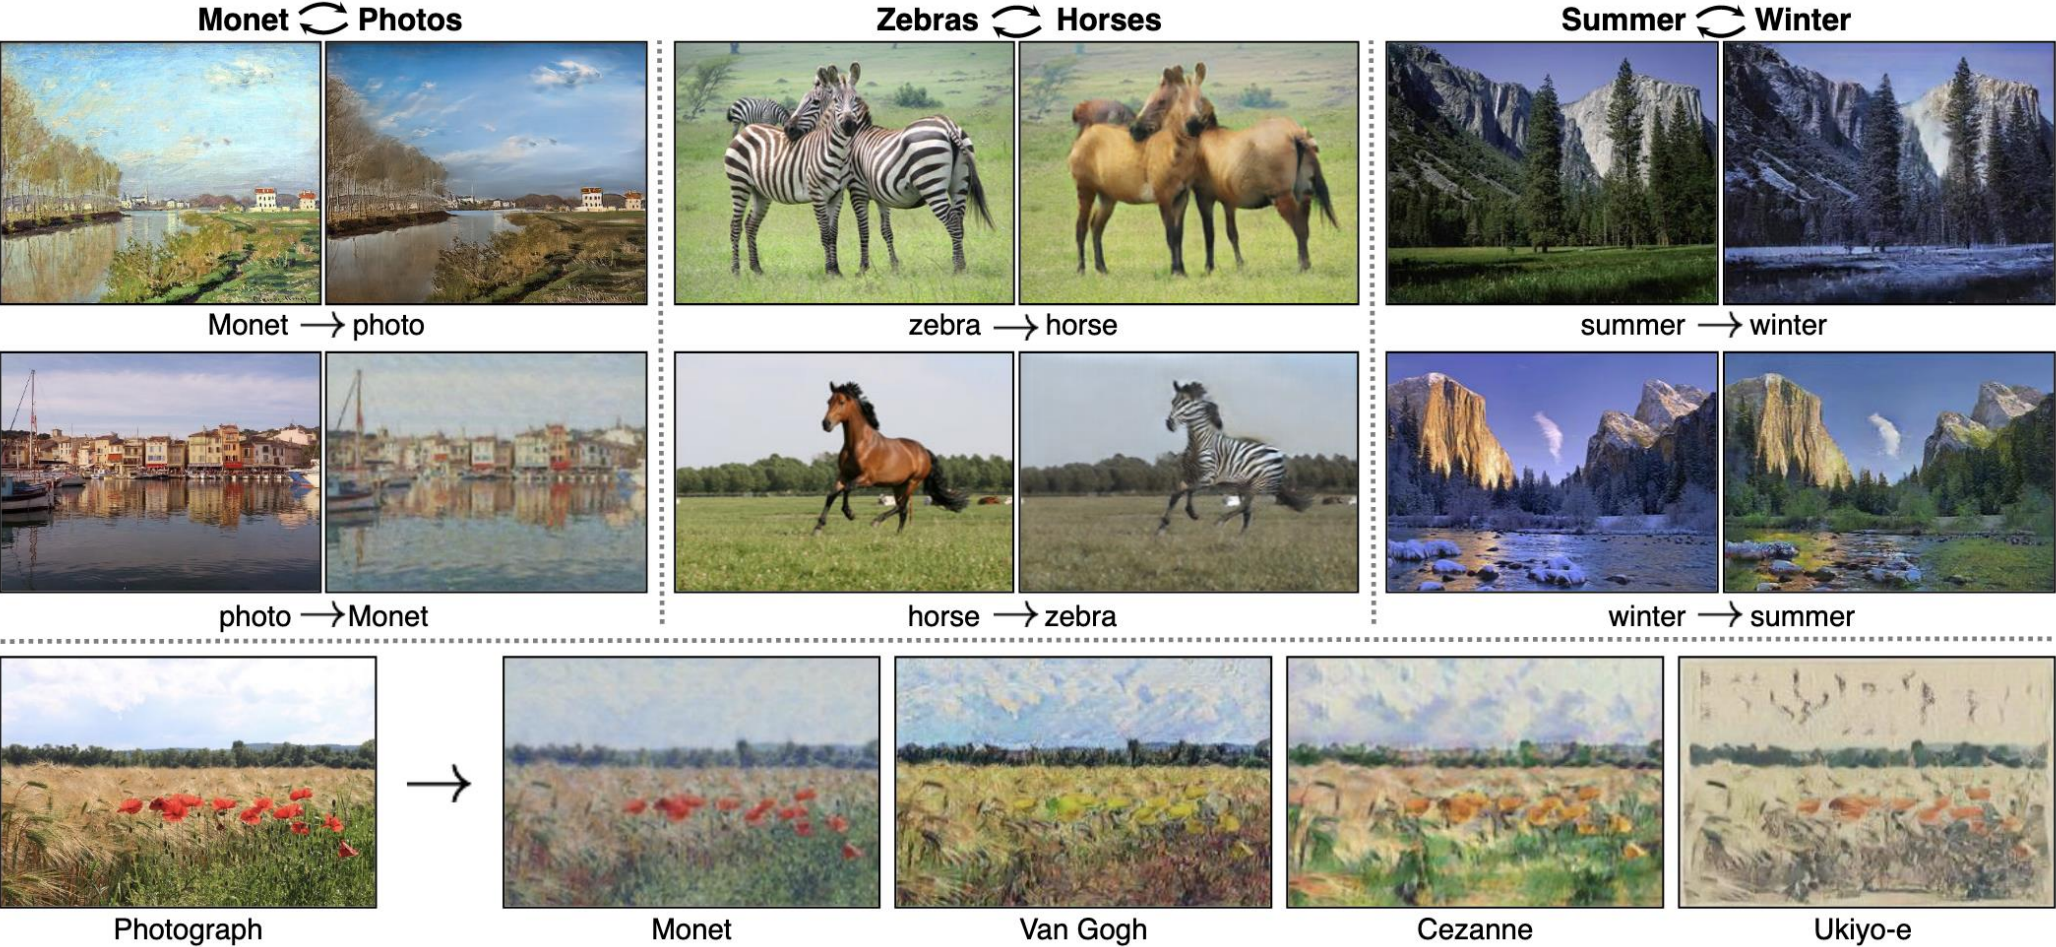
Horse => Zebra video:

<https://www.youtube.com/watch?v=9reHvktowLY>




Turning a horse video into a zebra video (by CycleGAN)

Cycle GAN - Results:



Cycle GAN - Results:



Thank you very much