

# PHYS52015 Core Ib: Introduction to High Performance Computing (HPC)

Session VI: Non-blocking P2P communication

Christopher Marcotte

Michaelmas term 2023

## Outline



Nonblocking point to point communication  
Collective operations

[pollev.com/christophermarcotte820](https://pollev.com/christophermarcotte820)

## Buffers

MPI distinguishes different types of buffers:

- ▶ variables
- ▶ user-level buffers
- ▶ hardware/system buffers

MPI implementations are excellent in tuning communication, i.e. avoid copying, but we have to assume that a message runs through all buffers, then through the network, and then bottom-up through all buffers again. This means that Send and Recv are expensive operations.

Even worse, two concurrent sends might deadlock (but only for massive message counts or extremely large messages).

## Buffers

MPI distinguishes different types of buffers:

- ▶ variables
- ▶ user-level buffers
- ▶ hardware/system buffers

MPI implementations are excellent in tuning communication, i.e. avoid copying, but we have to assume that a message runs through all buffers, then through the network, and then bottom-up through all buffers again. This means that Send and Recv are expensive operations.

Even worse, two concurrent sends might deadlock (but only for massive message counts or extremely large messages).

⇒ One way to deal with this is to allow MPI to optimize the messaging by giving both Send and Recv commands simultaneously — this is a `MPI_Sendrecv`.

```
int MPI_Sendrecv(  
    const void *sendbuf, int sendcount,  
    MPI_Datatype sendtype,  
    int dest, int sendtag,  
    void *recvbuf, int recvcount,  
    MPI_Datatype recvtype,  
    int source, int recvtag,  
    MPI_Comm comm, MPI_Status *status  
)
```

- ▶ Shortcut for send followed by receive
  - ▶ Allows MPI to optimise aggressively
  - ▶ Anticipates that many applications have dedicated compute and data exchange phases
- ⇒ Does not really solve our efficiency concerns, just weaken them

## MPI\_Sendrecv example

We have a program which sends an `nentries`-length buffer between two processes:

```
if (rank == 0) {  
    MPI_Send(sendbuf, nentries, MPI_INT, 1, 0, ...);  
    MPI_Recv(recvbuf, nentries, MPI_INT, 1, 0, ...);  
} else if (rank == 1) {  
    MPI_Send(sendbuf, nentries, MPI_INT, 0, 0, ...);  
    MPI_Recv(recvbuf, nentries, MPI_INT, 0, 0, ...);  
}
```

- Recall that `MPI_Send` behaves like `MPI_Bsend` when buffer space is available, and then behaves like `MPI_Ssend` when it is not.

## MPI\_Sendrecv example

We have a program which sends an `nentries`-length buffer between two processes:

```
if (rank == 0) {
    MPI_Send(sendbuf, nentries, MPI_INT, 1, 0, ...);
    MPI_Recv(recvbuf, nentries, MPI_INT, 1, 0, ...);
} else if (rank == 1) {
    MPI_Send(sendbuf, nentries, MPI_INT, 0, 0, ...);
    MPI_Recv(recvbuf, nentries, MPI_INT, 0, 0, ...);
}
```

- Recall that `MPI_Send` behaves like `MPI_Bsend` when buffer space is available, and then behaves like `MPI_Ssend` when it is not.

```
if (rank == 0) {
    MPI_Sendrecv(sendbuf, nentries, MPI_INT, 1, 0, /* Send */
                recvbuf, nentries, MPI_INT, 1, 0, /* Recv */ ...);
} else if (rank == 1) {
    MPI_Sendrecv(sendbuf, nentries, MPI_INT, 0, 0, /* Send */
                recvbuf, nentries, MPI_INT, 0, 0, /* Recv */ ...);
}
```

## Nonblocking P2P communication

- ▶ Non-blocking commands start with I (immediate return, e.g.)
- ▶ Non-blocking means that operation returns immediately though MPI might not have transferred data (might not even have started)
- ▶ Buffer thus is still in use and we may not overwrite it
- ▶ We explicitly have to validate whether message transfer has completed before we reuse or delete the buffer

```
// Create helper variable (handle)  
int a = 1;  
// trigger the send  
// do some work  
// check whether communication has completed  
a = 2;  
...
```

⇒ We now can overlap communication and computation.



## Why non-blocking...?

- ▶ Added flexibility of separating posting messages from receiving them.
- ⇒ MPI libraries often have optimisations to complete sends quickly if the matching receive already exists.
- ▶ Sending many messages to one process, which receives them all...

## Why non-blocking...?

- ▶ Added flexibility of separating posting messages from receiving them.
- ⇒ MPI libraries often have optimisations to complete sends quickly if the matching receive already exists.
- ▶ Sending many messages to one process, which receives them all...

```
int buf1, buf2;
// ...
for (int k=0; k<100; k++){
    if (rank==0){
        MPI_Recv(&buf2, 1, MPI_INT, k, k, COMM, &status);
    }else{
        MPI_Send(&buf1, 1, MPI_INT, 0, k, COMM);
    }
}
```

- ▶ The receiving process waits on each `MPI_Recv` before moving on, because it is blocking.
- ▶ If we used a non-blocking `MPI_Irecv`, then all can complete as each `MPI_Send` arrives and we just need to `MPI_Wait` for the results.

## Isend & Irecv

- ▶ Non-blocking variants of MPI\_Send and MPI\_Recv
- ▶ Returns immediately, but *buffer is not safe to reuse*

```
int MPI_Isend(const void *buffer, int count, MPI_Datatype dtype,  
             int dest, int tag, MPI_Comm comm, MPI_Request *request);  
int MPI_Irecv(void *buffer, int count, MPI_Datatype dtype,  
             int dest, int tag, MPI_Comm comm, MPI_Request *request);
```

- ▶ Note the request in the send, and the lack of status in recv
- ▶ We need to process that request before we can reuse the buffers

```
int MPI_Send(const void *buffer, int count, MPI_Datatype datatype,
             int dest, int tag, MPI_Comm comm
);

int MPI_Isend(const void *buf, int count, MPI_Datatype datatype,
              int dest, int tag, MPI_Comm comm,
              MPI_Request *request
);

int MPI_Test(MPI_Request *request, int *flag, MPI_Status *status);
int MPI_Wait(MPI_Request *request, MPI_Status *status);
```

- ▶ Pass additional pointer to object of type `MPI_Request`.
- ▶ Non-blocking, i.e. operation returns immediately.
- ▶ Check for send completion with `MPI_Wait` or `MPI_Test`.
- ▶ `MPI_Irecv` analogous.
- ▶ The status object is not required for the receive process, as we have to hand it over to wait or test later.

## Testing

If we have a request, we can check whether the message it corresponds to has been completed with `MPI_Test`:

```
int MPI_Test(MPI_Request *request, int *flag, MPI_Status *status);
```

## Testing

If we have a request, we can check whether the message it corresponds to has been completed with `MPI_Test`:

```
int MPI_Test(MPI_Request *request, int *flag, MPI_Status *status);
```

- ▶ `flag` will be true (an `int` of value 1) if the provided request has been completed, and false otherwise.
- ▶ If we don't want to test for completion, we can instead `MPI_Wait...`

## Waiting...

```
MPI_Request request1, request2;  
MPI_Status status;  
int buffer1[10];    int buffer2[10];  
  
MPI_Send(buffer1, 10, MPI_INT, right, 0, MPI_COMM_WORLD);  
MPI_Recv(buffer2, 10, MPI_INT, left, 0, MPI_COMM_WORLD, &status);  
buffer2[0] = 0;  
  
MPI_Isend(buffer2, 10, MPI_INT, right, 0, MPI_COMM_WORLD, &request1);  
MPI_Irecv(buffer1, 10, MPI_INT, left, 0, MPI_COMM_WORLD, &request2);  
buffer1[0] = 0;
```

There is an error in this code, what change do we need to make for it to be correct?

## Waiting...

```
MPI_Request request1, request2;
MPI_Status status;
int buffer1[10];    int buffer2[10];

MPI_Send(buffer1, 10, MPI_INT, right, 0, MPI_COMM_WORLD);
MPI_Recv(buffer2, 10, MPI_INT, left, 0, MPI_COMM_WORLD, &status);
buffer2[0] = 0;

MPI_Isend(buffer2, 10, MPI_INT, right, 0, MPI_COMM_WORLD, &request1);
MPI_Irecv(buffer1, 10, MPI_INT, left, 0, MPI_COMM_WORLD, &request2);
buffer1[0] = 0;
```

There is an error in this code, what change do we need to make for it to be correct? Before `buffer1[0] = 0;`:

```
MPI_Wait(&request1, &status);
MPI_Wait(&request2, &status);
```



## P2P communication in action

```
MPI_Request request1, request2;
MPI_Status status;
int buffer1[10];    int buffer2[10];
// Variant A
MPI_Recv(buffer1, 10, MPI_INT, left, 0, MPI_COMM_WORLD, &status);
MPI_Send(buffer2, 10, MPI_INT, right, 0, MPI_COMM_WORLD);
// Variant B
//MPI_Send(buffer2, 10, MPI_INT, right, 0, MPI_COMM_WORLD);
//MPI_Recv(buffer1, 10, MPI_INT, left, 0, MPI_COMM_WORLD, &status);
// Variant C
//MPI_Irecv(buffer1, 10, MPI_INT, left, 0, MPI_COMM_WORLD, &request1);
//MPI_Isend(buffer2, 10, MPI_INT, right, 0, MPI_COMM_WORLD, &request2);
//MPI_Wait(&request1, &status);
//MPI_Wait(&request2, &status);
```

- Does Variant A deadlock?

## P2P communication in action

```
MPI_Request request1, request2;
MPI_Status status;
int buffer1[10];    int buffer2[10];
// Variant A
//MPI_Recv(buffer1, 10, MPI_INT, left, 0, MPI_COMM_WORLD, &status);
//MPI_Send(buffer2, 10, MPI_INT, right, 0, MPI_COMM_WORLD);
// Variant B
MPI_Send(buffer2, 10, MPI_INT, right, 0, MPI_COMM_WORLD);
MPI_Recv(buffer1, 10, MPI_INT, left, 0, MPI_COMM_WORLD, &status);
// Variant C
//MPI_Irecv(buffer1, 10, MPI_INT, left, 0, MPI_COMM_WORLD, &request1);
//MPI_Isend(buffer2, 10, MPI_INT, right, 0, MPI_COMM_WORLD, &request2);
//MPI_Wait(&request1, &status);
//MPI_Wait(&request2, &status);
```

- ▶ Does Variant A deadlock? Yes! MPI\_Recv is always blocking.
- ▶ Does Variant B deadlock?

## P2P communication in action

```
MPI_Request request1, request2;
MPI_Status status;
int buffer1[10];  int buffer2[10];
// Variant A
//MPI_Recv(buffer1, 10, MPI_INT, left, 0, MPI_COMM_WORLD, &status);
//MPI_Send(buffer2, 10, MPI_INT, right, 0, MPI_COMM_WORLD);
// Variant B
//MPI_Send(buffer2, 10, MPI_INT, right, 0, MPI_COMM_WORLD);
//MPI_Recv(buffer1, 10, MPI_INT, left, 0, MPI_COMM_WORLD, &status);
// Variant C
MPI_Irecv(buffer1, 10, MPI_INT, left, 0, MPI_COMM_WORLD, &request1);
MPI_Isend(buffer2, 10, MPI_INT, right, 0, MPI_COMM_WORLD, &request2);
MPI_Wait(&request1, &status);
MPI_Wait(&request2, &status);
```

- ▶ Does Variant A deadlock? Yes! MPI\_Recv is always blocking.
- ▶ Does Variant B deadlock? Not for only 10 integers (if not too many messages sent before).
- ▶ Does Variant C deadlock? Is it correct? Is it fast? May we add additional operations before the first wait?

## Concept of building block

- ▶ Content
  - ▶ Introduce `sendrecv`
  - ▶ Introduce concept of non-blocking communication
  - ▶ Study variants of P2P communication w.r.t. blocking and call order
- ▶ Expected Learning Outcomes
  - ▶ The student knows difference of blocking and non-blocking operations
  - ▶ The student can explain the idea of non-blocking communication
  - ▶ The student can write MPI code where communication and computation overlap

## Definition: collective

**Collective operation:** A collective (MPI) operation is an operation involving many/all nodes/ranks.

- ▶ In MPI, a collective operation involves all ranks of one communicator (introduced later)
- ▶ For `MPI_COMM_WORLD`, a collective operation involves all ranks
- ▶ Collectives are blocking (though newer ( $\geq 3.1$ ) MPI standard introduces non-blocking collectives)
- ▶ Blocking collectives always synchronise all ranks, i.e. all ranks have to enter the same collective instruction before any rank proceeds

## A (manual) collective

```
double a;

if (myrank==0) {
    for (int i=1; i<mysize; i++) {
        double tmp;
        MPI_Recv(&tmp,1,MPI_DOUBLE, ...);
        a+=tmp;
    }
}
else {
    MPI_Send(&a,1,MPI_DOUBLE,0, ...);
}
```

What type of collective operation is realised here?

## A (manual) collective

```
double a;

if (myrank==0) {
    for (int i=1; i<mysize; i++) {
        double tmp;
        MPI_Recv(&tmp,1,MPI_DOUBLE, ..., ...);
        a+=tmp;
    }
}
else {
    MPI_Send(&a,1,MPI_DOUBLE,0, ...);
}
```

What type of collective operation is realised here?

```
double globalSum;
MPI_Reduce(&a, &globalSum, 1,
           MPI_DOUBLE, MPI_SUM, 0, MPI_COMM_WORLD);
```

## Flavours of collective operations in MPI

Type of collective	One-to-all	All-to-one	All-to-all
Synchronisation			
Communication			
Computation			

Insert the following MPI operations into the table (MPI prefix and signature neglected):

- ▶ Barrier
- ▶ Broadcast
- ▶ Reduce
- ▶ Allgather
- ▶ Scatter
- ▶ Gather
- ▶ Allreduce

⇒ Synchronisation as discussed is simplest kind of collective operation



## Flavours of collective operations in MPI


Type of collective	One-to-all	All-to-one	All-to-all
Synchronisation	Barrier		
Communication	Broadcast, Scatter	Gather	Allgather
Computation		Reduce	Allreduce

Insert the following MPI operations into the table (MPI prefix and signature neglected):

- ▶ Barrier
- ▶ Broadcast
- ▶ Reduce
- ▶ Allgather
- ▶ Scatter
- ▶ Gather
- ▶ Allreduce

⇒ Synchronisation as discussed is simplest kind of collective operation

## Good reasons to use MPI's collective



- ▶ Simplicity of code
- ▶ Performance through specialised implementations
- ▶ Support through dedicated hardware (cf. BlueGene's three network topologies: clique, fat tree, ring)

## MPI\_Barrier

- ▶ Simplest form of collective operation — synchronization of all ranks in `comm`.
  - ▶ Rarely used:
- ⇒ MPI\_Barrier doesn't synchronize non-blocking calls
- ⇒ Really meant for telling MPI about calls *outside* MPI, like IO

```
int rank, size;
MPI_Comm_rank(MPI_COMM_WORLD, &rank);
MPI_Comm_size(MPI_COMM_WORLD, &size);
for ( int ii = 0; ii < size; ++ii ) {
    if ( rank == ii ) {
        // my turn to write to the file
        writeStuffToTheFile();
    }
    MPI_Barrier(MPI_COMM_WORLD);
}
```

## MPI\_Bcast & MPI\_Scatter

- ▶ MPI\_Bcast sends the contents of a buffer from `root` to all other processes.
- ▶ MPI\_Scatter sends *parts* of a buffer from `root` to different processes.
- ▶ MPI\_Bcast is the inverse of MPI\_Reduce
- ▶ MPI\_Scatter is the inverse of MPI\_Gather

```
MPI_Comm comm;
int array[100];
int root=0;
//...
MPI_Bcast(array, 100, MPI_INT, root, comm);

int gsize, *sendbuf;
int rbuf[100];
//...
MPI_Comm_size(comm, &gsize);
sendbuf = (int *)malloc(gsize*100*sizeof(int));
//...
MPI_Scatter(sendbuf, 100, MPI_INT, rbuf, 100, MPI_INT, root, comm);
```

## MPI\_Reduce & MPI\_Gather

- ▶ MPI\_Reduce reduces a value across ranks to a single value on root using a prescribed reduction operator.
- ▶ MPI\_Gather concatenates the array pieces from all processes onto the root process.

```
MPI_Comm comm;
int sum, root=0;
//...
MPI_Reduce(sum, c, 1, MPI_INT, MPI_SUM, root, comm);

int gsize, sendarray[100];
int myrank, *rbuf;
//...
MPI_Comm_rank(comm, &myrank);
if (myrank == root) {
    MPI_Comm_size(comm, &gsize);
    rbuf = (int *)malloc(gsize*100*sizeof(int));
}
MPI_Gather(sendarray, 100, MPI_INT, rbuf, 100, MPI_INT, root, comm);
```

## MPI\_Allgather & MPI\_Allreduce

- ▶ MPI\_Allgather is an MPI\_Gather which concatenates the array pieces on all processes.
- ▶ MPI\_Allreduce is an MPI\_Reduce which reduces on all processes.

```
MPI_Comm comm;
int sum, root=0;
//...
MPI_Allreduce(sum, c, 1, MPI_INT, MPI_SUM, comm);

int gsize, sendarray[100];
int *rbuf;
//...
MPI_Comm_size(comm, &gsize);
rbuf = (int *)malloc(gsize*100*sizeof(int));
MPI_Allgather(sendarray, 100, MPI_INT, rbuf, 100, MPI_INT, comm);
```

- ▶ MPI\_Allgather is used a lot (for better or worse) debugging distributed arrays — serial checks work on one process!
- ▶ MPI\_Allreduce is particularly useful for *convergence* checks — we will see this when we return to the Jacobi iteration problem in MPI.

## Concept of building block

- ▶ Content
  - ▶ Classify different collectives from MPI
  - ▶ Understand use cases for main forms of MPI collectives
- ▶ Expected Learning Outcomes
  - ▶ The student knows which type of collectives do exist (\*)
  - ▶ The student can explain what collectives do (\*)
  - ▶ The student can identify collective code fragments (\*)
  - ▶ The student can use collectives or implement them manually